

Chapter 4: Distributions of random variables

OpenIntro Statistics, 4th Edition

Slides developed by Mine Çetinkaya-Rundel of OpenIntro.

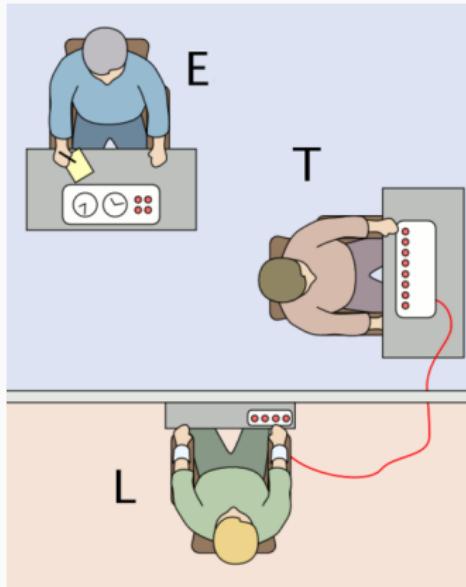
The slides may be copied, edited, and/or shared via the CC BY-SA license.

Some images may be included under fair use guidelines (educational purposes).

Geometric distribution

Milgram experiment

- Stanley Milgram, a Yale University psychologist, conducted a series of experiments on obedience to authority starting in 1963.
- Experimenter (E) orders the teacher (T), the subject of the experiment, to give severe electric shocks to a learner (L) each time the learner answers a question incorrectly.
- The learner is actually an actor, and the electric shocks are not real, but a prerecorded sound is played each time the teacher administers an electric shock.



[http://en.wikipedia.org/wiki/File:
Milgram_Experiment_v2.png](http://en.wikipedia.org/wiki/File:Milgram_Experiment_v2.png)

Milgram experiment (cont.)

- These experiments measured the willingness of study participants to obey an authority figure who instructed them to perform acts that conflicted with their personal conscience.
- Milgram found that about 65% of people would obey authority and give such shocks.
- Over the years, additional research suggested this number is approximately consistent across communities and time.

Bernoulli random variables

- Each person in Milgram's experiment can be thought of as a *trial*.
- A person is labeled a *success* if she refuses to administer a severe shock, and *failure* if she administers such shock.
- Since only 35% of people refused to administer a shock, *probability of success* is $p = 0.35$.
- When an individual trial has only two possible outcomes, it is called a *Bernoulli random variable*.

Geometric distribution

Dr. Smith wants to repeat Milgram's experiments but she only wants to sample people until she finds someone who will not inflict a severe shock. What is the probability that she stops after the first person?

$$P(1^{\text{st}} \text{ person refuses}) = 0.35$$

... the third person?

$$P(1^{\text{st}} \text{ and } 2^{\text{nd}} \text{ shock, } 3^{\text{rd}} \text{ refuses}) = \frac{S}{0.65} \times \frac{S}{0.65} \times \frac{R}{0.35} = 0.65^2 \times 0.35 \approx 0.15$$

... the tenth person?

$$P(9 \text{ shock, } 10^{\text{th}} \text{ refuses}) = \underbrace{\frac{S}{0.65} \times \cdots \times \frac{S}{0.65}}_{9 \text{ of these}} \times \frac{R}{0.35} = 0.65^9 \times 0.35 \approx 0.0072$$

Geometric distribution (cont.)

Geometric distribution describes the waiting time until a success for *independent and identically distributed (iid)* Bernoulli random variables.

- independence: outcomes of trials don't affect each other
- identical: the probability of success is the same for each trial

Geometric probabilities

If p represents probability of success, $(1 - p)$ represents probability of failure, and n represents number of independent trials

$$P(\text{success on the } n^{\text{th}} \text{ trial}) = (1 - p)^{n-1} p$$

Can we calculate the probability of rolling a 6 for the first time on the 6th roll of a die using the geometric distribution? Note that what was a success (rolling a 6) and what was a failure (not rolling a 6) are clearly defined and one or the other must happen for each trial.

- (a) no, on the roll of a die there are more than 2 possible outcomes
- (b) *yes, why not*

$$P(6 \text{ on the } 6^{\text{th}} \text{ roll}) = \left(\frac{5}{6}\right)^5 \left(\frac{1}{6}\right) \approx 0.067$$

Expected value

How many people is Dr. Smith expected to test before finding the first one that refuses to administer the shock?

The expected value, or the mean, of a geometric distribution is defined as $\frac{1}{p}$.

$$\mu = \frac{1}{p} = \frac{1}{0.35} = 2.86$$

She is expected to test 2.86 people before finding the first one that refuses to administer the shock.

But how can she test a non-whole number of people?

Expected value and its variability

Mean and standard deviation of geometric distribution

$$\mu = \frac{1}{p} \quad \sigma = \sqrt{\frac{1-p}{p^2}}$$

- Going back to Dr. Smith's experiment:

$$\sigma = \sqrt{\frac{1-p}{p^2}} = \sqrt{\frac{1-0.35}{0.35^2}} = 2.3$$

- Dr. Smith is expected to test 2.86 people before finding the first one that refuses to administer the shock, give or take 2.3 people.
- These values only make sense in the context of repeating the experiment many many times.

Binomial distribution

Suppose we randomly select four individuals to participate in this experiment. What is the probability that exactly 1 of them will refuse to administer the shock?

Let's call these people Allen (A), Brittany (B), Caroline (C), and Damian (D). Each one of the four scenarios below will satisfy the condition of "exactly 1 of them refuses to administer the shock":

Suppose we randomly select four individuals to participate in this experiment. What is the probability that exactly 1 of them will refuse to administer the shock?

Let's call these people Allen (A), Brittany (B), Caroline (C), and Damian (D). Each one of the four scenarios below will satisfy the condition of "exactly 1 of them refuses to administer the shock":

Scenario 1: $\frac{0.35}{(\text{A}) \text{refuse}} \times \frac{0.65}{(\text{B}) \text{shock}} \times \frac{0.65}{(\text{C}) \text{shock}} \times \frac{0.65}{(\text{D}) \text{shock}} = 0.0961$

Scenario 2: $\frac{0.65}{(\text{A}) \text{shock}} \times \frac{0.35}{(\text{B}) \text{refuse}} \times \frac{0.65}{(\text{C}) \text{shock}} \times \frac{0.65}{(\text{D}) \text{shock}} = 0.0961$

Scenario 3: $\frac{0.65}{(\text{A}) \text{shock}} \times \frac{0.65}{(\text{B}) \text{shock}} \times \frac{0.35}{(\text{C}) \text{refuse}} \times \frac{0.65}{(\text{D}) \text{shock}} = 0.0961$

Scenario 4: $\frac{0.65}{(\text{A}) \text{shock}} \times \frac{0.65}{(\text{B}) \text{shock}} \times \frac{0.65}{(\text{C}) \text{shock}} \times \frac{0.35}{(\text{D}) \text{refuse}} = 0.0961$

The probability of exactly one of 4 people refusing to administer the shock is the sum of all of these probabilities.

$$0.0961 + 0.0961 + 0.0961 + 0.0961 = 4 \times 0.0961 = 0.3844$$

Binomial distribution

The question from the prior slide asked for the probability of given number of successes, k , in a given number of trials, n , ($k = 1$ success in $n = 4$ trials), and we calculated this probability as

$$\# \text{ of scenarios} \times P(\text{single scenario})$$

- $\# \text{ of scenarios}$: there is a less tedious way to figure this out, we'll get to that shortly...
- $P(\text{single scenario}) = p^k (1 - p)^{(n-k)}$

probability of success to the power of number of successes, probability of failure to the power of number of failures

The *Binomial distribution* describes the probability of having exactly k successes in n independent Bernouilli trials with probability of success p .

Counting the # of scenarios

Earlier we wrote out all possible scenarios that fit the condition of exactly one person refusing to administer the shock. If n was larger and/or k was different than 1, for example, $n = 9$ and $k = 2$:

RR SSSSSSS

S RR SSSSSS

SS RR SSSSS

...

SS R SS R SSS

...

SSSSSSS RR

writing out all possible scenarios would be incredibly tedious and prone to errors.

Calculating the # of scenarios

Choose function

The *choose function* is useful for calculating the number of ways to choose k successes in n trials.

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}$$

- $k = 1, n = 4: \binom{4}{1} = \frac{4!}{1!(4-1)!} = \frac{4 \times 3 \times 2 \times 1}{1 \times (3 \times 2 \times 1)} = 4$
- $k = 2, n = 9: \binom{9}{2} = \frac{9!}{2!(9-1)!} = \frac{9 \times 8 \times 7!}{2 \times 1 \times 7!} = \frac{72}{2} = 36$

Note: You can also use R for these calculations:

```
> choose(9,2)  
[1] 36
```

Properties of the choose function

Which of the following is false?

- (a) There are n ways of getting 1 success in n trials, $\binom{n}{1} = n$.
- (b) There is only 1 way of getting n successes in n trials, $\binom{n}{n} = 1$.
- (c) There is only 1 way of getting n failures in n trials, $\binom{n}{0} = 1$.
- (d) *There are $n - 1$ ways of getting $n - 1$ successes in n trials,* $\binom{n}{n-1} = n - 1$.

Binomial distribution (cont.)

Binomial probabilities

If p represents probability of success, $(1 - p)$ represents probability of failure, n represents number of independent trials, and k represents number of successes

$$P(k \text{ successes in } n \text{ trials}) = \binom{n}{k} p^k (1 - p)^{(n-k)}$$

Which of the following is not a condition that needs to be met for the binomial distribution to be applicable?

- (a) the trials must be independent
- (b) the number of trials, n , must be fixed
- (c) each trial outcome must be classified as a *success* or a *failure*
- (d) *the number of desired successes, k , must be greater than the number of trials*
- (e) the probability of success, p , must be the same for each trial

A 2012 Gallup survey suggests that 26.2% of Americans are obese. Among a random sample of 10 Americans, what is the probability that exactly 8 are obese?

- (a) pretty high
- (b) *pretty low*

Gallup: <http://www.gallup.com/poll/160061/obesity-rate-stable-2012.aspx>, January 23, 2013.

A 2012 Gallup survey suggests that 26.2% of Americans are obese. Among a random sample of 10 Americans, what is the probability that exactly 8 are obese?

- (a) $0.262^8 \times 0.738^2$
- (b) $\binom{8}{10} \times 0.262^8 \times 0.738^2$
- (c) $\binom{10}{8} \times 0.262^8 \times 0.738^2 = 45 \times 0.262^8 \times 0.738^2 = 0.0005$
- (d) $\binom{10}{8} \times 0.262^2 \times 0.738^8$

The birthday problem

What is the probability that 2 randomly chosen people share a birthday?

Pretty low, $\frac{1}{365} \approx 0.0027$.

What is the probability that at least 2 people out of 366 people share a birthday?

Exactly 1! (Excluding the possibility of a leap year birthday.)

The birthday problem (cont.)

What is the probability that at least 2 people (1 match) out of 121 people share a birthday?

Somewhat complicated to calculate, but we can think of it as the complement of the probability that there are no matches in 121 people.

$$\begin{aligned}P(\text{no matches}) &= 1 \times \left(1 - \frac{1}{365}\right) \times \left(1 - \frac{2}{365}\right) \times \cdots \times \left(1 - \frac{120}{365}\right) \\&= \frac{365 \times 364 \times \cdots \times 245}{365^{121}} \\&= \frac{365!}{365^{121} \times (365 - 121)!} \\&= \frac{121! \times \binom{365}{121}}{365^{121}} \approx 0\end{aligned}$$

$$P(\text{at least 1 match}) \approx 1$$

Expected value

A 2012 Gallup survey suggests that 26.2% of Americans are obese.

Among a random sample of 100 Americans, how many would you expect to be obese?

- Easy enough, $100 \times 0.262 = 26.2$.
- Or more formally, $\mu = np = 100 \times 0.262 = 26.2$.
- But this doesn't mean in every random sample of 100 people exactly 26.2 will be obese. In fact, that's not even possible. In some samples this value will be less, and in others more. How much would we expect this value to vary?

Expected value and its variability

Mean and standard deviation of binomial distribution

$$\mu = np \qquad \sigma = \sqrt{np(1 - p)}$$

- Going back to the obesity rate:

$$\sigma = \sqrt{np(1 - p)} = \sqrt{100 \times 0.262 \times 0.738} \approx 4.4$$

- We would expect 26.2 out of 100 randomly sampled Americans to be obese, with a standard deviation of 4.4.

Note: Mean and standard deviation of a binomial might not always be whole numbers, and that is alright, these values represent what we would expect to see on average.

Unusual observations

Using the notion that *observations that are more than 2 standard deviations away from the mean are considered unusual* and the mean and the standard deviation we just computed, we can calculate a range for the plausible number of obese Americans in random samples of 100.

$$26.2 \pm (2 \times 4.4) = (17.4, 35)$$

An August 2012 Gallup poll suggests that 13% of Americans think home schooling provides an excellent education for children. Would a random sample of 1,000 Americans where only 100 share this opinion be considered unusual?

(a) No

(b) Yes

	Excellent	Good	Only fair	Poor	Total excellent/good
	%	%	%	%	%
Independent private school	31	47	13	2	78
Parochial or church-related schools	21	48	18	5	69
Charter schools	17	43	23	5	60
Home schooling	13	33	30	14	46
Public schools	5	32	42	19	37

Gallup, Aug. 9-12, 2012

An August 2012 Gallup poll suggests that 13% of Americans think home schooling provides an excellent education for children. Would a random sample of 1,000 Americans where only 100 share this opinion be considered unusual?

$$\mu = np = 1,000 \times 0.13 = 130$$

$$\sigma = \sqrt{np(1-p)} = \sqrt{1,000 \times 0.13 \times 0.87} \approx 10.6$$

Method 1: Range of usual observations: $130 \pm 2 \times 10.6 = (108.8, 151.2)$

100 is outside this range, so would be considered unusual.

Method 2: Z-score of observation: $Z = \frac{x - \text{mean}}{\text{SD}} = \frac{100 - 130}{10.6} = -2.83$

100 is more than 2 SD below the mean, so would be considered unusual.

<http://www.gallup.com/poll/156974/private-schools-top-marks-educating-children.aspx>

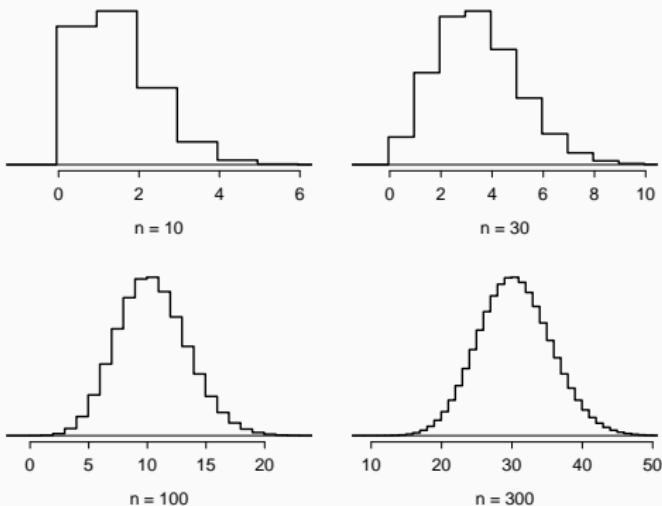
Shapes of binomial distributions

For this activity you will use a web applet. Go to https://gallery.shinyapps.io/dist_calc/ and choose Binomial coin experiment in the drop down menu on the left.

- Set the number of trials to 20 and the probability of success to 0.15. Describe the shape of the distribution of number of successes.
- Keeping p constant at 0.15, determine the minimum sample size required to obtain a unimodal and symmetric distribution of number of successes. Please submit only one response per team.
- Further considerations:
 - What happens to the shape of the distribution as n stays constant and p changes?
 - What happens to the shape of the distribution as p stays constant and n changes?

Distributions of number of successes

Hollow histograms of samples from the binomial model where $p = 0.10$ and $n = 10, 30, 100$, and 300 . What happens as n increases?



Low large is large enough?

The sample size is considered large enough if the expected number of successes and failures are both at least 10.

$$np \geq 10 \quad \text{and} \quad n(1 - p) \geq 10$$

$$10 \times 0.13 = 1.3; 10 \times (1 - 0.13) = 8.7$$

Below are four pairs of Binomial distribution parameters. Which distribution can be approximated by the normal distribution?

- (a) $n = 100, p = 0.95$
- (b) $n = 25, p = 0.45 \rightarrow 25 \times 0.45 = 11.25; 25 \times 0.55 = 13.75$
- (c) $n = 150, p = 0.05$
- (d) $n = 500, p = 0.015$

An analysis of Facebook users

A recent study found that “Facebook users get more than they give”.

For example:

- 40% of Facebook users in our sample made a friend request, but 63% received at least one request
- Users in our sample pressed the like button next to friends' content an average of 14 times, but had their content “liked” an average of 20 times
- Users sent 9 personal messages, but received 12
- 12% of users tagged a friend in a photo, but 35% were themselves tagged in a photo

Any guesses for how this pattern can be explained?

Power users contribute much more content than the typical user.

This study also found that approximately 25% of Facebook users are considered power users. The same study found that the average Facebook user has 245 friends. What is the probability that the average Facebook user with 245 friends has 70 or more friends who would be considered power users? Note any assumptions you must make.

We are given that $n = 245$, $p = 0.25$, and we are asked for the probability $P(K \geq 70)$. To proceed, we need independence, which we'll assume but could check if we had access to more Facebook data.

$$\begin{aligned}P(X \geq 70) &= P(K = 70 \text{ or } K = 71 \text{ or } K = 72 \text{ or } \dots \text{ or } K = 245) \\&= P(K = 70) + P(K = 71) + P(K = 72) + \dots + P(K = 245)\end{aligned}$$

This seems like an awful lot of work...

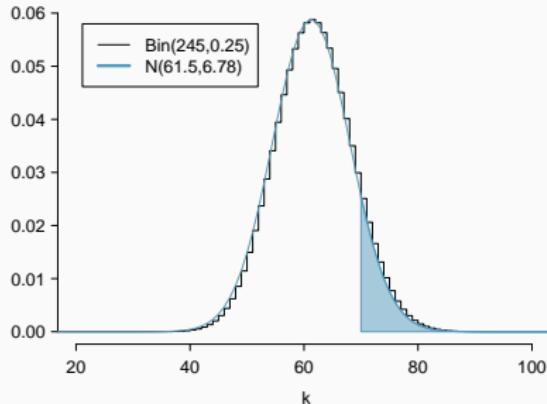
Normal approximation to the binomial

When the sample size is large enough, the binomial distribution with parameters n and p can be approximated by the normal model with parameters $\mu = np$ and $\sigma = \sqrt{np(1 - p)}$.

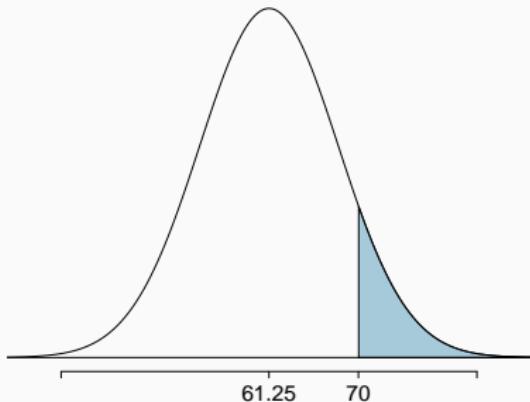
- In the case of the Facebook power users, $n = 245$ and $p = 0.25$.

$$\mu = 245 \times 0.25 = 61.25 \quad \sigma = \sqrt{245 \times 0.25 \times 0.75} = 6.78$$

- $Bin(n = 245, p = 0.25) \approx N(\mu = 61.25, \sigma = 6.78)$.



What is the probability that the average Facebook user with 245 friends has 70 or more friends who would be considered power users?



$$Z = \frac{obs - mean}{SD} = \frac{70 - 61.25}{6.78} = 1.29$$

$$P(Z > 1.29) = 1 - 0.9015 = 0.0985$$

```
> pnorm(1.29)  
[1] 0.9014747
```

The normal approximation breaks down on small intervals

- The normal approximation to the binomial distribution tends to perform poorly when estimating the probability of a small range of counts, even when the conditions are met.
- This approximation for intervals of values is usually improved if cutoff values are extended by 0.5 in both directions.
- The tip to add extra area when applying the normal approximation is most often useful when examining a range of observations. While it is possible to also apply this correction when computing a tail area, the benefit of the modification usually disappears since the total interval is typically quite wide.

Negative binomial distribution

Negative binomial distribution

- The *negative binomial distribution* describes the probability of observing the k^{th} success on the n^{th} trial.
- The following four conditions are useful for identifying a negative binomial case:
 1. The trials are independent.
 2. Each trial outcome can be classified as a success or failure.
 3. The probability of success (p) is the same for each trial.
 4. The last trial must be a success.

Note that the first three conditions are common to the binomial distribution.

Negative binomial distribution

$$P(k^{th} \text{ success on the } n^{th} \text{ trial}) = \binom{n-1}{k-1} p^k (1-p)^{n-k},$$

where p is the probability that an individual trial is a success. All trials are assumed to be independent.

A college student working at a psychology lab is asked to recruit 10 couples to participate in a study. She decides to stand outside the student center and ask every 5th person leaving the building whether they are in a relationship and, if so, whether they would like to participate in the study with their significant other. Suppose the probability of finding such a person is 10%. What is the probability that she will need to ask 30 people before she hits her goal?

Given: $p = 0.10$, $k = 10$, $n = 30$. We are asked to find the probability of 10th success on the 30th trial, therefore we use the negative binomial distribution.

$$\begin{aligned}P(10^{\text{th}} \text{ success on the } 30^{\text{th}} \text{ trial}) &= \binom{29}{9} \times 0.10^{10} \times 0.90^{20} \\&= 10,015,005 \times 0.10^{10} \times 0.90^{20} \\&= 0.00012\end{aligned}$$

Binomial vs. negative binomial

How is the negative binomial distribution different from the binomial distribution?

- In the binomial case, we typically have a fixed number of trials and instead consider the number of successes.
- In the negative binomial case, we examine how many trials it takes to observe a fixed number of successes and require that the last observation be a success.

Practice

Which of the following describes a case where we would use the negative binomial distribution to calculate the desired probability?

- (a) Probability that a 5 year old boy is taller than 42 inches.
- (b) Probability that 3 out of 10 softball throws are successful.
- (c) Probability of being dealt a straight flush hand in poker.
- (d) Probability of missing 8 shots before the first hit.
- (e) *Probability of hitting the ball for the 3rd time on the 8th try.*