

DSCI351-351m-451: Class 11a-p Sample Size, Effect Size and Stat Power

Profs: R. H. French, L. S. Bruckman, P. Leu, K. Davis, S. Cirlos

TAs: W. Oltjen, K. Hernandez, M. Li, M. Li, D. Colvin

08 November, 2022

Contents

10.2.3.1	Class Readings, Assignments, Syllabus Topics	1
10.2.3.1.1	Reading, Lab Exercises, SemProjects	1
10.2.3.2	Power Analysis, Sample Size, Effect Size	
	2
10.2.3.2.1	Review of NHST	
	2
10.2.3.2.2	Controversy surrounding null hypothesis significance testing	
	4
10.2.3.3	Implementing power analysis with the pwr package	
	6
10.2.3.3.1	t-tests	
	6
10.2.3.3.2	ANOVA	
	9
10.2.3.3.3	Correlations	
	9
10.2.3.3.4	Linear Models	
	10
10.2.3.3.5	Tests of Propotions	
	12
10.2.3.3.6	χ^2 tests	
	13
10.2.3.3.7	Samples Sizes in Anova Analysis	
	16
10.2.3.4	Creating Power Analysis Plots	
	16
10.2.3.5	Links	
	18

10.2.3.1 Class Readings, Assignments, Syllabus Topics

10.2.3.1.1 Reading, Lab Exercises, SemProjects

- Readings:
 - For today: OIS 8
 - For next class: ISLR 1, 2.1, 2.2
- Laboratory Exercises:

- LE 5: is due Thursday November 10th
- LE 6: Will be posted Thursday November 10th
- Office Hours: (Class Canvas Calendar for Zoom Link)
 - Wednesday @ 4:00 PM to 5:00 PM, Will Oltjen
 - Saturday @ 3:00 PM to 4:00 PM, Kristen Hernandez
 - **Office Hours are on Zoom, and recorded**
- Semester Projects
 - DSCI 451 Students Biweekly Update #6 Due Friday Nov. 18th
 - DSCI 451 Students
 - * Next **Report Out 3 is Due Friday November 25th**
 - All DSCI 351/351M/451 Students:
 - * **Peer Grading of Report Out #2 is Due November 8th**
 - Exams
 - * Final: Monday December 19, 2022, 12:00PM - 3:00PM, Nord 356 or remote

10.2.3.2 Power Analysis, Sample Size, Effect Size

Power analysis allows you to

- determine the sample size required
 - to detect an effect of a given “effect size”
 - with a given degree of confidence.

Conversely, it allows you to determine

- the probability of detecting an effect
 - of a given effect size
- with a given level of confidence,
 - under sample size constraints.

If the probability is unacceptably low,

- you’d be wise to alter or abandon the experiment.

Lets see how to conduct power analyses

- for a variety of statistical tests,
 - including tests of proportions,
 - t-tests,
 - chi-square tests,
 - balanced one-way ANOVA,
 - tests of correlations,
 - and linear models.

Because power analysis applies to hypothesis testing situations,

- we’ll start with a review of null hypothesis significance testing (NHST).

Then we’ll review conducting power analyses within R,

- focusing primarily on the `pwr` package.

Finally, we’ll consider other approaches to power analysis available with R.

10.2.3.2.1 Review of NHST

In statistical hypothesis testing,

- you specify a hypothesis about a population parameter
 - (your null hypothesis, or H_0).

You then draw a sample from this population

- and calculate a statistic that's used
 - to make inferences about the population parameter.
- Assuming that the null hypothesis is true,
 - you calculate the probability of obtaining
 - the observed sample statistic or one more extreme.
- If the probability is sufficiently small,
 - you reject the null hypothesis in favor of its opposite
 - (referred to as the alternative or research hypothesis, H_1).

An example will clarify the process.

- Say you're interested in evaluating
 - the impact of cell phone use on driver reaction time.
- Your null hypothesis is $H_0: \mu_1 - \mu_2 = 0$,
 - where μ_1 is the mean response time for drivers using a cell phone
 - and μ_2 is the mean response time for drivers that are cell phone free
- (here, $\mu_1 - \mu_2$ is the population parameter of interest).

If you reject this null hypothesis,

- you're left with the alternate or research hypothesis,
 - namely $H_1: \mu_1 - \mu_2 \neq 0$.
- This is equivalent to $\mu_1 \neq \mu_2$,
 - that the mean reaction times for the two conditions are not equal.

A sample of individuals

- is selected
 - and randomly assigned to one of two conditions.
- In the first condition,
 - participants react to a series of driving challenges in a simulator
 - while talking on a cell phone.
- In the second condition, participants complete the same series of challenges
 - but without a cell phone.
- Overall reaction time is assessed for each individual.

Based on the sample data, you can calculate the statistic

- $(\bar{X}_1 - \bar{X}_2)/(s/\sqrt{n})$
 - where \bar{X}_1 and \bar{X}_2
 - are the sample reaction time means in the two conditions,
- s is the pooled sample standard deviation,
- and n is the number of participants in each condition.

If the null hypothesis is true

- and you can assume that reaction times are normally distributed,
- this sample statistic will follow
 - a t distribution
 - with $2n-2$ degrees of freedom.

Using this fact, you can calculate

- the probability of obtaining a sample statistic this large or larger.

If the probability (p) is smaller than some predetermined cutoff

- (say $p < .05$),
- you reject the null hypothesis
 - in favor of the alternate hypothesis.

This predetermined cutoff (0.05)

- is called the significance level of the test.

Note that you use sample data

- to make an inference about the population it's drawn from.

Your null hypothesis is

- that the mean reaction time of all drivers talking on cell phones
 - isn't different from
- the mean reaction time of all drivers who aren't talking on cell phones,
 - not just those drivers in your sample.

The four possible outcomes from your decision are as follows:

- 1. If the null hypothesis is false
 - and the statistical test leads you to reject it,
 - you've made a correct decision.
 - You've correctly determined that reaction time
 - is affected by cell phone use.
- 2. If the null hypothesis is true and you don't reject it,
 - again you've made a correct decision.
 - Reaction time isn't affected by cell phone use.
- 3. If the null hypothesis is true but you reject it,
 - you've committed a Type I error.
 - You've concluded that cell phone use affects reaction time
 - when it doesn't.
- 4. If the null hypothesis is false and you fail to reject it,
 - you've committed a Type II error.
 - Cell phone use affects reaction time,
 - but you've failed to discern this.

10.2.3.2.2 Controversy surrounding null hypothesis significance testing

Null hypothesis significance testing

- isn't without controversy
 - and detractors have raised numerous concerns about the approach,
 - particularly as practiced in the field of psychology.
- They point to
 - a widespread misunderstanding of p values,
 - reliance on statistical significance over practical significance,
 - the fact that the null hypothesis is never exactly true
 - and will always be rejected for sufficient sample sizes,
 - and a number of logical inconsistencies in NHST practices.

An interesting discussion is in 3-readings/2-articles

- [1]Bruce Thompson, Statistical Significance and Effect Size Reporting: Portrait of a Possible Future, Research in the Schools. 5 (1998) 33–38.
- [2]A. Gelman, E. Loken, The statistical crisis in science: data-dependent analysis—a "garden of forking paths"—explains why many statistically significant comparisons don't hold up, American Scientist. 102 (2014) 460–466. <https://doi.org/10.1511/2014.111.460>.

- [3]A. Gelman, The Problems With P-Values are not Just With P-Values, The American Statistician. (2016) 2.
- [4]B.B. McShane, D. Gal, A. Gelman, C. Robert, J.L. Tackett, Abandon Statistical Significance, The American Statistician. 73 (2019) 235–245. <https://doi.org/10.1080/00031305.2018.1527253>.

Researchers typically pays special attention to four quantities

- **Sample size** refers to the number of observations
 - in each condition/group of the experimental design or study protocol
- The **significance level** (also referred to as alpha)
 - defined as the probability of making a Type I error.
 - The significance level can also be thought of as
 - the probability of finding an effect that is not there.
- **Statistical Power** defined as
 - one minus the probability of making a Type II error.
 - Power can be thought of as
 - the probability of finding an effect that is there.
- **Effect size** is the magnitude of the effect
 - under the alternate or research hypothesis.
 - The formula for effect size depends on
 - the statistical methodology employed in the hypothesis testing.




Figure 10.1. Four primary quantities considered in a study design power analysis. Given any three, you can calculate the fourth.

Figure 1: Fig. 10.1

Although the sample size and significance level

- are under the direct control of the researcher,
 - power and effect size are affected more indirectly.

- For example, as you relax the significance level
 - (in other words, make it easier to reject the null hypothesis),
 - power increases.
- Similarly, increasing the sample size increases power.

Your research goal is typically

- to maximize the power of your statistical tests
 - while maintaining an acceptable significance level
 - and employing as small a sample size as possible.
- That is, you want to maximize the chances of finding a real effect
 - and minimize the chances of finding an effect that isn't really there,
 - while keeping study costs within reason.

The four quantities

- **sample size, significance level, power, and effect size**
 - have an intimate relationship.
- Given any three,
 - you can determine the fourth.

You'll use this fact to carry out various power analyses

- as we move forward,
- we'll look at ways of implementing power analyses
 - using the R package `pwr`.
- And we'll briefly look at some highly specialized power functions
 - that are used in biology and genetics.

10.2.3.3 Implementing power analysis with the `pwr` package

The `pwr` package implements power analysis as outlined by Cohen (1988).

Some of the more important functions are listed in table 10.1.

For each function,

- the user can specify three of the four quantities
 - (sample size, significance level, power, effect size),
- and the fourth will be calculated.

Of the four quantities, effect size is often the most difficult to specify.

Calculating effect size typically requires

- some experience with the measures involved
 - and knowledge of past research.
- But what can you do if you have no clue
 - what effect size to expect in a given study?
- You'll look at this difficult question in the next section.

Here we'll apply `pwr` package functions

- to common statistical tests.

10.2.3.3.1 t-tests

When the statistical test to be used is a t-test,

- the `pwr.t.test()` function
 - provides a number of useful power analysis options.

Table 10.1 pwr package functions

Function	Power calculations for ...
<code>pwr.2p.test</code>	Two proportions (equal n)
<code>pwr.2p2n.test</code>	Two proportions (unequal n)
<code>pwr.anova.test</code>	Balanced one way ANOVA
<code>pwr.chisq.test</code>	Chi-square test
<code>pwr.f2.test</code>	General linear model
<code>pwr.p.test</code>	Proportion (one sample)
<code>pwr.r.test</code>	Correlation
<code>pwr.t.test</code>	t-tests (one sample, two samples, paired)
<code>pwr.t2n.test</code>	t-test (two samples with unequal n)

Figure 2: table 10.1

```
pwr.t.test(n=, d=, sig.level=, power=, alternative=)
```

where

- `n` is the sample size.
- `d` is the effect size defined as the standardized mean difference.

$$d = \frac{|\mu_1 - \mu_2|}{\sigma}$$

μ_1 = mean of group 1
 μ_2 = mean of group 2
 σ^2 = common error variance

- `sig.level` is the significance level (0.05 is the default).
- `power` is the power level.
- `type` is two-sample t-test ("`two.sample`"), a one-sample t-test ("`one.sample`"), or a dependent sample t-test ("`paired`"). A two-sample test is the default.
- `alternative` indicates whether the statistical test is two-sided ("`two.sided`") or one-sided ("`less`" or "`greater`"). A two-sided test is the default.

Figure 3: t-test

The format is

Let's work through an example.

- Continuing the experiment from section 10.1
 - involving cell phone use and driving reaction time,
- assume that you'll be using a two-tailed independent sample t-test
 - to compare the mean reaction time
 - for participants in the cell phone condition
- with the mean reaction time for participants driving unencumbered.

Let's assume that you know from past experience

- that reaction time has a standard deviation of 1.25 seconds.
- Also suppose that a 1-second difference in reaction time
 - is considered an important difference.
- You'd therefore like to conduct a study
 - in which you're able to detect
 - an effect size of $d = 1/1.25 = 0.8$ or larger.
- Additionally, you want to be 90% sure
 - to detect such a difference if it exists,
- and 95% sure that you won't declare a difference
 - to be significant when it's actually due to random variability.
- How many participants will you need in your study?

Entering this information in the `pwr.t.test()` function, you have the following:

```
# t-tests
library(pwr)
pwr.t.test(
  d = .8,
  sig.level = .05,
  power = .9,
  type = "two.sample",
  alternative = "two.sided"
)

##
##      Two-sample t test power calculation
##
##              n = 33.82555
##              d = 0.8
##      sig.level = 0.05
##      power = 0.9
##      alternative = two.sided
##
## NOTE: n is number in *each* group
```

The results suggest that

- you need 34 participants in each group
 - (for a total of 68 participants)
- in order to detect an effect size of 0.8
 - with 90% certainty
 - and no more than a 5% chance of erroneously concluding
- that a difference exists
 - when, in fact, it doesn't.

Let's alter the question. Assume that in comparing the two conditions

- you want to be able to detect
 - a 0.5 standard deviation difference in population means.
- You want to limit the chances of falsely declaring
 - the population means to be different to 1 out of 100.
- Additionally, you can only afford
 - to include 40 participants in the study.
- What's the probability that you'll be able to detect
 - a difference between the population means that's this large,
 - given the constraints outlined?

Assuming that an equal number of participants

- will be placed in each condition, you have

```
pwr.t.test(
  n = 20,
  d = .5,
  sig.level = .01,
  type = "two.sample",
  alternative = "two.sided"
)
```

```
##
##      Two-sample t test power calculation
##
##              n = 20
##              d = 0.5
##      sig.level = 0.01
##      power = 0.1439551
##      alternative = two.sided
##
## NOTE: n is number in *each* group
```

10.2.3.3.2 ANOVA

You can also use the `pwr` package for ANOVA problems

```
pwr.anova.test(
  k = 5,
  f = .25,
  sig.level = .05,
  power = .8
)
```

```
##
##      Balanced one-way analysis of variance power calculation
##
##              k = 5
##              n = 39.1534
##              f = 0.25
##      sig.level = 0.05
##      power = 0.8
##
## NOTE: n is number in each group
```

10.2.3.3.3 Correlations

The `pwr.r.test()` function

- provides a power analysis for tests of correlation coefficients.

The format is as follows

- `pwr.r.test(n=, r=, sig.level=, power=, alternative=)`
- where
 - `n` is the number of observations,
 - `r` is the effect size (as measured by a linear correlation coefficient),
 - `sig.level` is the significance level,
 - `power` is the power level,
- and `alternative`
 - specifies a two-sided (`two.sided`)
 - or a one-sided (`less` or `greater`)
 - significance test.

For example, let's assume that you're studying

- the relationship between depression and loneliness.

Your null and research hypotheses are

- $H_0: \rho = 0.25$
- versus $H_1: \rho > 0.25$
 - where ρ is the population correlation
 - between these two psychological variables.

You've set your significance level to 0.05,

- and you want to be 90% confident
 - that you'll reject H_0 if it's false.
- How many observations will you need?

This code provides the answer:

```
# correlations
pwr.r.test(
  r = .25,
  sig.level = .05,
  power = .90,
  alternative = "greater"
)

##
##      approximate correlation power calculation (arctangh transformation)
##
##              n = 133.2803
##              r = 0.25
##      sig.level = 0.05
##      power = 0.9
##      alternative = greater
```

Thus, you need to assess depression and loneliness

- in 134 participants
- in order to be 90% confident
 - that you'll reject the null hypothesis
 - if it's false.

10.2.3.3.4 Linear Models

For linear models (such as multiple regression),

- the `pwr.f2.test()` function can be used
 - to carry out a power analysis.
- The format is
 - `pwr.f2.test(u=, v=, f2=, sig.level=, power=)`
- where `u` and `v` are
 - the numerator and denominator degrees of freedom
- and `f2` is the effect size.

$$f^2 = \frac{R^2}{1 - R^2} \quad \text{where} \quad R^2 = \text{population multiple correlation}$$

$$f^2 = \frac{R_{AB}^2 - R_A^2}{1 - R_{AB}^2} \quad \text{where} \quad R_A^2 = \text{variance accounted for in the population by variable set A}$$

$$R_{AB}^2 = \text{variance accounted for in the population by variable set A and B together}$$

The

first formula for `f2`

- is appropriate when you're evaluating
 - the impact of a set of predictors on an outcome.
- The second formula is appropriate
 - when you're evaluating the impact of one set of predictors
 - above and beyond a second set of predictors (or covariates).

Let's say you're interested in whether a boss's leadership style

- impacts workers' satisfaction
 - above and beyond the salary and perks associated with the job.
- Leadership style is assessed by four variables,
 - and salary and perks are associated with three variables.
- Past experience suggests that salary and perks
 - account for roughly 30% of the variance in worker satisfaction.

From a practical standpoint,

- it would be interesting if leadership style
 - accounted for at least 5% above this figure.
- Assuming a significance level of 0.05,
 - how many subjects would be needed
 - to identify such a contribution with 90% confidence?

In this example

- `sig.level=0.05`
- `power=0.90`
- `u=3`
 - (total number of predictors minus the number of predictors in set B),
- and the effect size is
 - `f2 = (.35 - .30)/(1 - .35) = 0.0769`.

Entering this into the function yields the following:

```
# linear models
pwr.f2.test(
  u = 3,
  f2 = 0.0769,
  sig.level = 0.05,
  power = 0.90
)

##
##      Multiple regression power calculation
##
##              u = 3
##              v = 184.2426
##              f2 = 0.0769
##      sig.level = 0.05
##              power = 0.9
```

10.2.3.3.5 Tests of Propotions

The `pwr.2p.test()` function

- can be used to perform a power analysis
 - when comparing two proportions.

The format is `pwr.2p.test(h=, n=, sig.level=, power=)`

- where `h` is the effect size
- and `n` is the common sample size in each group.

The effect size `h` is defined as

$$h = 2 \arcsin\left(\sqrt{p_1}\right) - 2 \arcsin\left(\sqrt{p_2}\right)$$

Figure 4: `h`

and can be calculated with the function `ES.h(p1, p2)`.

For unequal `ns`,

- the desired function is
- `pwr.2p2n.test(h =, n1 =, n2 =, sig.level=, power=)`

The `alternative=` option

- can be used to specify
 - a two-tailed (`two.sided`)
 - or one-tailed (`less` or `greater`) test.
- A two-tailed test is the default.

Let's say that you suspect

- that a popular medication relieves symptoms in 60% of users.
- A new (and more expensive) medication will be marketed
 - if it improves symptoms in 65% of users.
- How many participants will you need to include in a study
 - comparing these two medications

- if you want to detect a difference this large?

Assume that you want to be 90% confident in a conclusion

- that the new drug is better
 - and 95% confident that you won't reach this conclusion erroneously.
- You'll use a one-tailed test
 - because you're only interested in assessing
 - whether the new drug is better than the standard.

```
# tests of proportions
pwr.2p.test(
  h = ES.h(.65, .6),
  sig.level = .05,
  power = .9,
  alternative = "greater"
)
```

```
##
##      Difference of proportion power calculation for binomial distribution (arcsine transformation)
##
##              h = 0.1033347
##              n = 1604.007
##      sig.level = 0.05
##      power = 0.9
##      alternative = greater
##
## NOTE: same sample sizes
```

Based on these results,

- you'll need to conduct a study with 1,605 individuals
 - receiving the new drug
- and 1,605 receiving the existing drug
- in order to meet the criteria.

10.2.3.3.6 χ^2 tests

χ^2 tests are often used

- to assess the relationship between two categorical variables.

The null hypothesis is typically

- that the variables are independent
 - versus a research hypothesis that they aren't.
- The `pwr.chisq.test()` function can be used
 - to evaluate the power, effect size, or requisite sample size
 - when employing a χ^2 test.

The format is

- `pwr.chisq.test(w = , N = , df = , sig.level = , power =)`
- where `w` is the effect size,
- `N` is the total sample size,
- and `df` is the degrees of freedom.

Here, effect size `w` is defined as

The summation goes from 1 to `m`,

- where `m` is the number of cells in the contingency table.

$$w = \sqrt{\sum_{i=1}^m \frac{(p0_i - p1_i)^2}{p0_i}} \quad \text{where} \quad \begin{array}{l} p0_i = \text{cell probability in the } i\text{th cell under } H_0 \\ p1_i = \text{cell probability in the } i\text{th cell under } H_1 \end{array}$$

Figure 5: effect size for chi-squared test

The function `ES.w2(P)` can be used

- to calculate the effect size
 - corresponding to the alternative hypothesis
 - in a two-way contingency table.
- Here, `P` is a hypothesized two-way probability table.

As a simple example, let's assume that you're looking at

- the relationship between ethnicity and promotion.
- You anticipate that
 - 70% of your sample will be Caucasian,
 - 10% will be African American,
 - and 20% will be Hispanic.
- Further, you believe that
 - 60% of Caucasians tend to be promoted,
 - compared with 30% for African Americans
 - and 50% for Hispanics.

Your research hypothesis is

- that the probability of promotion
- follows the values in this table.

Table 10.2 Proportion of individuals expected to be promoted

Ethnicity	Promoted	Not Promoted
Caucasian	0.42	0.28
African American	0.03	0.07
Hispanic	0.10	0.10

Figure 6: promoted

For example, you expect that

- 42% of the population will be promoted Caucasians
 - (.42 = .70 × .60)
- and 7% of the population will be non-promoted African Americans
 - (.07 = .10 × .70).

Let's assume a significance level of 0.05

- and that the desired power level is 0.90.

The degrees of freedom in a two-way contingency table

- are $(r - 1) \times (c - 1)$,
 - where r is the number of rows
 - and c is the number of columns.

You can calculate the hypothesized effect size as follows

```
# chi-square tests
prob <- matrix(c(.42, .28, .03, .07, .10, .10),
               byrow = TRUE,
               nrow = 3)
ES.w2(prob)
```

```
## [1] 0.1853198
```

Using this information, you can calculate the necessary sample size by

```
pwr.chisq.test(
  w = .1853,
  df = 2,
  sig.level = .05,
  power = .9
)
```

```
##
##      Chi squared power calculation
##
##              w = 0.1853
##              N = 368.5317
##              df = 2
##      sig.level = 0.05
##              power = 0.9
##
## NOTE: N is the number of observations
```

The results suggest that

- a study with 369 participants
 - will be adequate to detect
 - a relationship between ethnicity and promotion
- given the
 - effect size,
 - power,
 - and significance level specified.

```
# Sample sizes for detecting significant effects in a one-way ANOVA
library(pwr)
es <-
  seq(.1, .5, .01) # Sets the range of correlations and power values
nes <- length(es)

samsize <-
  NULL
for (i in 1:nes) { # Obtains sample sizes

  result <- pwr.anova.test(
    k = 5,
```

```

    f = es[i],
    sig.level = .05,
    power = .9
  )
  samsize[i] <- ceiling(result$n)
}

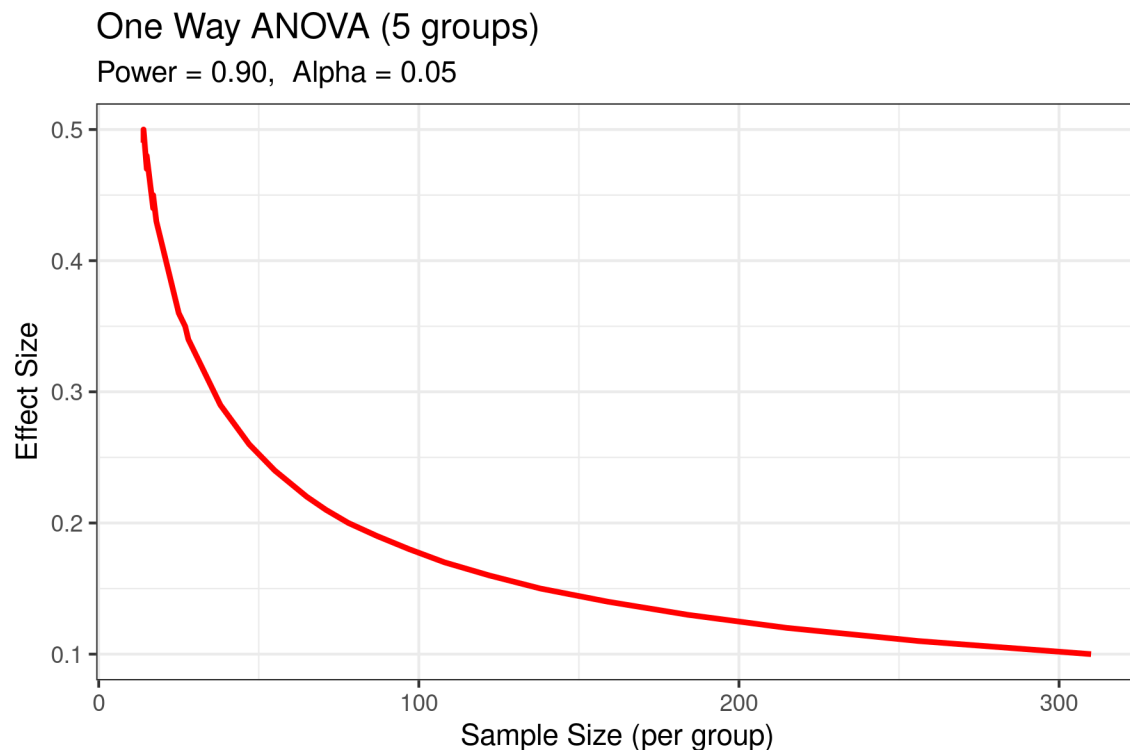
plotdata <- data.frame(es, samsize)

library(ggplot2) # plot power curves
ggplot(plotdata, aes(x = samsize, y = es)) +
  geom_line(color = "red", size = 1) +
  theme_bw() +
  labs(
    title = "One Way ANOVA (5 groups)",
    subtitle = "Power = 0.90, Alpha = 0.05",
    x = "Sample Size (per group)",
    y = "Effect Size"
  )

```

10.2.3.3.7 Samples Sizes in Anova Analysis

Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
 ## i Please use `linewidth` instead.



10.2.3.4 Creating Power Analysis Plots

Let's look at a more involved graphing example.

Suppose you'd like to see the sample size necessary

- to declare a correlation coefficient statistically significant
- for a range of
 - effect sizes
 - and power levels.

You can use the `pwr.r.test()` function

- and for loops to accomplish this task

Here we use the `seq()` function

- to generate a range of effect sizes `r`
 - (correlation coefficients under H_1)
- and power levels `p`.
- The `expand.grid()` function is used
 - to create a data frame
 - with every combination of these two variables.
- A for loop then cycles through rows of the data frame,
 - calculating the sample size (`n`)
 - for that row's correlation and power level,
 - and saving the result.
- The `ggplot2` package is then used
 - to plot a sample size vs. correlation curve for each power level.

```
# Sample size curves for detecting correlations of various sizes
library(pwr)
r <- seq(.1, .5, .01) # Sets the range of correlations and power values
p <- seq(.4, .9, .1)


df <- expand.grid(r, p)
colnames(df) <- c("r", "p")

for (i in 1:nrow(df)) { # obtain sample sizes
  result <- pwr.r.test(
    r = df$r[i],
    sig.level = .05,
    power = df$p[i],
    alternative = "two.sided"
  )
  df$n[i] <- ceiling(result$n)
}

library(ggplot2) # plot power curves
ggplot(data = df,
       aes(x = r, y = n, color = factor(p))) +
  geom_line(size = 1) +
  theme_bw() +
  labs(
    title = "Sample Size Estimation for Correlation Studies",
    subtitle = "Sig=0.05 (Two-tailed)",
    x = "Correlation Coefficient (r)",
    y = "Sample Size (n)",
    color = "Power"
  )
}
```

Sample Size Estimation for Correlation Studies

Sig=0.05 (Two-tailed)



As you can see from the graph,

- you'd need a sample size of approximately 75
 - to detect a correlation of 0.20
 - with 40% confidence.
- You'd need approximately 185 additional observations
 - (n = 260)
 - to detect the same correlation with 90% confidence.
- With simple modifications, the same approach can be used
 - to create sample size and power curve graphs
 - for a wide range of statistical tests.

10.2.3.5 Links

Robert I. Kabacoff, R in Action, 3rd Edition, Manning Publications