Data science with **R**

January 2017

Hadley Wickham
@hadleywickham
Chief Scientist, **RStudio**


My goal is to make
a **pit of success**
http://blog.codinghorror.com/falling-into-the-pit-of-success/

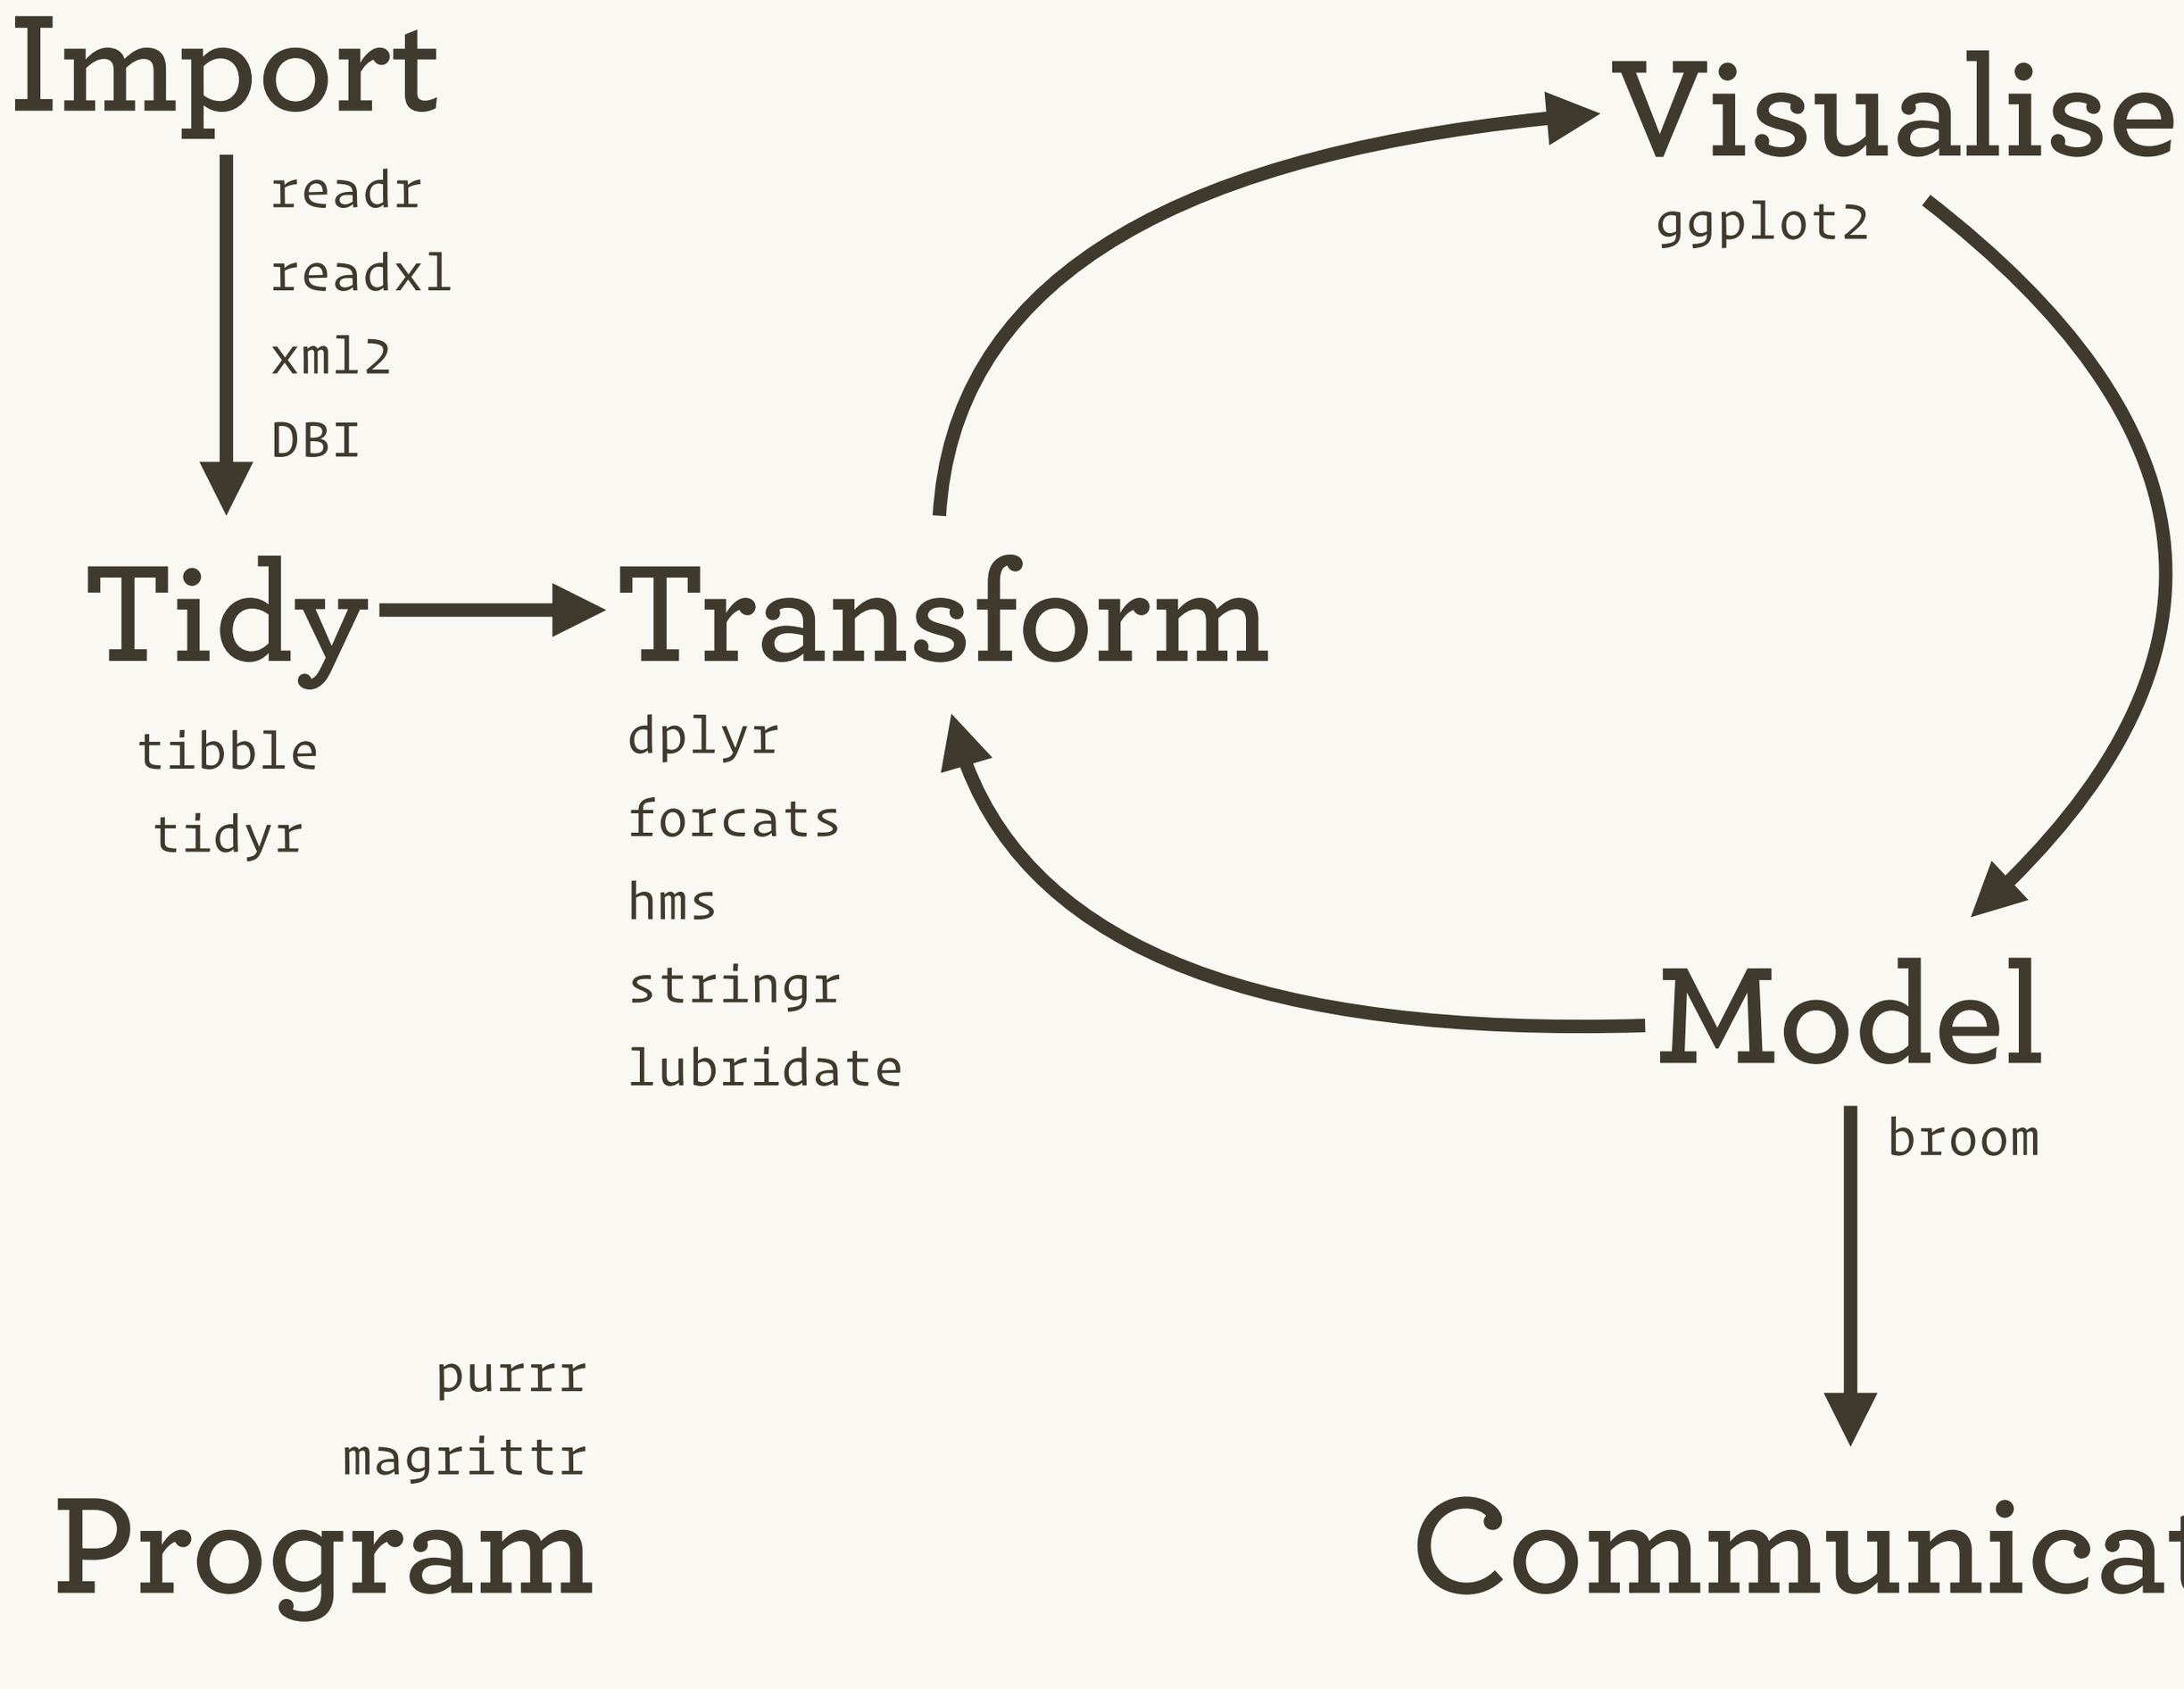
Import
readr
readxl
xml2
DBI
Visualise
ggplot2
Tidy
tibble
tidyr
Transform
dplyr
forcats
hms
stringr
lubridate
Model
broom
purrr
magrittr
Program
Communicate

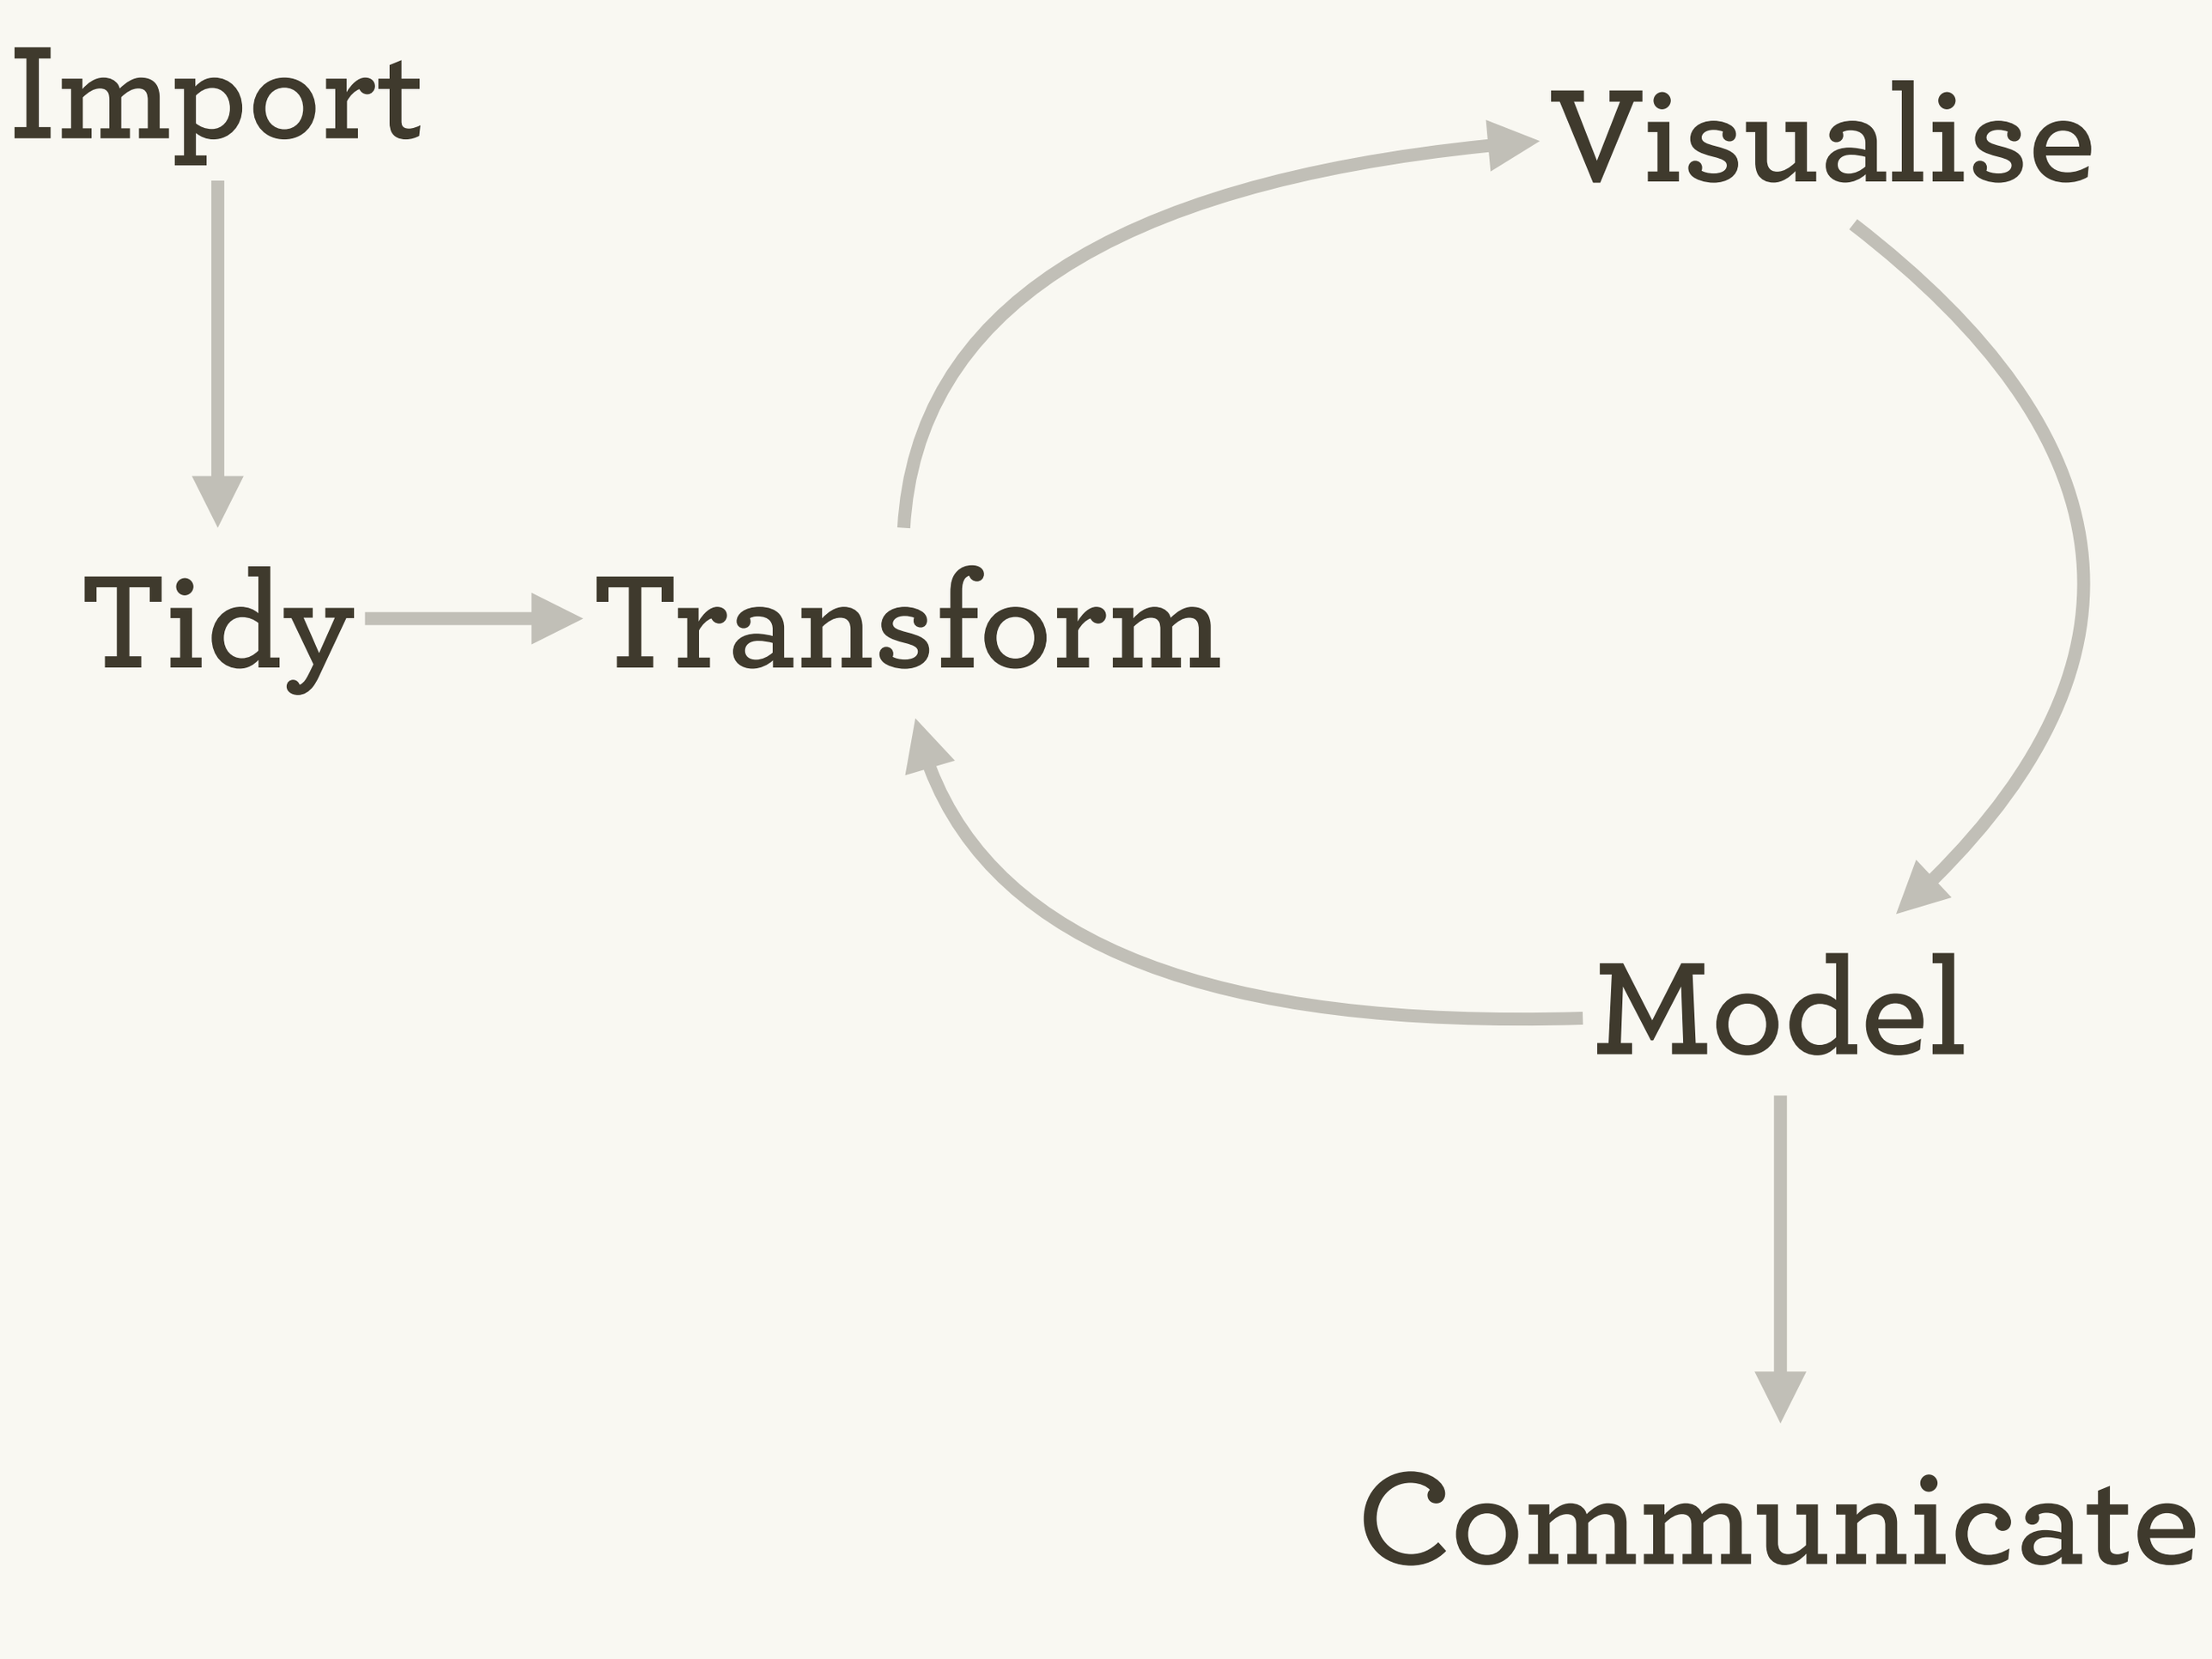
Import
Visualise
Tidy
Transform
Model
Communicate

No matter how complex and polished the individual operations are, it is often the quality of the glue that most directly determines the power of the system.

— *Hal Abelson*

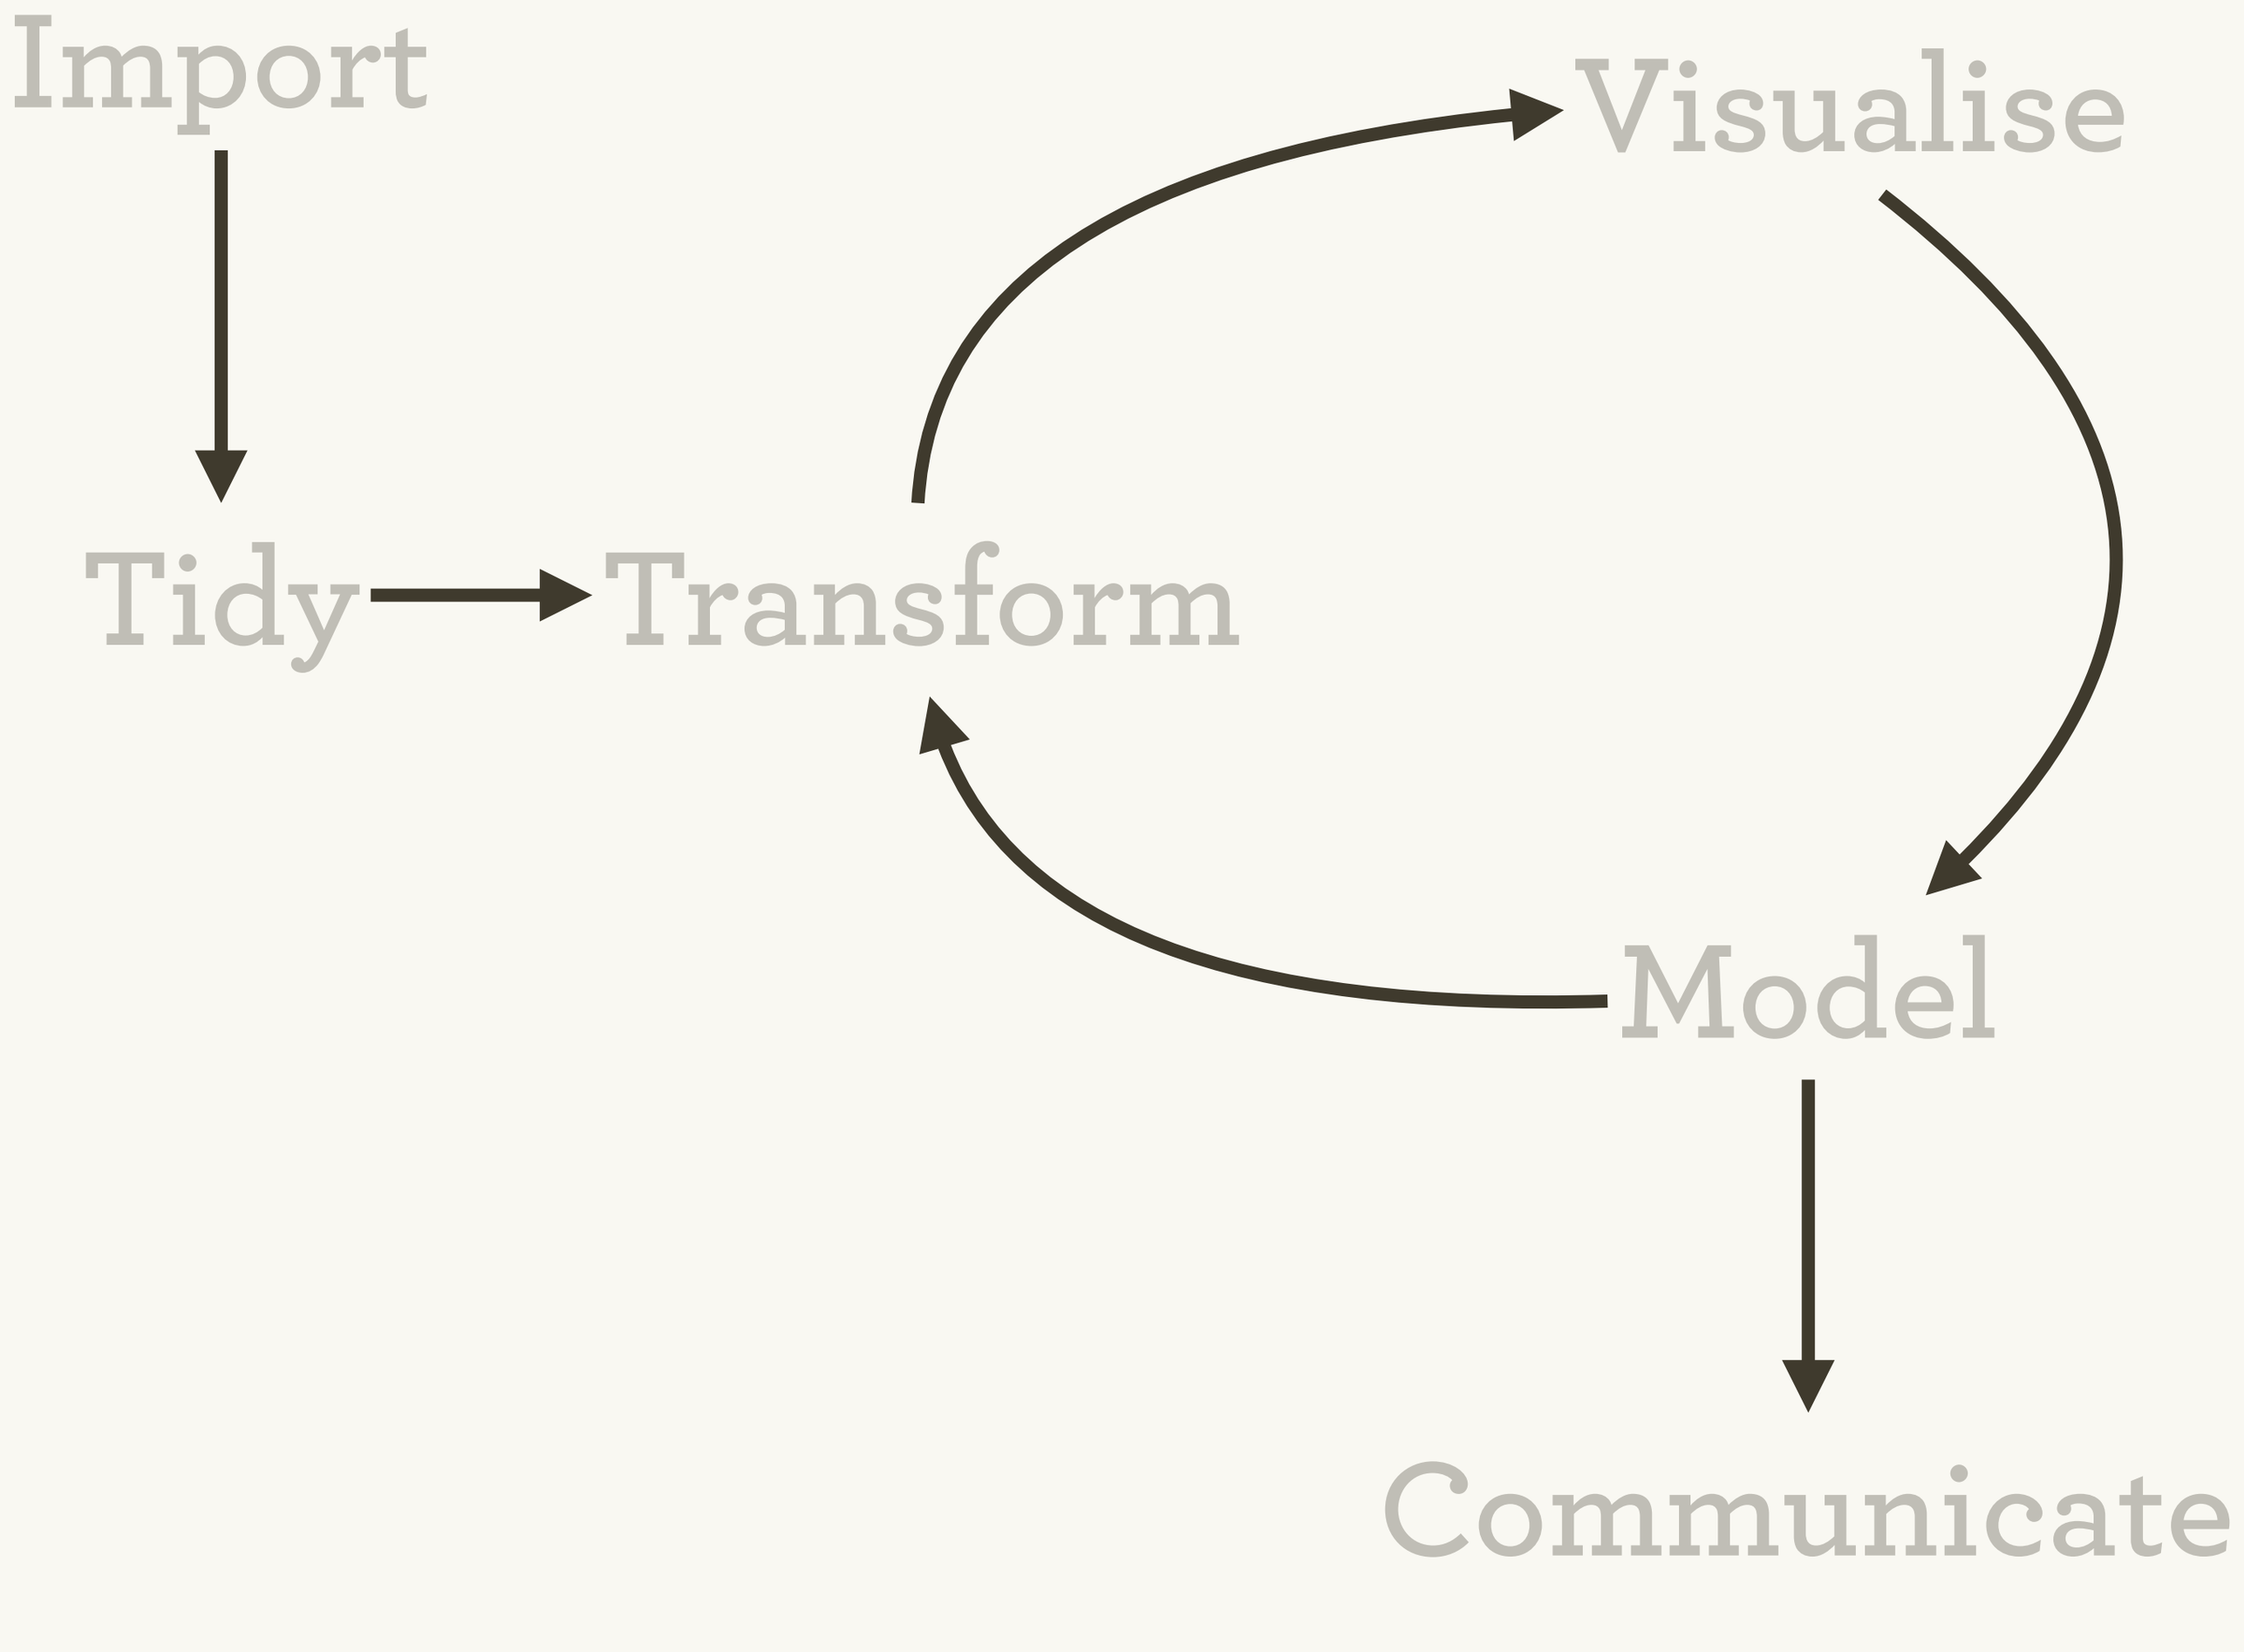
Import
Visualise
Tidy
Transform
Model
Communicate

The tidyverse

https://unsplash.com/photos/gkbo_YnBY1U

Four practical things you should know

1. It exists
2. It has a website
3. It has a package
4. It has a book

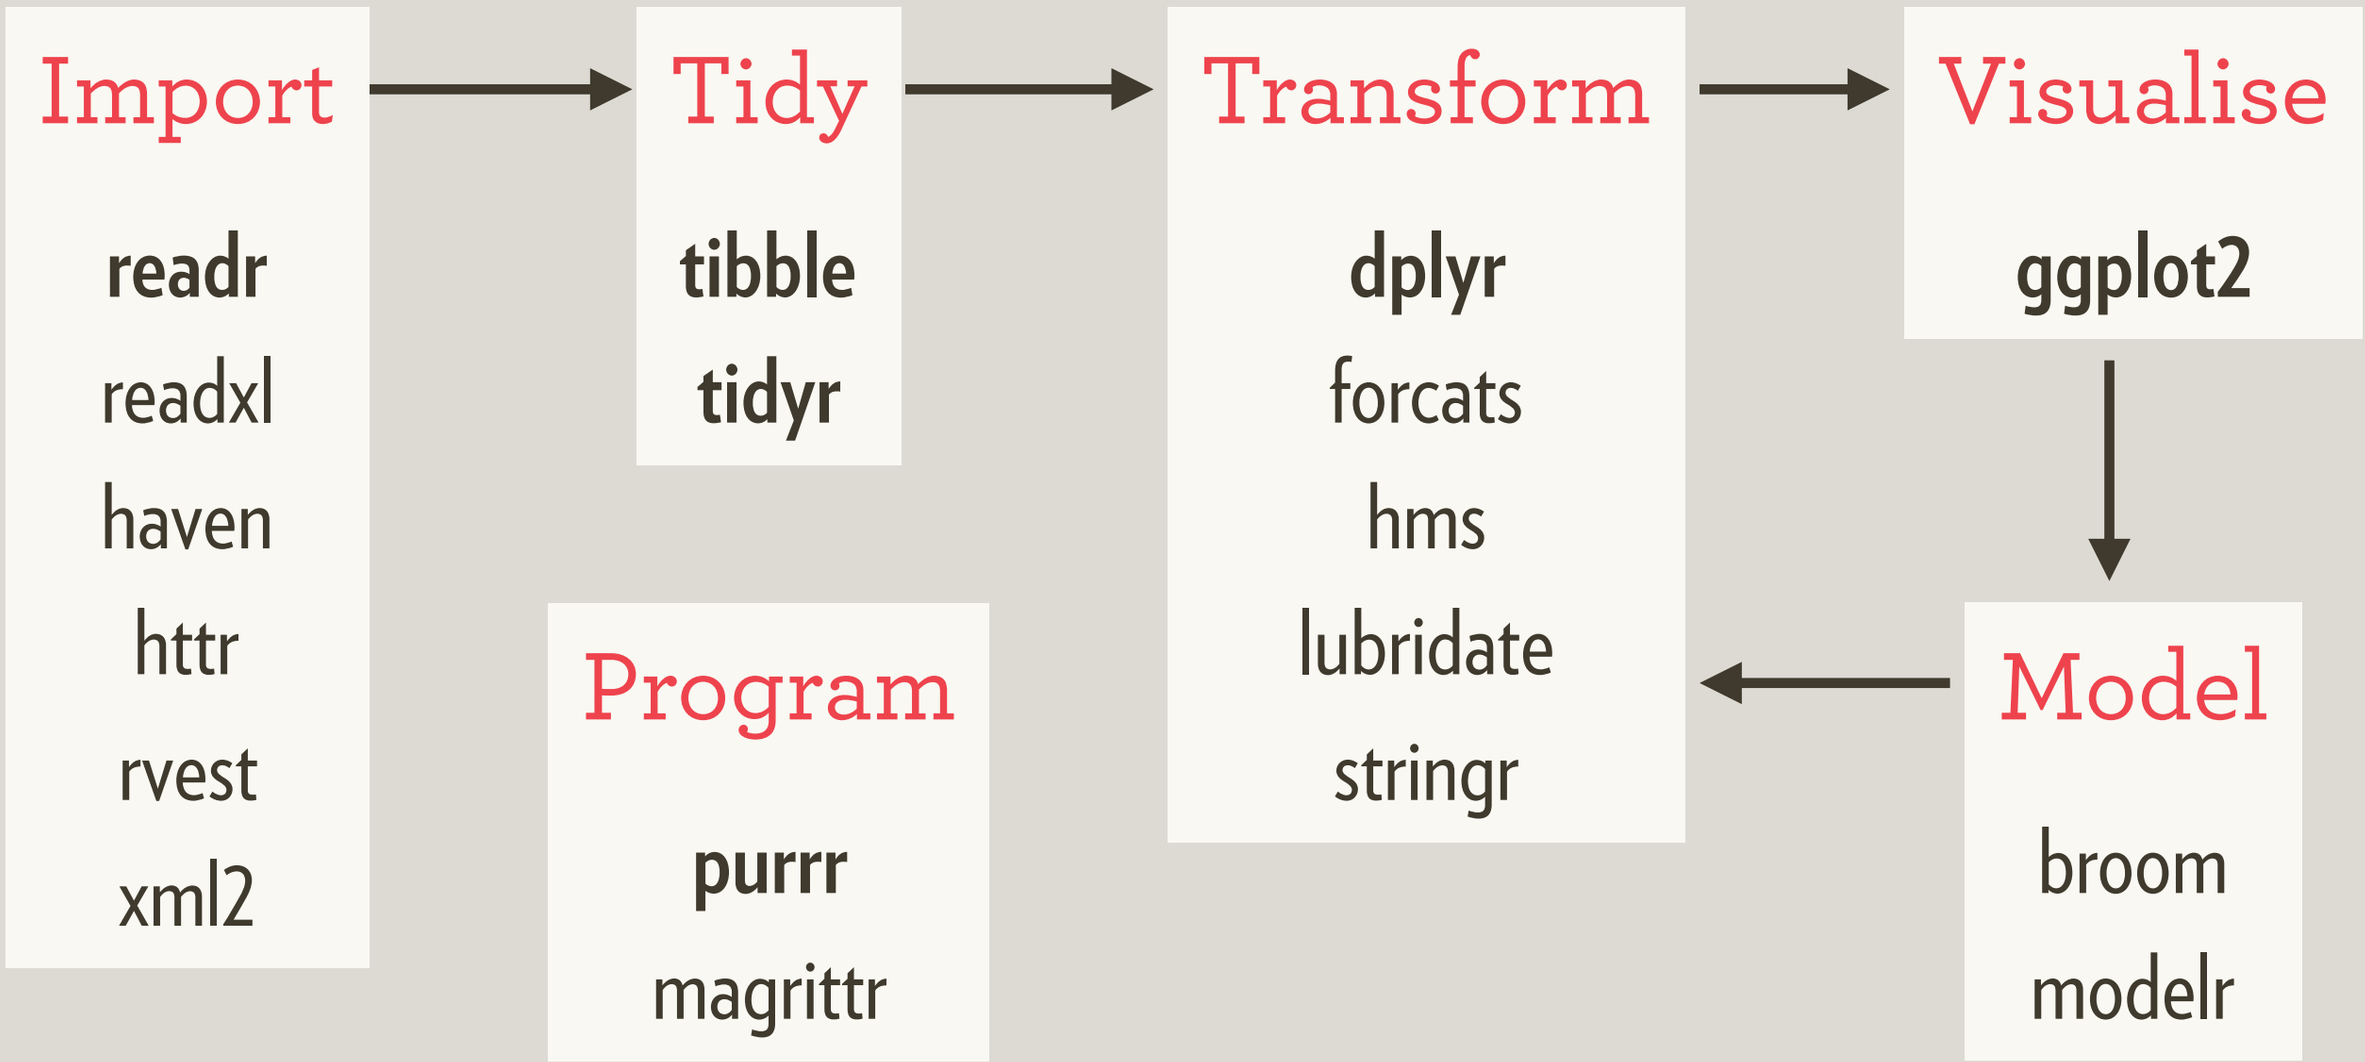
Import
readr
readxl
haven
httr
rvest
xml2
Tidy
tibble
tidyr
Transform
dplyr
forcats
hms
lubridate
stringr
Visualise
ggplot2
Model
broom
modelr
Program
purrr
magrittr

tidyverse.org

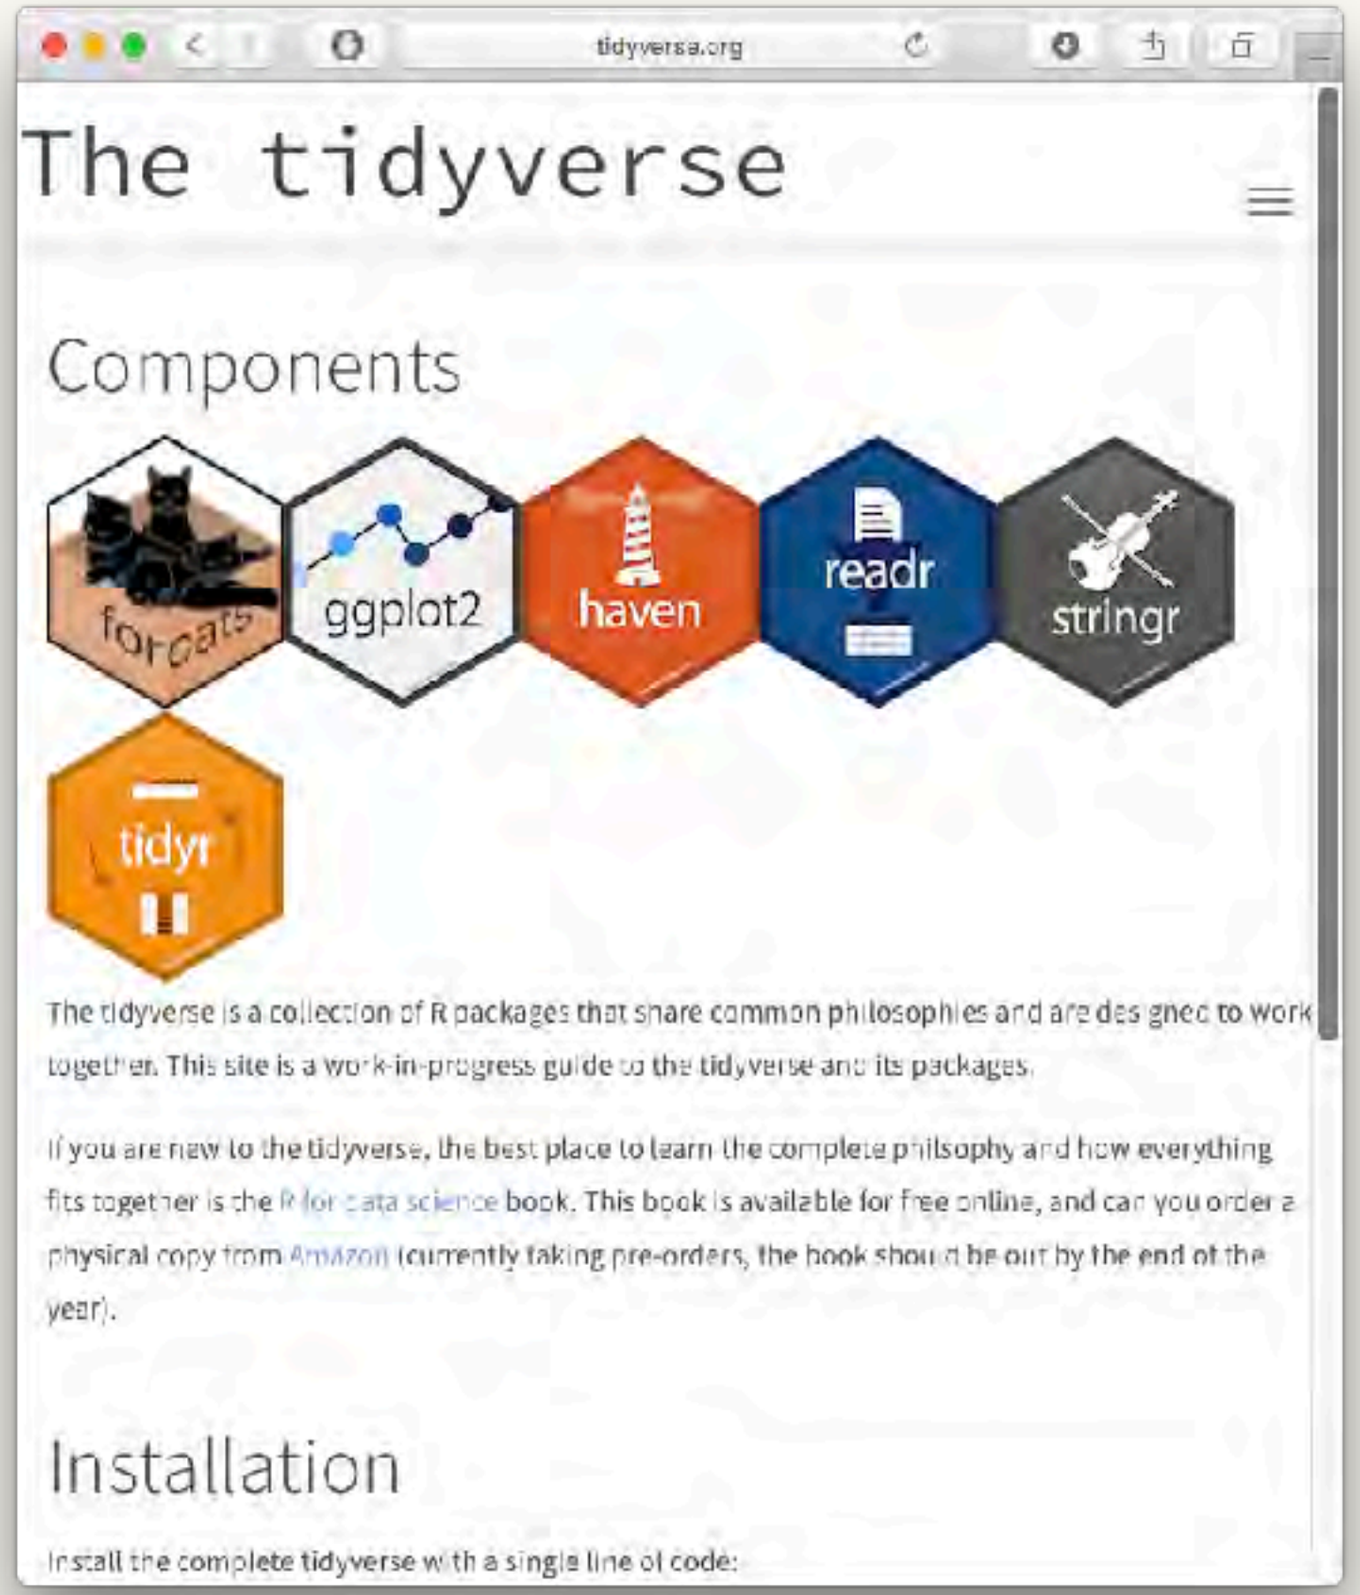

What am I missing out on?

How can I learn more?

How can I get help?

Installing tidyverse installs everything

```
install.packages("tidyverse")

# Instead of
install.packages(c(
  "broom", "dplyr", "feather",
  "forcats","ggplot2", "haven",
  "httr", "hms", "jsonlite",
  "lubridate", "magrittr",
  "modelr", "purrr", "readr",
  "readxl", "stringr", "tibble",
  "rvest", "tidyr", "xml2"
))
```

Loading it loads the **core** tidyverse

```
library(tidyverse)

# Instead of:
library(ggplot2)
library(tibble)
library(tidyr)
library(readr)
library(purrr)
library(dplyr)

# These are the packages you use in almost
# every analysis
```

Read online:

r4ds.had.co.nz

O'Reilly discount:

AUTHD

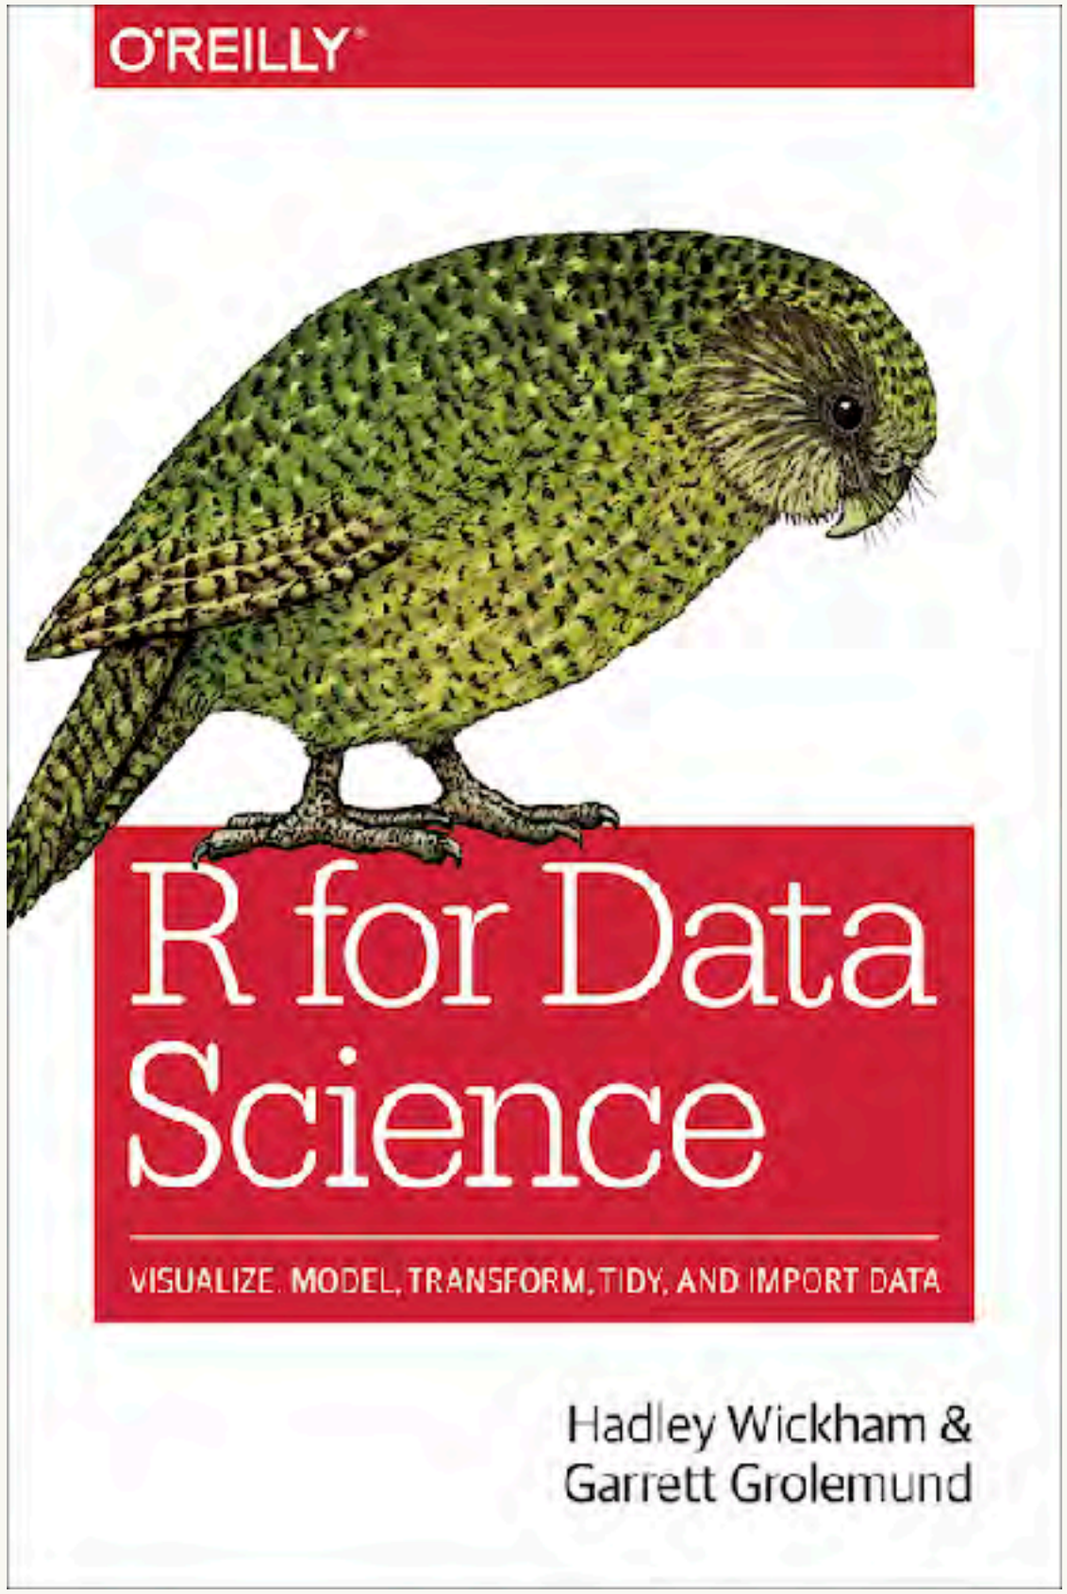

https://unsplash.com/photos/JxLf4KYVT-A

Goal: Solve complex problems by combining simple, uniform pieces.

Consistent functions

The tidyverse separates **commands** and **queries**

A **command** function performs an action

A **query** function computes a value

Closely ideas: *referential transparency* and *functional purity*

Which is which?

mutate()

write_csv()

print()

summarise()

+ geom_line()

<-

plot()

What's the difference?

```
# Command
print()
plot()
write_csv()
<-

# Query
summarise()
mutate()
+ geom_line()
```

Base R generally sticks to this principle

```
mod <- lm(mpg ~ wt, data = mtcars)
summary(mod)
#> Coefficients:
#>             Estimate Std. Error t value Pr(>|t|)
#> (Intercept)   37.285      1.878   19.86  < 2e-16 ***
#> wt            -5.344      0.559   -9.56  1.3e-10 ***
#> ---
#>
#> Residual standard error: 3.05 on 30 degrees of freedom
#> Multiple R-squared:  0.753, Adjusted R-squared:  0.745
#> F-statistic: 91.4 on 1 and 30 DF,  p-value: 1.29e-10
```

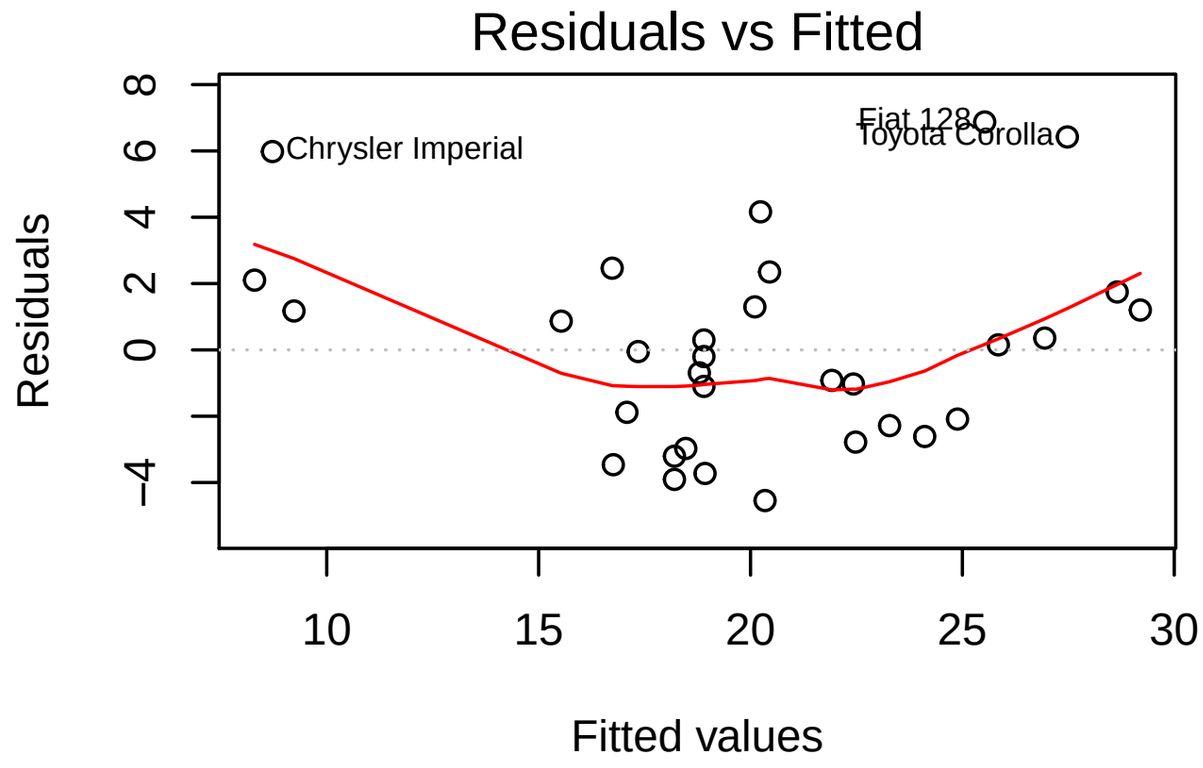
Residuals vs Fitted
Residuals
Fitted values
Chrysler Imperial
Fiat 128
Toyota Corolla
−4
0
2
4
6
8
10
15
20
25
30

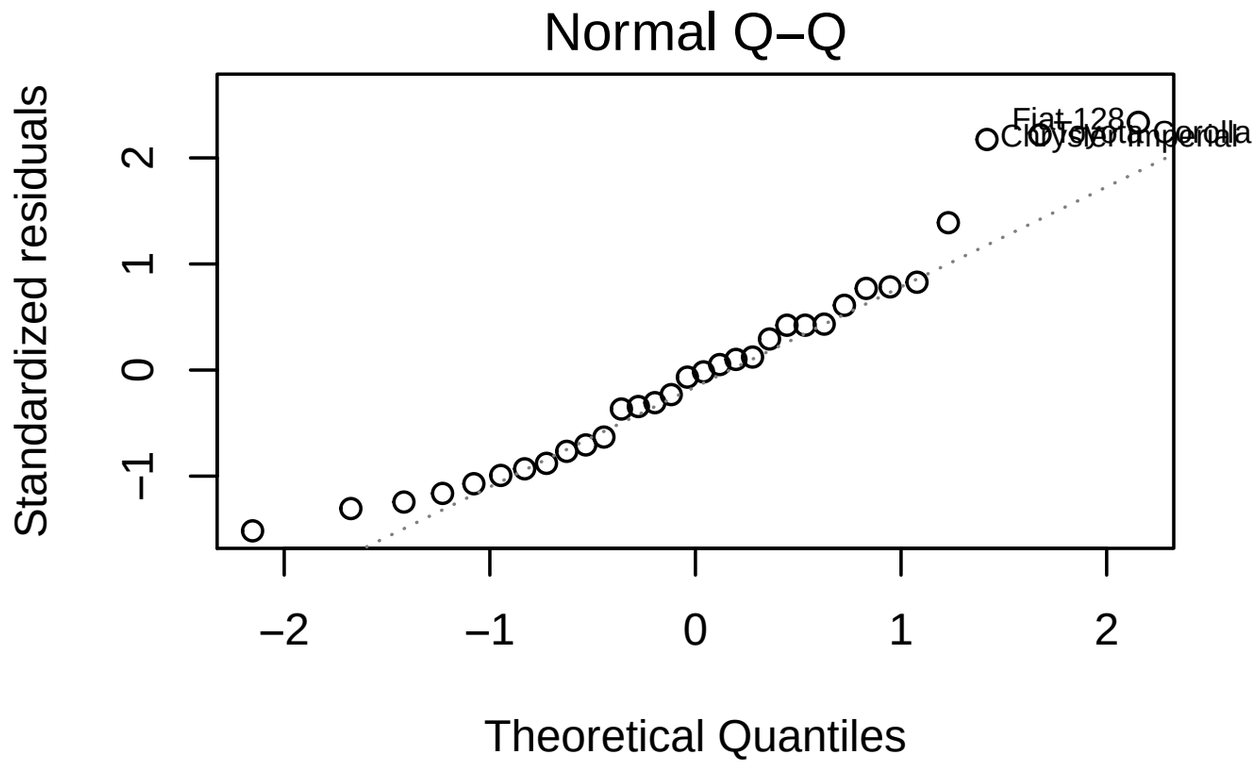
Normal Q–Q
Standardized residuals
Theoretical Quantiles
Fiat 128
Toyota Corolla
Chrysler Imperial
−1
0
1
2
−2
−1
0
1
2

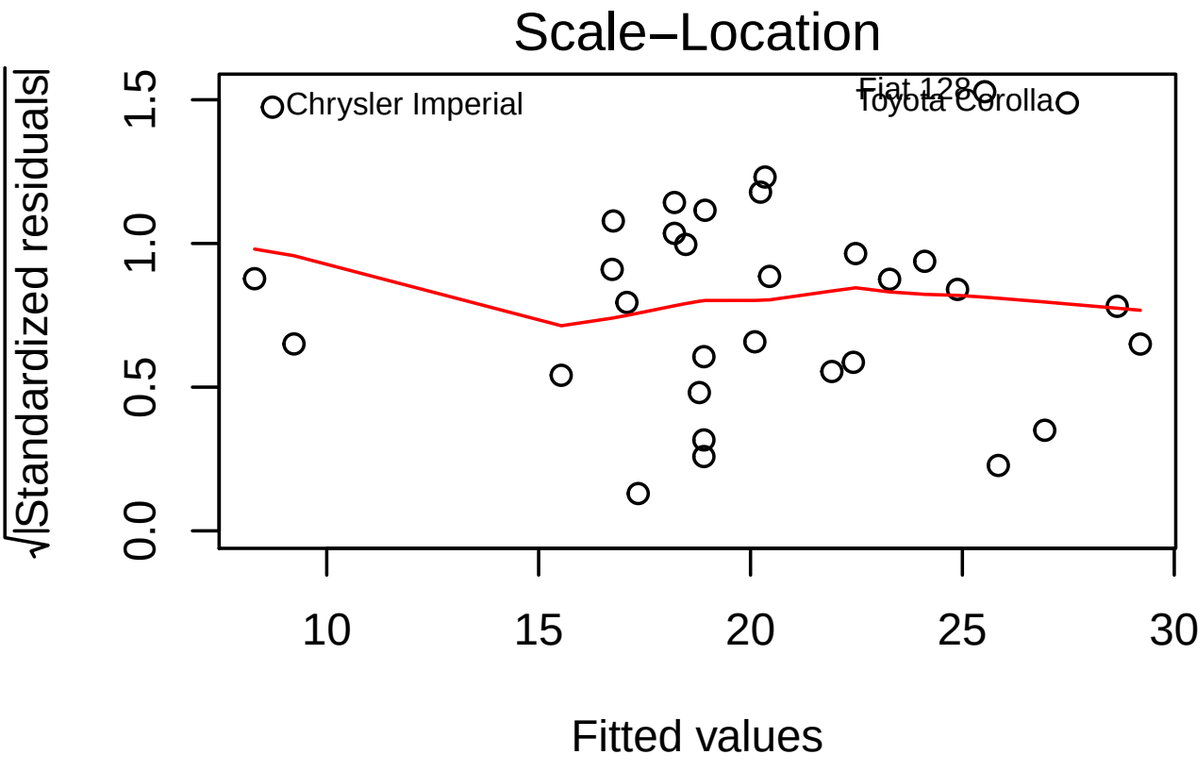
Scale–Location
√|Standardized residuals|
Fitted values
Chrysler Imperial
Fiat 128
Toyota Corolla
0.0
0.5
1.0
1.5
10
15
20
25
30

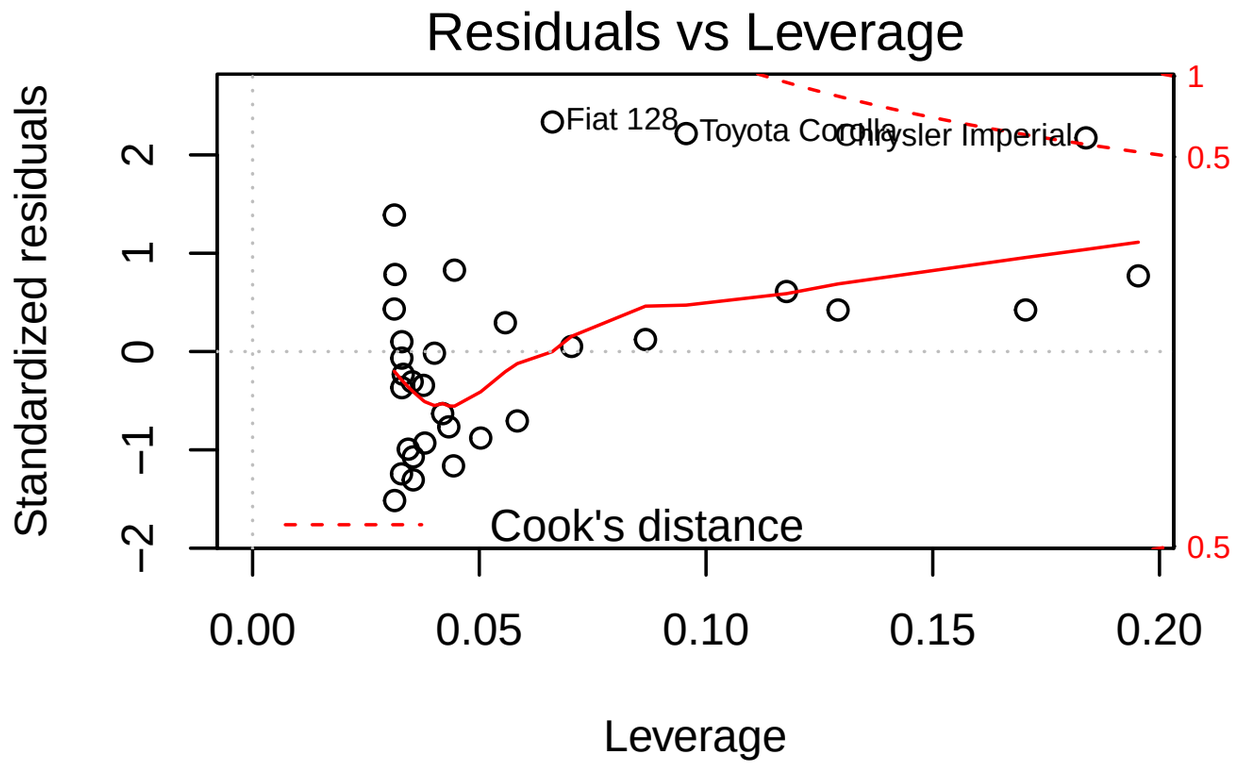
Residuals vs Leverage
Standardized residuals
Leverage
Fiat 128
Toyota Corolla
Chrysler Imperial
Cook's distance
1
0.5
0.5
−2
−1
0
1
2
0.00
0.05
0.10
0.15
0.20

Query functions are like legos

(As long as you pick a consistent data structure)

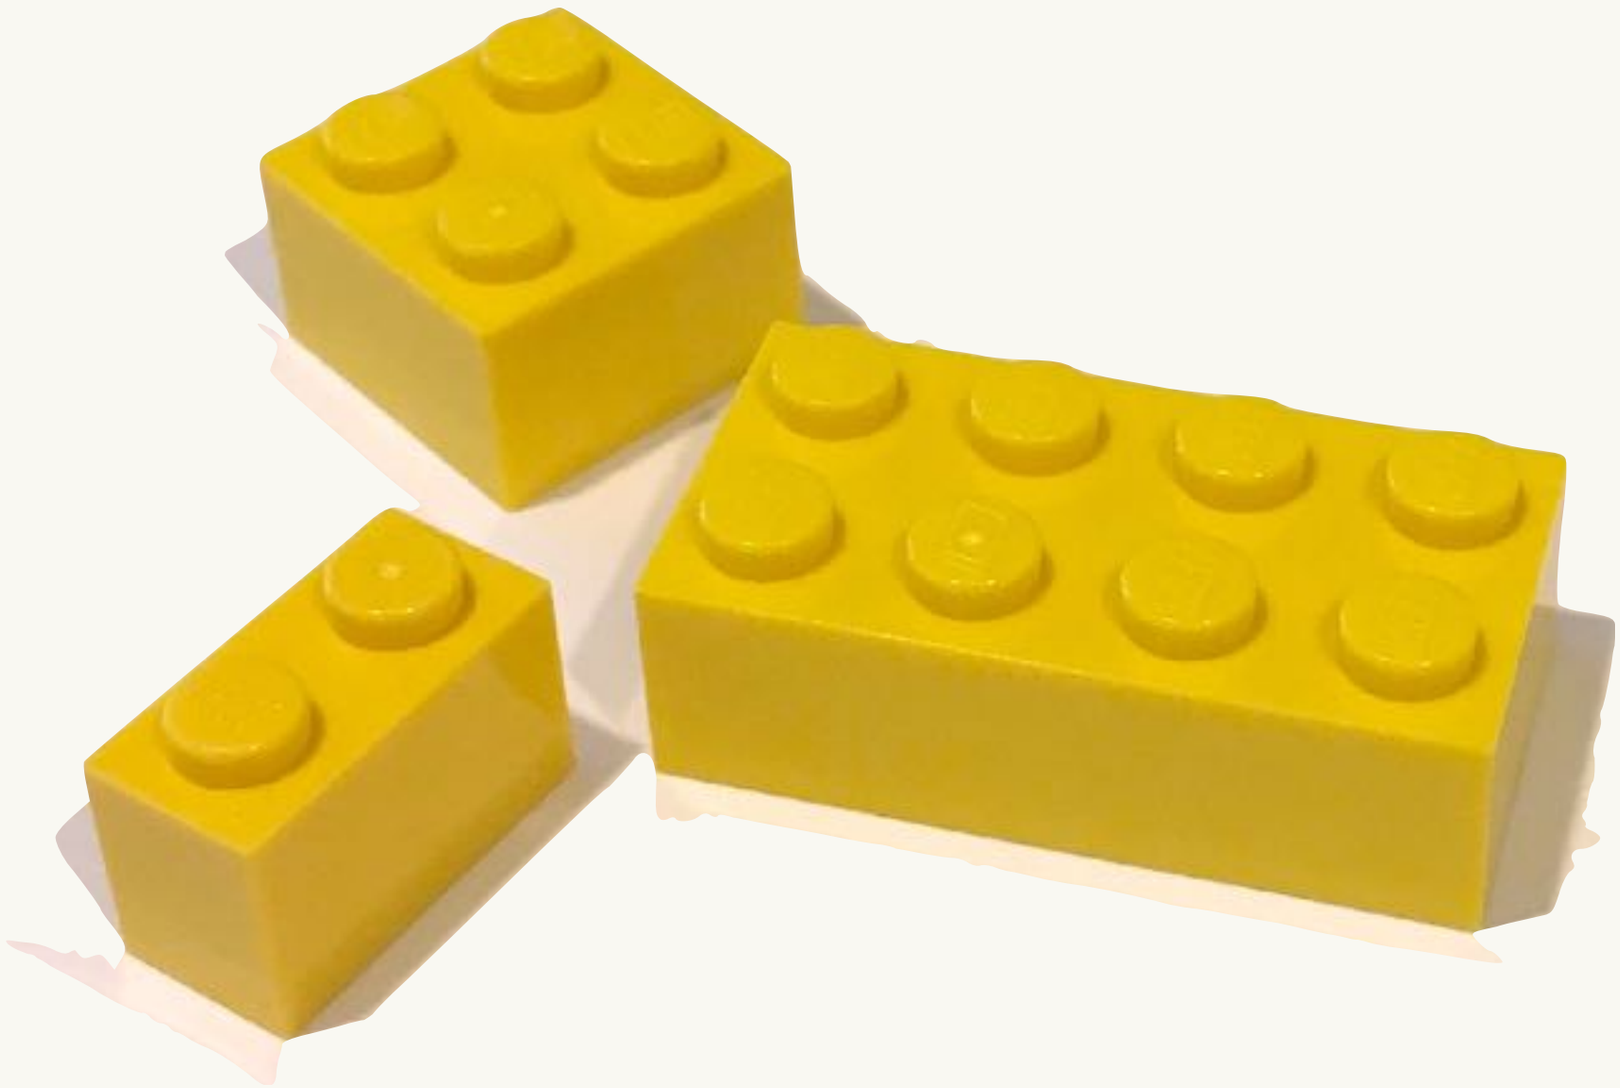

Command functions are like playmobil

https://unsplash.com/photos/0VNVxhEnkII

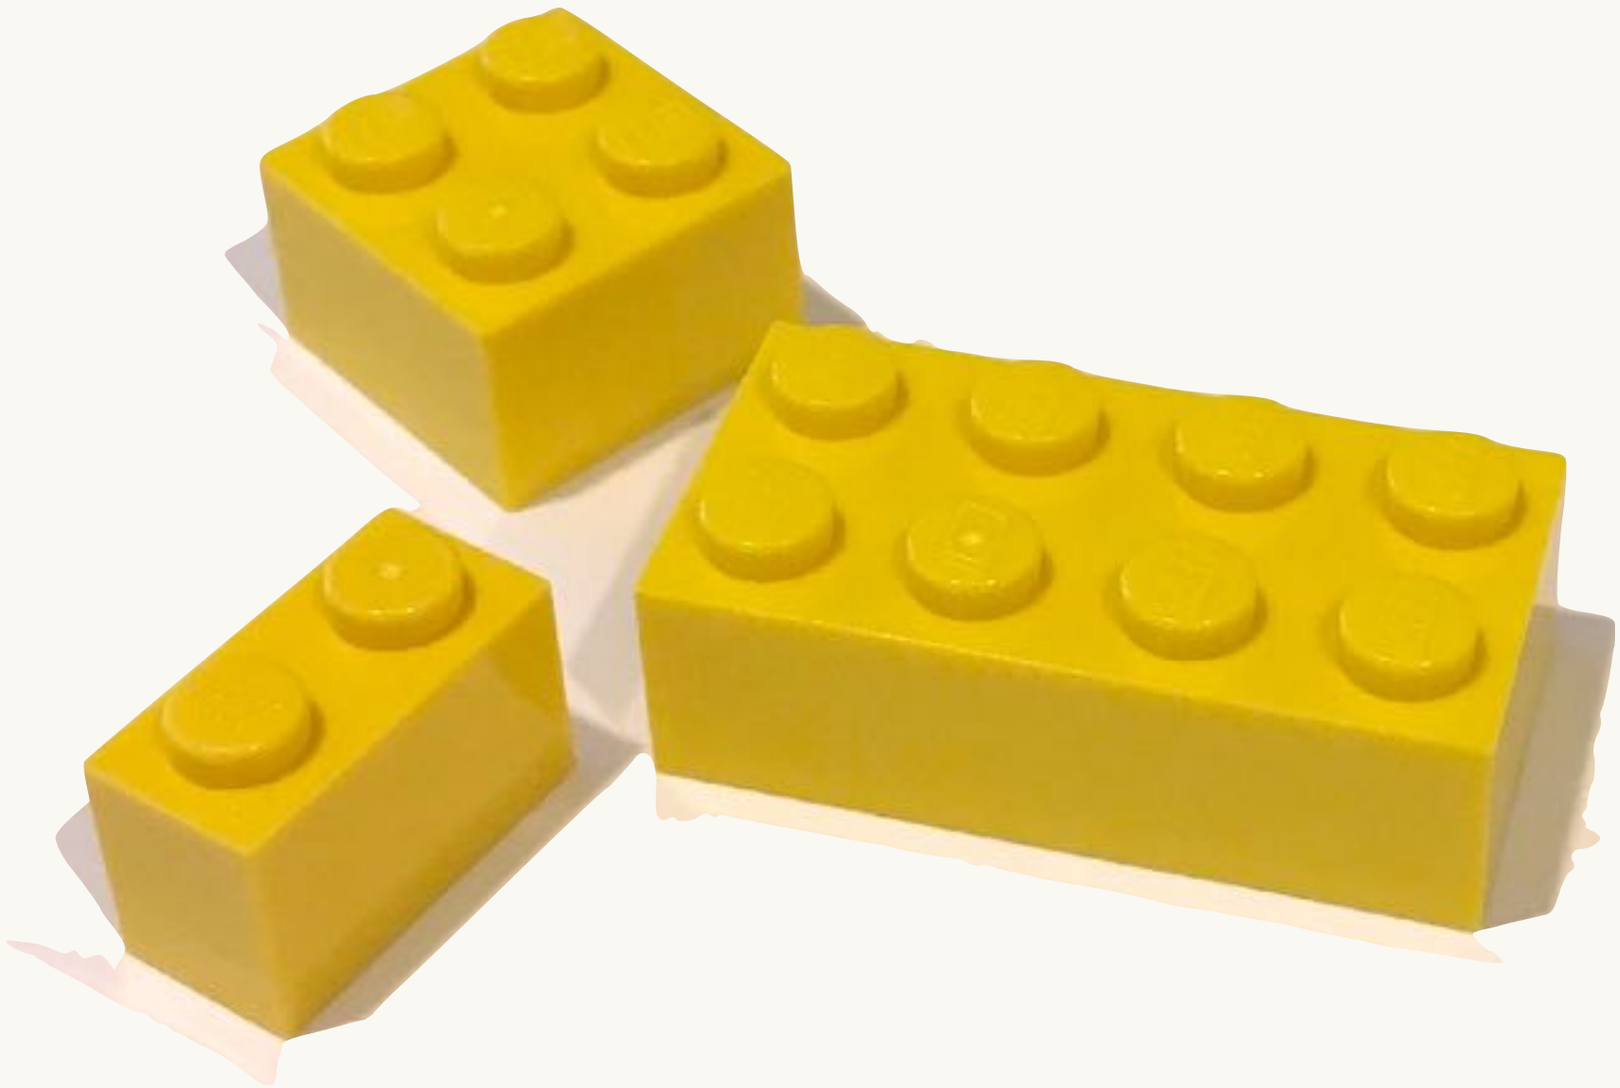

http://brickartist.com/gallery/pc-magazine-computer/. CC-BY-NC

The pipe

Goal: Solve complex problems by **combining** uniform pieces.

https://unsplash.com/photos/hjEesK4KSDs

We already have ways to combine functions

```
by_dest <- group_by(flights, dest)
dest_delay <- summarise(by_dest,
  delay = mean(dep_delay, na.rm = TRUE),
  n = n()
)
big_dest <- filter(dest_delay, n > 100)
arrange(big_dest, desc(delay))
```

But naming is hard work

```
foo <- group_by(flights, dest)
foo <- summarise(foo,
  delay = mean(dep_delay, na.rm = TRUE),
  n = n()
)
foo <- filter(foo, n > 100)
arrange(foo, desc(delay))
```

But naming is hard work

```
foo1 <- group_by(flights, dest)
foo2 <- summarise(foo1,
  delay = mean(dep_delay, na.rm = TRUE),
  n = n()
)
foo3 <- filter(foo2, n > 100)
arrange(foo3, desc(delay))
```

You *could* nest function calls

```
arrange(
  filter(
    summarise(
      group_by(flights, dest),
      delay = mean(dep_delay, na.rm = TRUE),
      n = n()
    ),
    n > 100
  ),
  desc(delay)
)
```

magrittr::

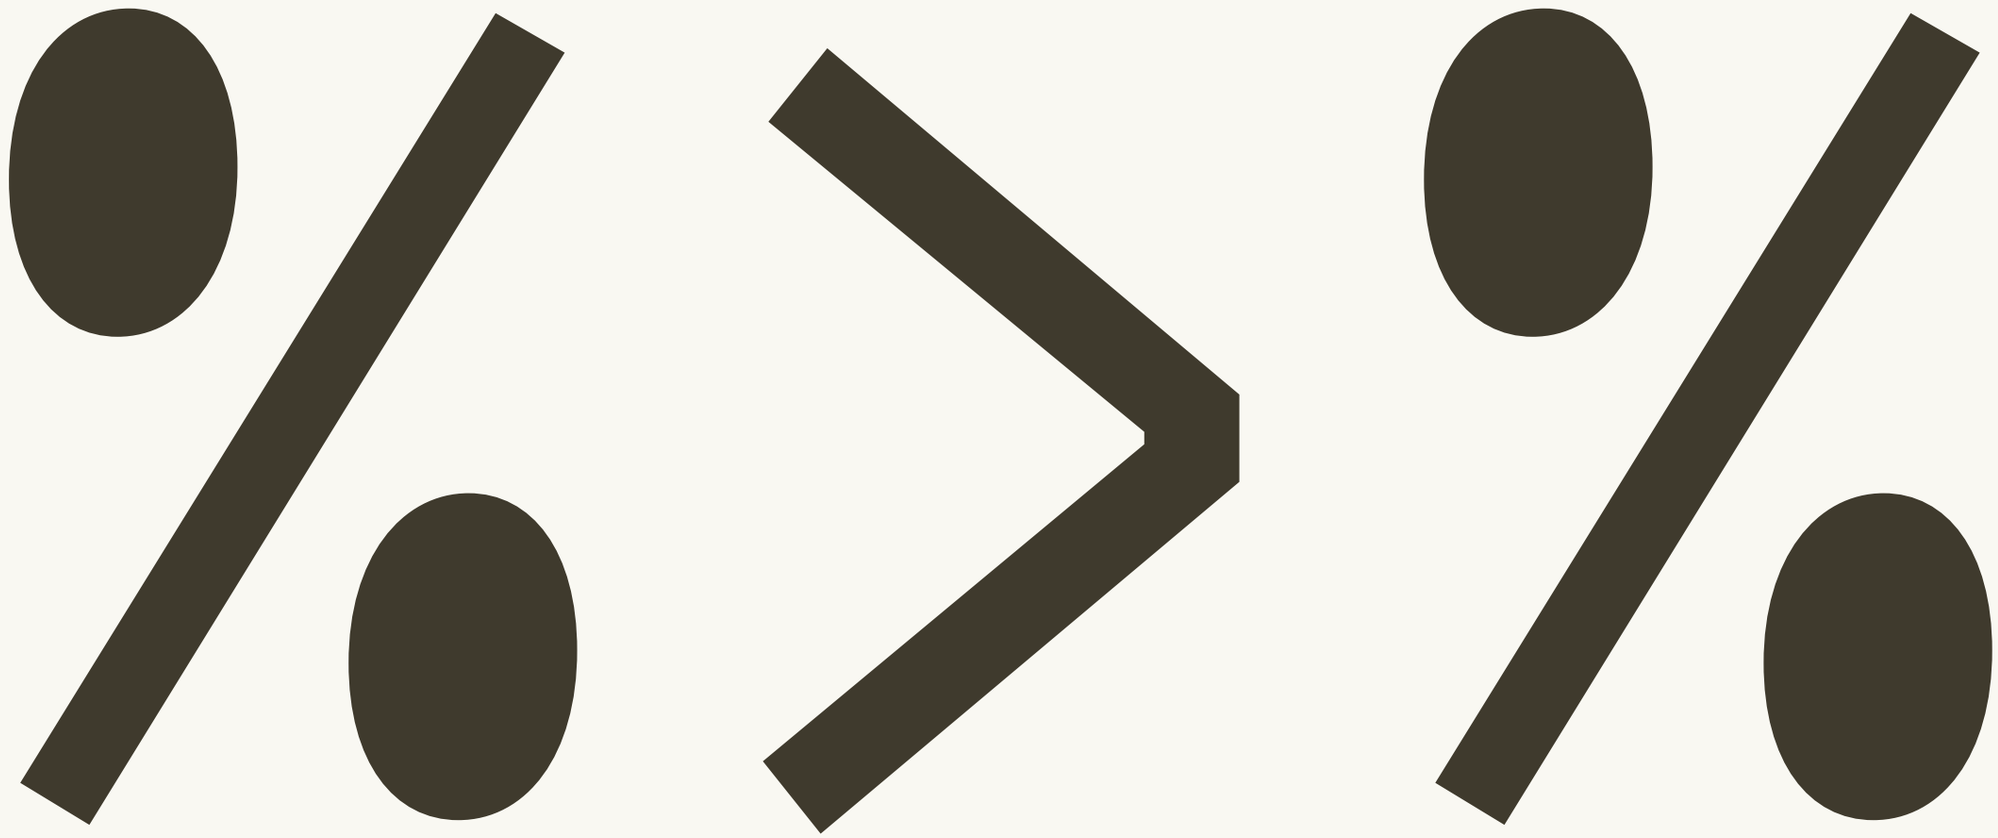

This is easy to read & doesn't require naming

```
flights %>%
  group_by(dest) %>%
  summarise(
    delay = mean(dep_delay, na.rm = TRUE),
    n = n()
) %>%
filter(n > 100) %>%
arrange(desc(delay))
```

Writing readable code with pipes

Bob Rudis

Friday • 3:10pm

ggplot2 was written before the pipe

```
flights %>%
  group_by(date) %>%
  summarise(n = n()) %>%
  ggplot(aes(date, n)) +
  geom_line()
```

And is inconsistent

```
ggsave(
  flights %>%
    group_by(date) %>%
    summarise(n = n()) %>%
    ggplot(aes(date, n)) +
    geom_line(),
  "my-plot.pdf"
)
```

The command-query distinction is useful for pipes

The body is made up of **queries**

Every pipe is ended by a **command**

Where is the command function?

```
flights %>%
  group_by(dest) %>%
  summarise(
    delay = mean(dep_delay, na.rm = TRUE),
    n = n()
) %>%
filter(n > 100) %>%
arrange(desc(delay))
```

In the absence of a command, R prints

```
flights %>%
  group_by(dest) %>%
  summarise(
    delay = mean(dep_delay, na.rm = TRUE),
    n = n()
) %>%
filter(n > 100) %>%
arrange(desc(delay)) %>%
print()
```

Another common command is **assign**

```
flights %>%
  group_by(dest) %>%
  summarise(
    delay = mean(dep_delay, na.rm = TRUE),
    n = n()
) %>%
filter(n > 100) %>%
arrange(desc(delay)) ->
dest_delays
```

But leading with assignment improves readability

```
dest_delays <- flights %>%
  group_by(dest) %>%
  summarise(
    delay = mean(dep_delay, na.rm = TRUE),
    n = n()
  ) %>%
  filter(n > 100) %>%
  arrange(desc(delay))
```

Functions fit best into a pipe when:

1. The first argument is the “data”
2. The data is the same type across a family of functions

Tidy data

Goal: Solve complex problems by combining simple, **uniform pieces**.

Tidy data is a consistent way of storing data

1. Each **dataset** goes in a **data frame**.

2. Each **variable** goes in a **column**.

Happy families are all alike;
every unhappy family is
unhappy in its own way

— *Leo Tolstoy*

Tidy datasets are all alike;
every messy dataset is
messy in its own way

— *Hadley Wickham*

Messy data has a varied shape

```
# A tibble: 5,769 × 22
   iso2  year   m04  m514  m014 m1524 m2534 m3544 m4554 m5564   m65    mu   f04  f514  f014 f1524
   <chr> <int> <int> <int> <int> <int> <int> <int> <int> <int> <int> <int> <int> <int> <int> <int>
1  AD    1989    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA
2  AD    1990    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA
3  AD    1991    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA
4  AD    1992    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA
5  AD    1993    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA
6  AD    1994    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA    NA
7  AD    1996    NA    NA     0     0     0     4     1     0     0    NA    NA    NA     0     1
8  AD    1997    NA    NA     0     0     1     2     2     1     6    NA    NA    NA     0     1
9  AD    1998    NA    NA     0     0     0     1     0     0     0    NA    NA    NA    NA    NA
10 AD    1999    NA    NA     0     0     0     1     1     0     0    NA    NA    NA     0     0
11 AD    2000    NA    NA     0     0     1     0     0     0     0    NA    NA    NA    NA    NA
12 AD    2001    NA    NA     0    NA    NA     2     1    NA    NA    NA    NA    NA    NA    NA
13 AD    2002    NA    NA     0     0     0     1     0     0     0    NA    NA    NA     0     1
14 AD    2003    NA    NA     0     0     0     1     2     0     0    NA    NA    NA     0     1
15 AD    2004    NA    NA     0     0     0     1     1     0     0    NA    NA    NA     0     0
16 AD    2005     0     0     0     0     1     1     0     0     0     0     0     0     0     1
17 AD    2006     0     0     0     1     1     2     0     1     1     0     0     0     0     0
# ... with 5,752 more rows, and 6 more variables: f2[illegible]
#   f5564 <int>, f65 <int>, fu <int>
```

What are the variables in this dataset?

(Hint: f = female, u = unknown, 1524 = 15-24)

Tidy data has a uniform shape

```
# A tibble: 35,750 × 5
   country  year   sex   age       n
     <chr> <int> <chr> <chr> <int>
1       AD  1996     f   014       0
2       AD  1996     f  1524       1
3       AD  1996     f  2534       1
4       AD  1996     f  3544       0
5       AD  1996     f  4554       0
6       AD  1996     f  5564       1
7       AD  1996     f    65       0
8       AD  1996     m   014       0
9       AD  1996     m  1524       0
10      AD  1996     m  2534       0
# ... with 35,740 more rows
```

tidytext

by Julia Silge & David Robinson

The family of Dashwood had long been settled in Sussex. Their estate was large, and their residence was at Norland Park, in the centre of their property, where, for many generations, they had lived in so respectable a manner as to engage the general good opinion of their surrounding acquaintance.

— *Sense & Sensibility*, **Jane Austen**

tidytext provides an answer

```
# A tibble: 724,880 × 4
                  book linenumber chapter     word
                <fctr>      <int>   <int>    <chr>
1  Sense & Sensibility         10       1  chapter
2  Sense & Sensibility         10       1        1
3  Sense & Sensibility         13       1      the
4  Sense & Sensibility         13       1   family
5  Sense & Sensibility         13       1       of
6  Sense & Sensibility         13       1 dashwood
7  Sense & Sensibility         13       1      had
8  Sense & Sensibility         13       1     long
9  Sense & Sensibility         13       1     been
10 Sense & Sensibility         13       1  settled
# ... with 724,870 more rows
```

Sentiment of Jane Austen books

Sense & Sensibility | Pride & Prejudice | Mansfield Park

Emma | Northanger Abbey | Persuasion

sentiment

50
25
0
−25
−50

0 50 100 150

0 50 100

0 50 100 150

0 50 100 150

0 20 40 60 80

0 20 40 60 80

Text mining, the tidy way

Julia Silge

Friday • 3:51pm

http://tidytextmining.com

list-cols

Tidy tibbles are better than tidy data frames

1. Each **dataset** goes in a tibble.

2. Each **variable** goes in a column.

Tibbles are data frames that are **lazy** & **surly**

```
df <- tibble(xyz = "a")

df$x
#> Warning: Unknown column 'x'
#> NULL

df$xyz
#> [1] "a"
```

But also have better support for list-cols

```
data.frame(x = list(1:2, 3:5))
#> Error: arguments imply differing number
#> of rows: 2, 3
```

```
tibble(x = list(1:2, 3:5))
#> # A tibble: 2 x 1
#>           x
#>      <list>
#> 1 <int [2]>
#> 2 <int [3]>
```

List-columns keep related things together

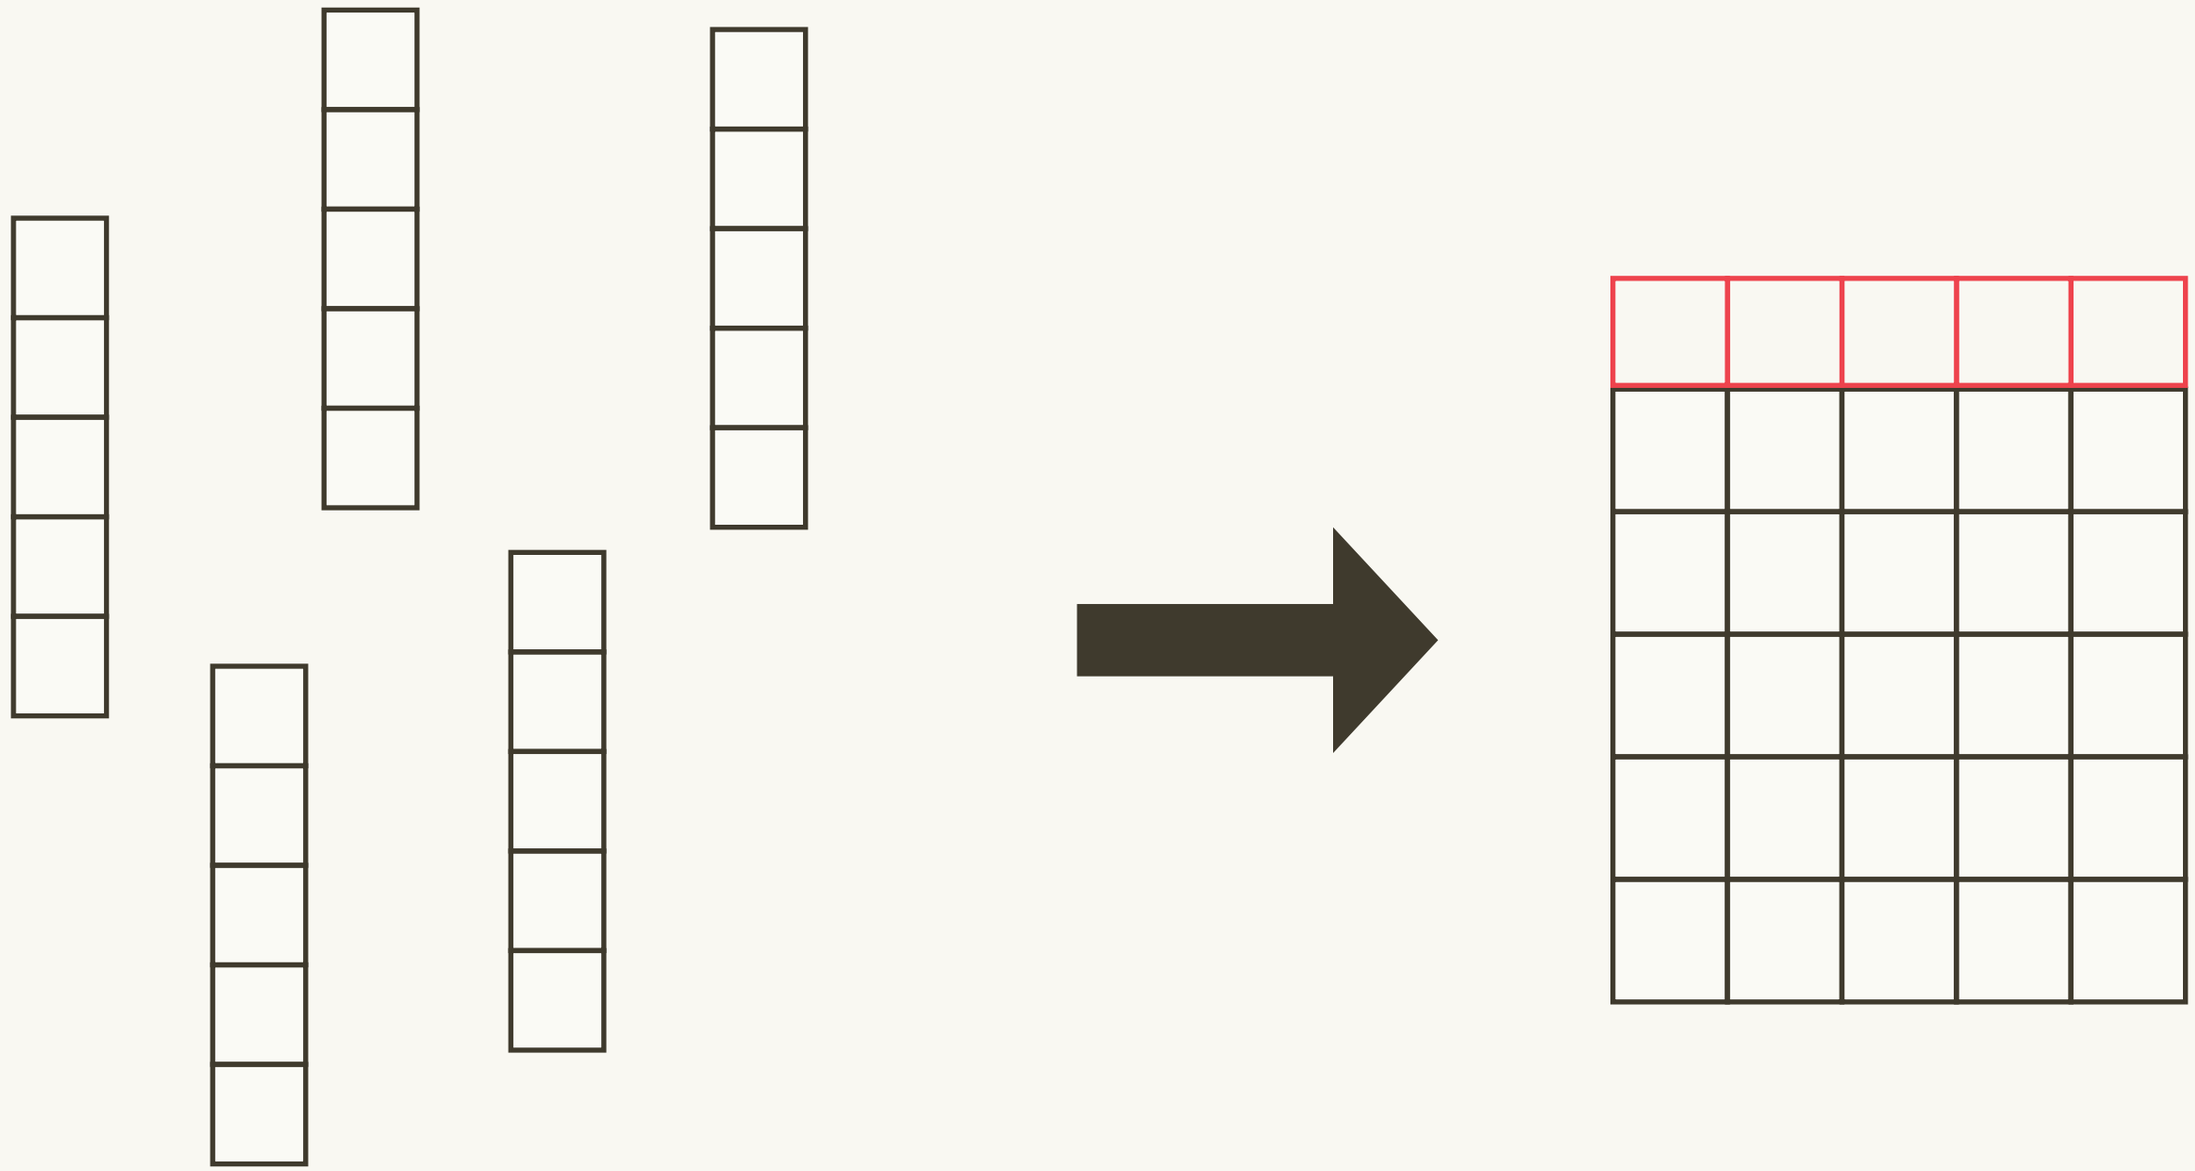

Anything can go in a list & a list can go in a data frame

sf (successor to sp) uses list-cols

by Edzer Pebesma

```
nc <- sf::st_read(system.file("shape/nc.shp", package = "sf"))

nc %>%
  as_tibble() %>%
  select(NAME, FIPS, AREA, geometry)
#> # A tibble: 100 × 4
#>          NAME   FIPS  AREA          geometry
#>        <fctr> <fctr> <dbl>  <simple_feature>
#> 1         Ashe  37009 0.114 <MULTIPOLYGON...>
#> 2    Alleghany  37005 0.061 <MULTIPOLYGON...>
#> 3        Surry  37171 0.143 <MULTIPOLYGON...>
#> 4    Currituck  37053 0.070 <MULTIPOLYGON...>
#> 5  Northampton  37131 0.153 <MULTIPOLYGON...>
#> 6     Hertford  37091 0.097 <MULTIPOLYGON...>
#> 7       Camden  37029 0.062 <MULTIPOLYGON...>
#> 8        Gates  37073 0.091 <MULTIPOLYGON...>
#> 9       Warren  37185 0.118 <MULTIPOLYGON...>
#> 10      Stokes  37169 0.124 <MULTIPOLYGON...>
#> # ... with 90 more rows
```

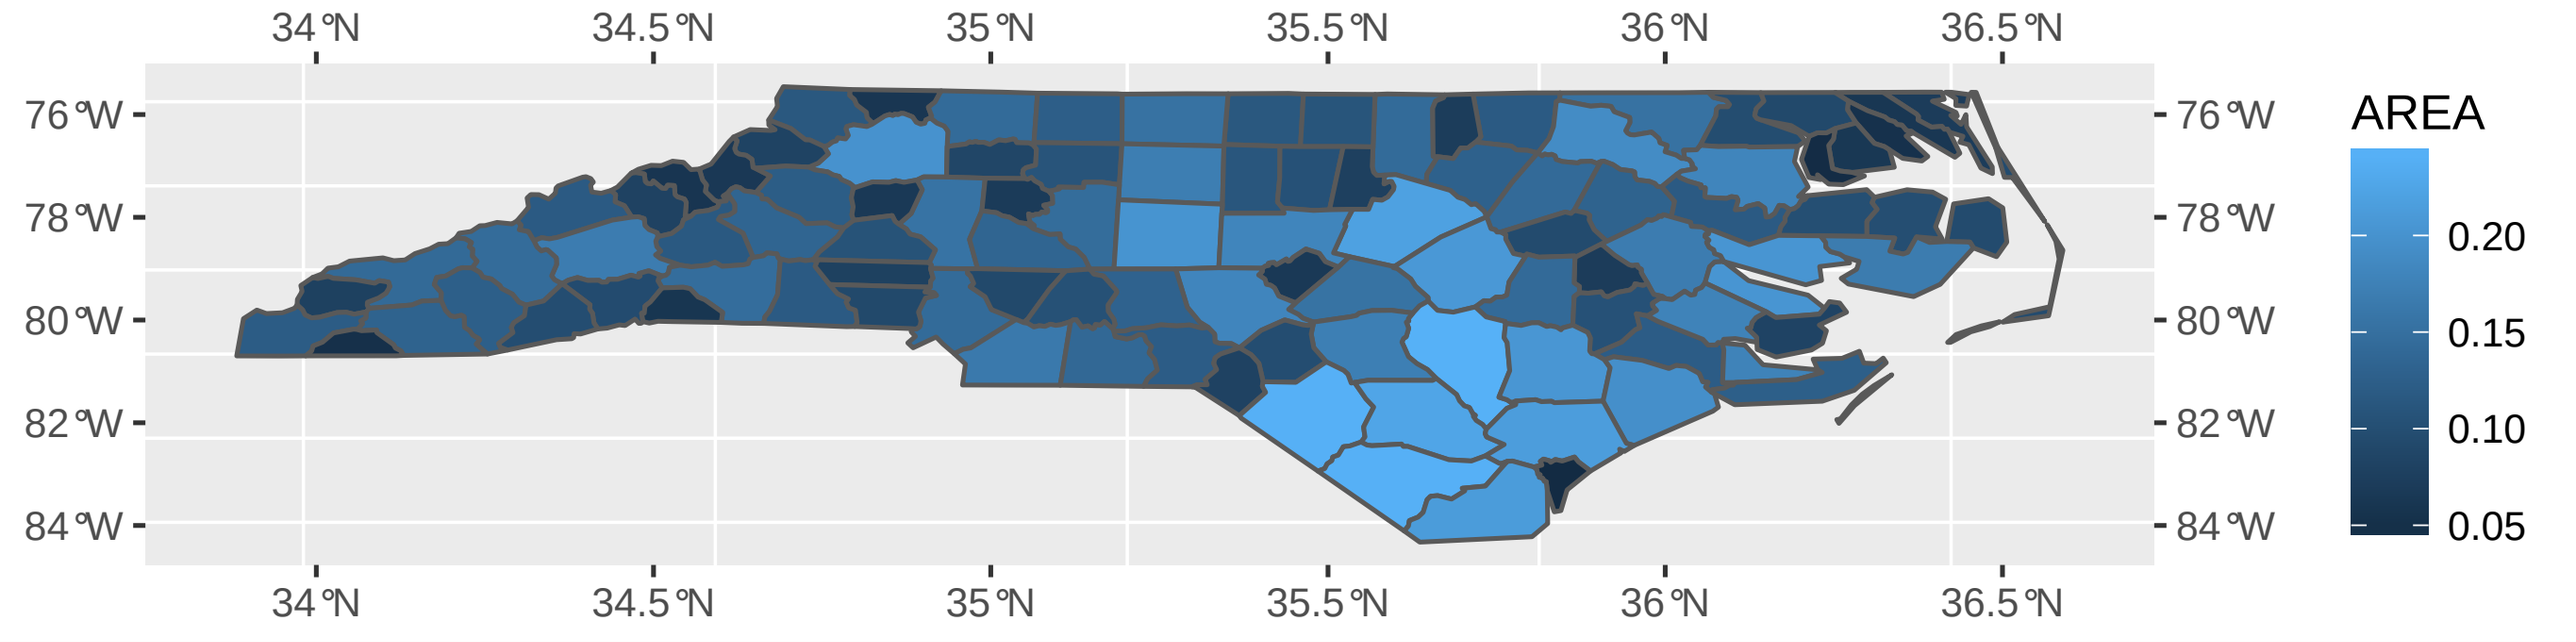


```
nc %>%
  ggplot(aes(geometry = geometry, fill = AREA)) +
    geom_sf() +
    coord_sf(crs = sf::st_crs(nc))
```

list-cols are a beautiful fit to cross-validation

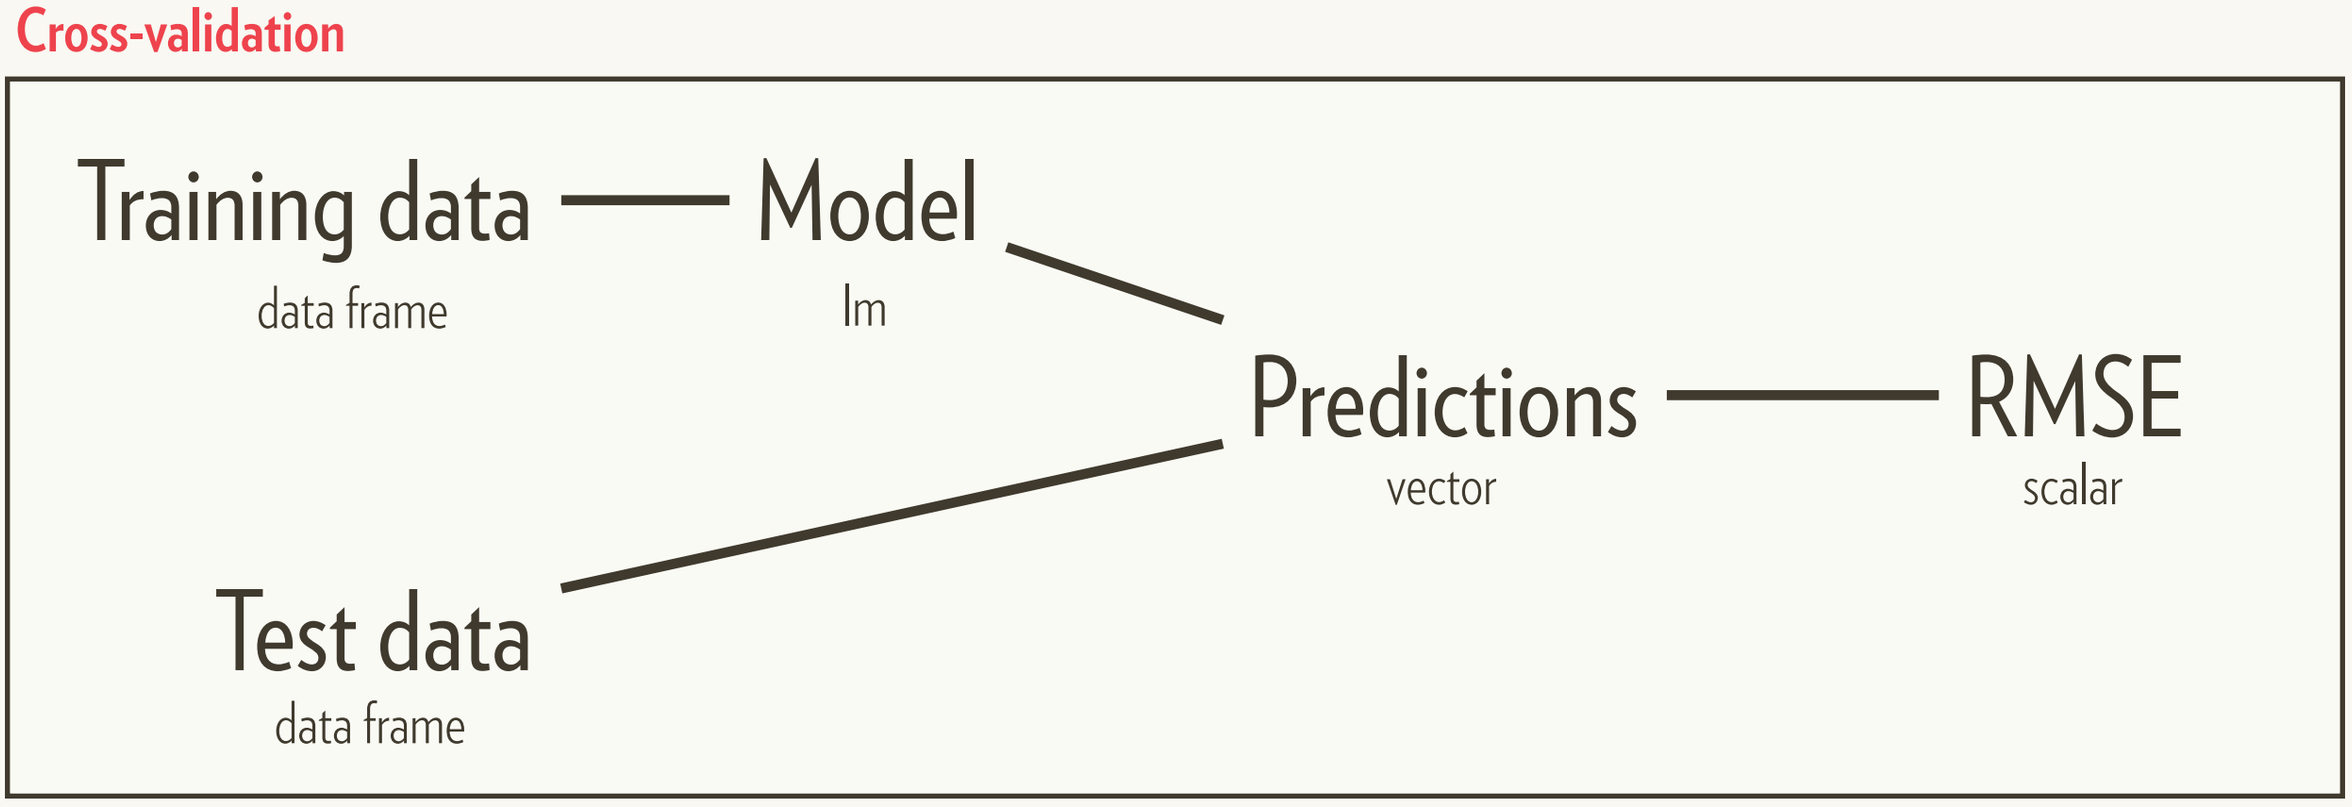

Each resample becomes one row

```
# A tibble: 100 x 5
            train              test   .id      mod      rmse
           <list>            <list> <chr>   <list>     <dbl>
1  <S3: resample> <S3: resample>   001 <S3: lm> 0.5661605
2  <S3: resample> <S3: resample>   002 <S3: lm> 0.2399357
3  <S3: resample> <S3: resample>   003 <S3: lm> 3.5482986
4  <S3: resample> <S3: resample>   004 <S3: lm> 0.2396810
5  <S3: resample> <S3: resample>   005 <S3: lm> 0.1591336
6  <S3: resample> <S3: resample>   006 <S3: lm> 0.1934869
7  <S3: resample> <S3: resample>   007 <S
8  <S3: resample> <S3: resample>   008 <S
9  <S3: resample> <S3: resample>   009 <S
10 <S3: resample> <S3: resample>   010 <S
... with 90 more rows
```

Putting square pegs in round holes
Jenny Bryan
Friday • 3:31pm

Conclusion

Four important facts:

1. It exists
2. It has a website
3. It has a package
4. It has a book

Four underlying principles:

1. Each function encapsulates one task
2. And is either a query or a command
3. Functions are composed with %>%
4. And use tidy tibbles as primary data structure

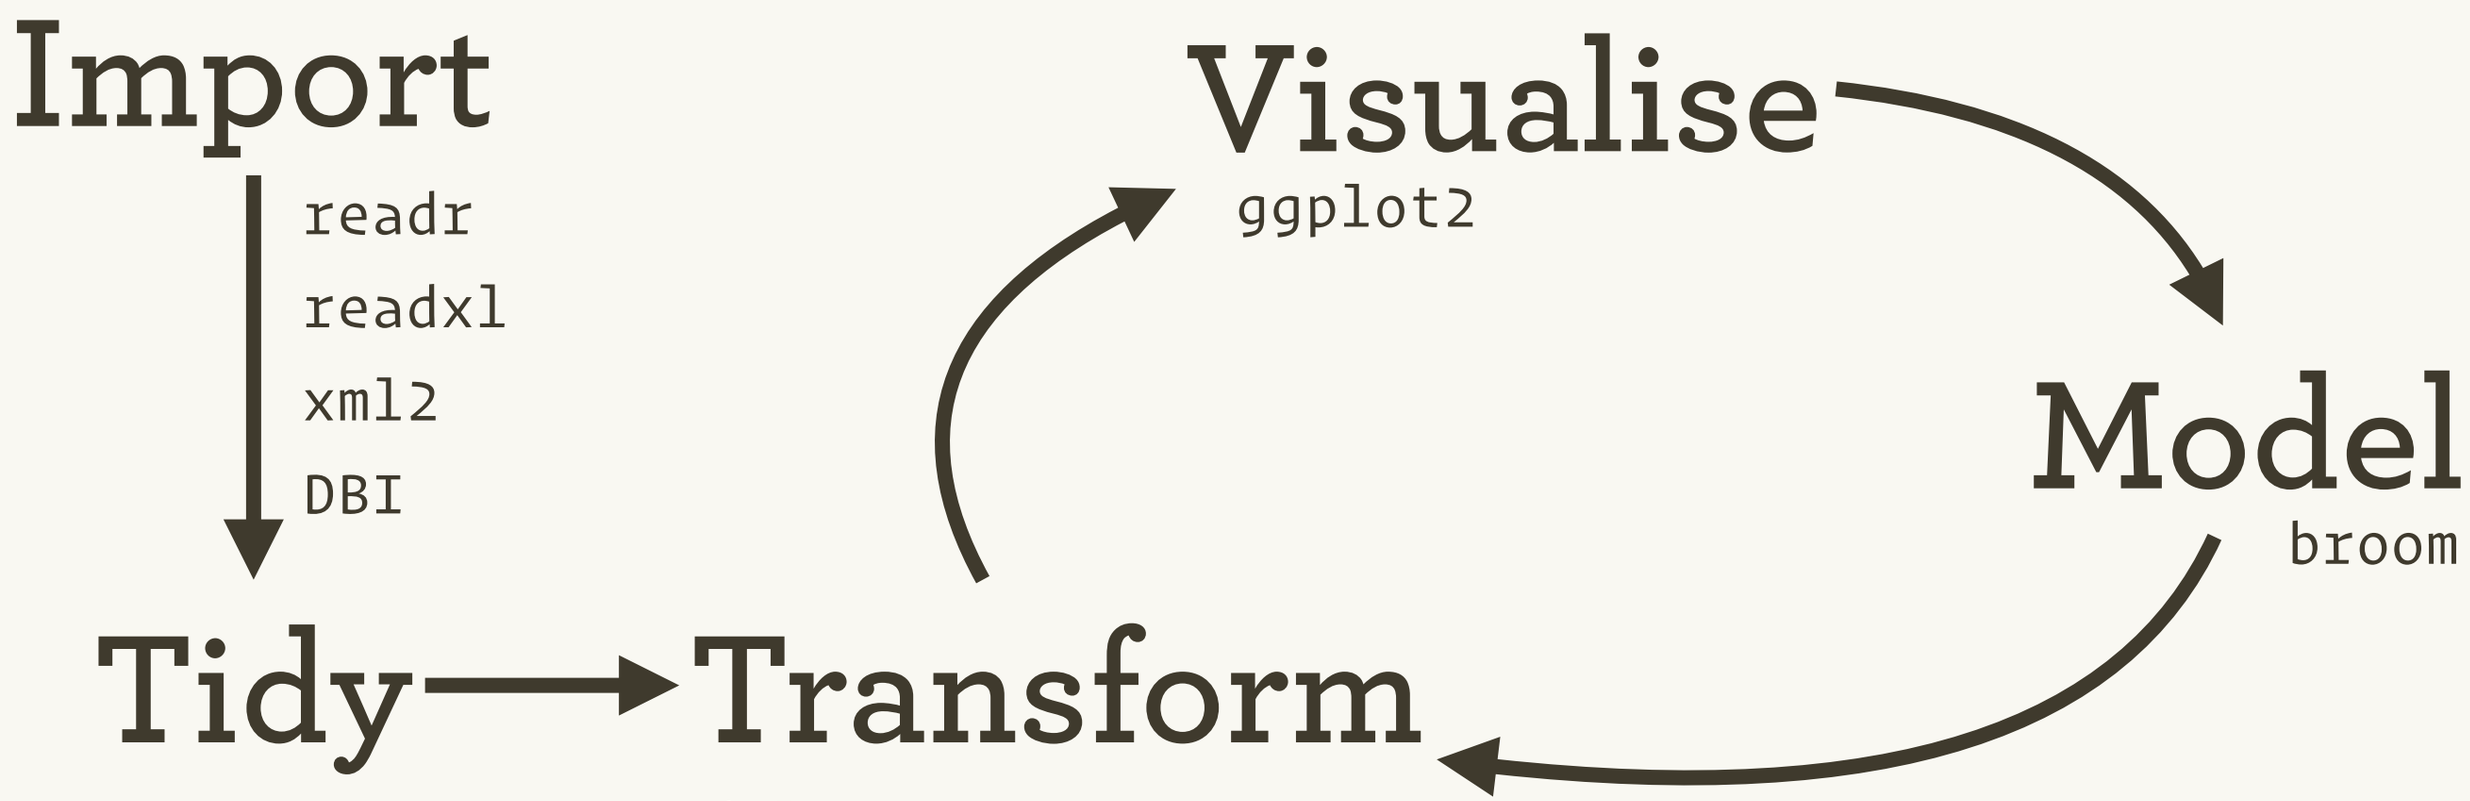

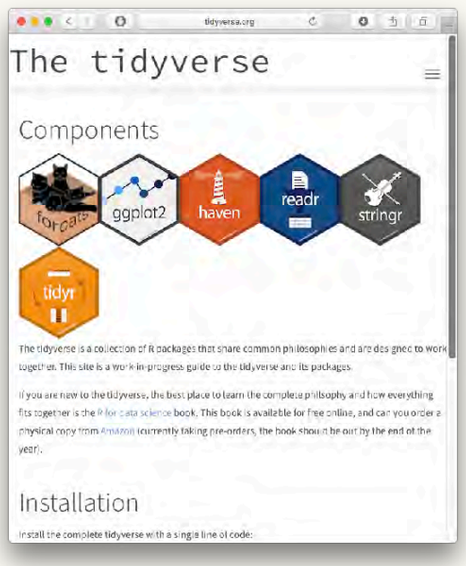

tidyverse.org

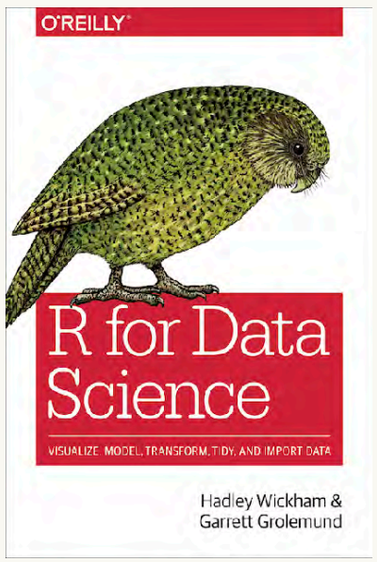

r4ds.had.co.nz

Tutorials

Functional programming

Charlotte Wickham

```
install.packages("tidyverse")
```

Shiny dashboards

Winston Chang & Joe Cheng

```
devtools::install_github("jcheng5/dashtutorial")
```