

# Resampling, Cross-validation & Bootstrap

ISLR2 Chapter 05

## Cross Validation

- The Validation Set Approach
- Leave-One-Out Cross Validation
- K-fold Cross Validation
- Bias-Variance Trade-off for k-fold Cross Validation
- Cross Validation on Classification Problems

# What are resampling methods?

---

## Tools that involves repeatedly drawing samples

- from a training set
- and refitting a model of interest on each sample
- in order to obtain more information about the fitted model

### Model Assessment:

- estimate test error rates

### Model Selection:

- select the appropriate level of model flexibility

## They are computationally expensive!

- But these days we have powerful computers 😊

## Two resampling methods:

- Cross Validation
- Bootstrapping

## 5.1.1 Typical Approach: The Validation Set Approach

**Suppose that we would like to find a set of variables**


- that give the lowest test (not training) error rate

**If we have a large data set,**

- we can achieve this goal by randomly splitting the data
- into training and validation(testing) parts

**We would then use the training part**

- to build each possible model
  - (i.e. the different combinations of variables)
- and choose the model that gave the lowest error rate
  - when applied to the validation data



## Example: Auto Data

---

Suppose that we want to predict mpg from horsepower

Two models:

- $\text{mpg} \sim \text{horsepower}$
- $\text{mpg} \sim \text{horsepower} + \text{horsepower}^2$

Which model gives a better fit?

- Randomly split Auto data set  
into training (196 obs.) and validation data (196 obs.)
- Fit both models using the training data set
- Then, evaluate both models  
using the validation data set
- The model with the lowest validation (testing) MSE is the winner!

## Results: Auto Data



**Left: Validation error rate for a single split**

**Right: Validation method repeated 10 times,**

- each time the split is done randomly!

**There is a lot of variability among the MSE's...**

- **Not good! We need more stable methods!**



## Advantages:

- Simple
- Easy to implement

## Disadvantages:

- The validation MSE can be highly variable
- Only a subset of observations are used to fit the model (training data).  
Statistical methods tend to perform worse  
when trained on fewer observations

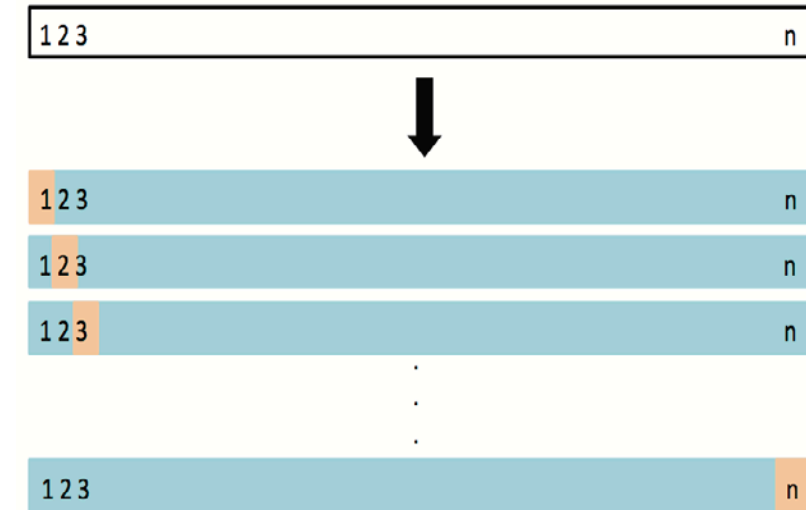
## 5.1.2 Leave-One-Out Cross Validation (LOOCV)

**This method is similar to the Validation Set Approach,  
but it tries to address the latter's disadvantages**

**For each suggested model, do:**

- Split the data set of size  $n$  into  
Training data set (blue) size:  $n - 1$   
Validation data set (beige) size: 1
- Fit the model using the training data
- Validate model using the validation data,  
and compute the corresponding MSE
- Repeat this process  $n$  times
- The MSE for the model is computed as follows:

$$CV_{(n)} = \frac{1}{n} \sum_{i=1}^n MSE_i.$$





# LOOCV vs. the Validation Set Approach

---

## LOOCV has less bias

- We repeatedly fit the statistical learning method using training data that contains  $n-1$  obs., i.e. almost all the data set is used

## LOOCV produces a less variable MSE

- The validation approach produces different MSE when applied repeatedly due to randomness in the splitting process,
- while performing LOOCV multiple times will always yield the same results, because we split based on 1 obs. each time

## LOOCV is computationally intensive (disadvantage)

- We fit the each model  $n$  times!

## 5.1.3 k-fold Cross Validation

---

**LOOCV is computationally intensive,**

- so we can run k-fold Cross Validation instead

**With k-fold Cross Validation,**

- we divide the data set into K different parts
- (e.g.  $K = 5$ , or  $K = 10$ , etc.)

**We then remove the first part,**

- fit the model on the remaining  $K-1$  parts,
- and see how good the predictions are on the left out part
- (i.e. compute the MSE on the first part)

**We then repeat this K different times**


- taking out a different part each time

**By averaging the K different MSE's**

- we get an estimated validation (test) error rate
- for new observations

$$CV_{(k)} = \frac{1}{k} \sum_{i=1}^k MSE_i.$$

# K-fold Cross Validation



**Left: LOOCV error curve**

**Right: 10-fold CV was run many times,**



- and the figure shows the slightly different CV error rates

**LOOCV is a special case of k-fold,**

- where  $k = n$

**They are both stable,**

- but LOOCV is more computationally intensive!





# Auto Data: Validation Set Approach vs. K-fold CV Approach

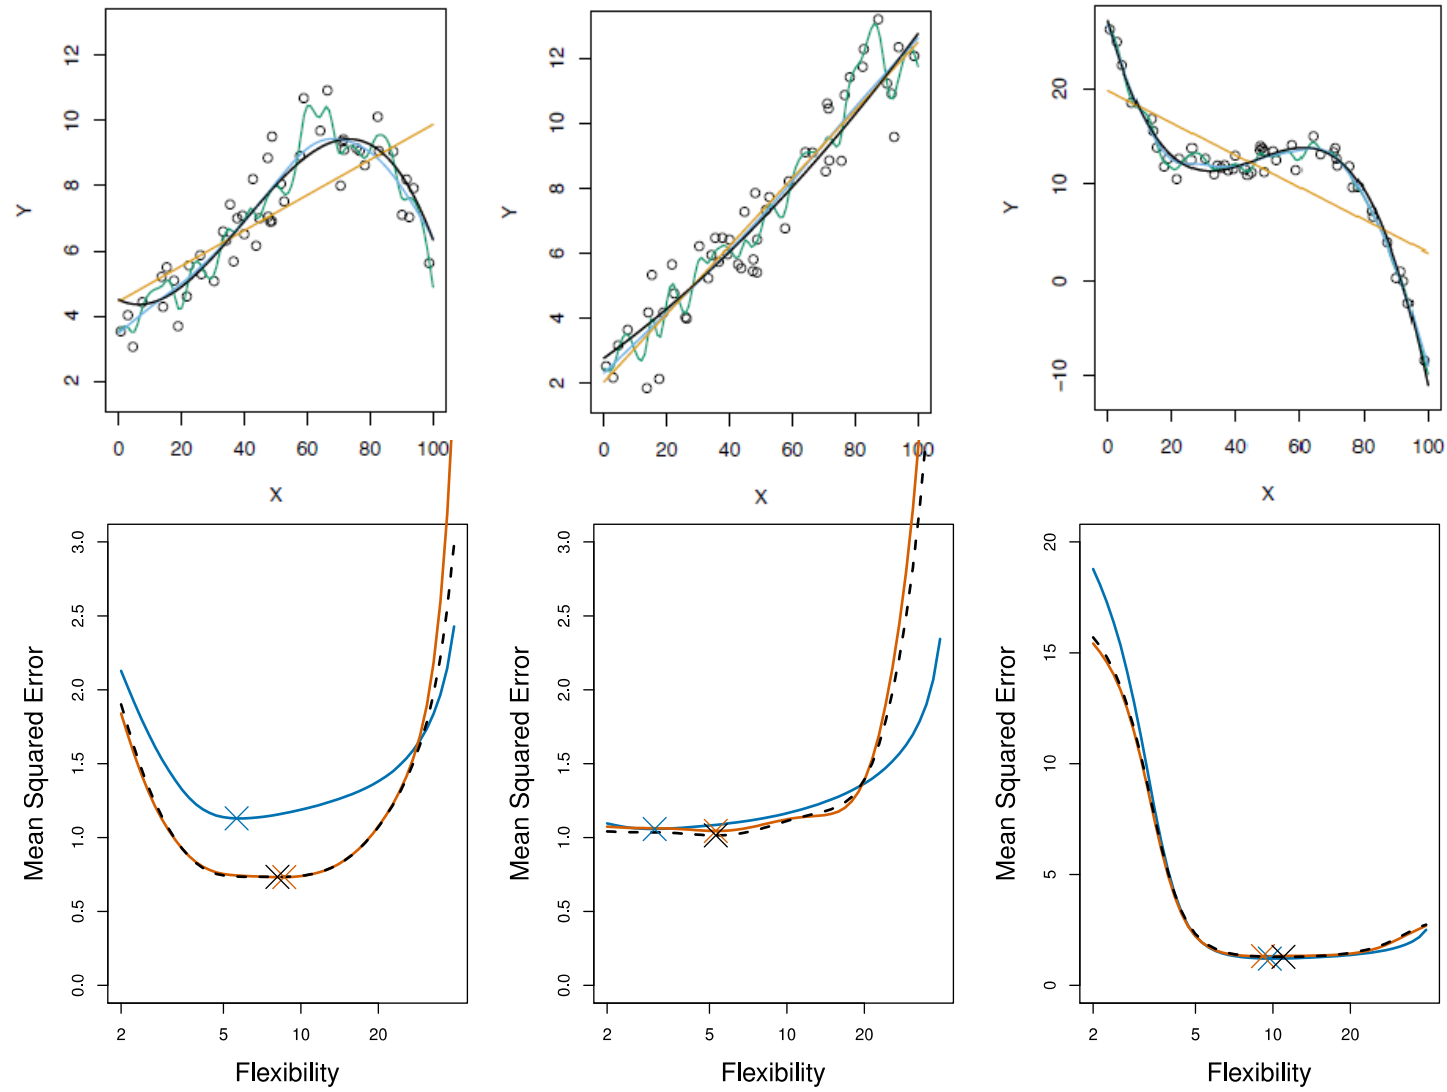
**Left: Validation Set Approach**

**Right: 10-fold Cross Validation Approach**

**Indeed, 10-fold CV is more stable!**



# K-fold Cross Validation on Three Simulated Data



**Blue: True Test MSE**

**Black: LOOCV MSE**

**Orange: 10-fold MSE**

**Refer to chapter 2 for the top graphs, Fig 2.9, 2.10, and 2.11**

## 5.1.4 Bias- Variance Trade-off for k-fold CV

---

**Putting aside that LOOCV is more computationally intensive than k-fold CV...**

**Which is better LOOCV or K-fold CV?**

- LOOCV is less biased than k-fold CV (when  $k < n$ )
- But, LOOCV has higher variance than k-fold CV (when  $k < n$ )
- Thus, there is a trade-off between what to use

**Conclusion:**

- We tend to use k-fold CV with ( $K = 5$  and  $K = 10$ )  
These are the magical K's
- It has been empirically shown that they yield test error rate estimates that suffer neither from excessively high bias, nor from very high variance

## 5.1.5 Cross Validation on Classification Problems

---

**So far, we have been dealing with CV on regression problems**

**We can use cross validation in a classification situation in a similar manner**


- Divide data into  $K$  parts
- Hold out one part,  
fit using the remaining data  
and compute the error rate on the hold out data
- Repeat  $K$  times
- CV error rate is the average over the  $K$  errors we have computed



# CV to Choose Order of Polynomial

The data set used is simulated (refer to Fig 2.13)

The purple dashed line is the Bayes' boundary



Bayes' Error Rate: 0.133


Naive Bayes classifiers are a family of simple "probabilistic classifiers" based on applying Bayes' theorem with strong (naïve) independence assumptions between the features (see Bayes classifier). They are among the simplest Bayesian network models, but coupled with kernel density estimation, they can achieve higher accuracy levels.

Naïve Bayes classifiers are highly scalable, requiring a number of parameters linear in the number of variables (features/predictors) in a learning problem. Maximum-likelihood training can be done by evaluating a closed-form expression, which takes linear time, rather than by expensive iterative approximation as used for many other types of classifiers.

## Linear Logistic regression (Degree 1)


- is not able to fit the Bayes' decision boundary

## Quadratic Logistic regression does better than linear




## Using cubic and quartic predictors,

- the accuracy of the model improves

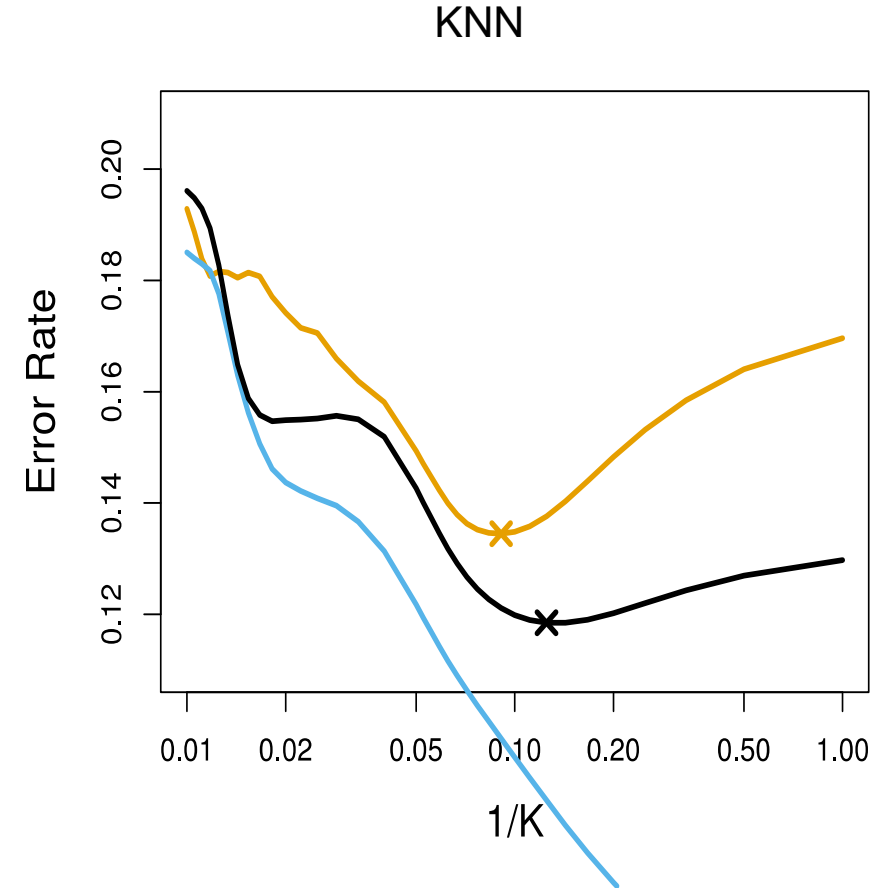



Error Rate: 0.160



Error Rate: 0.162

# CV to Choose the Order



**Brown: Test Error**

**Blue: Training Error**

**Black: 10-fold CV Error**