

DSCI353-353m-453: Class 04a p Multilevel Modeling

Profs: R. H. French, L. S. Bruckman, P. Leu, K. Davis, S. Cirlos

TAs: W. Oltjen, K. Hernandez, M. Li, M. Li, D. Colvin

09 February, 2023

Contents

4.1.2.1	Reading, Homeworks, Projects, SemProjects	1
4.1.2.2	Textbooks	2
4.1.2.2.1	Introduction to R and Data Science	2
4.1.2.2.2	Textbooks for this class	2
4.1.2.3	Tidyverse Cheatsheets, Functions and Reading Your Code	2
4.1.2.4	Syllabus	3
4.1.2.5	Intro to Frequentist (Multilevel) Generalised Linear Models (GLM) in R with glm and lme4	3
4.1.2.5.1	More readings	3
4.1.2.5.2	The packages being used (all installed on Markov and ODS Desktop)	3
4.1.2.6	Introduction to GLM	5
4.1.2.6.1	Recall that in a linear regression model,	5
4.1.2.6.2	The GLM is the genearlised version of linear regression	5
4.1.2.7	Thai Educational Data	6
4.1.2.8	Data Preparation	7
4.1.2.8.1	Load necessary packages	7
4.1.2.8.2	Import Data	8
4.1.2.8.3	Data Processing	8
4.1.2.8.4	Inspect Missing Data	9
4.1.2.9	Binomial Logistic Regression	9
4.1.2.9.1	Explore Data: number of REPEAT by SEX and PPED	9
4.1.2.9.2	Fit a Binary Logistic Regression Model	10
4.1.2.9.3	Interpretation	10
4.1.2.9.4	Visualisation of Parameter Effects	11
4.1.2.9.5	Model Evaluation: Goodness of Fit	13
4.1.2.9.6	AUC (area under the curve).	15
4.1.2.10	Binomial Logistic Regression	16
4.1.2.10.1	Transform Data	16
4.1.2.10.2	Explore Data	17
4.1.2.10.3	Fit a Binomial Logistic Regression Model	17
4.1.2.10.4	Interpretation	18
4.1.2.11	Multilevel Binary Logistic Regression	19
4.1.2.11.1	Center Variables	21
4.1.2.11.2	Intercept Only Model	21
4.1.2.11.3	Full Model	23
4.1.2.12	Other Family (Distribution) and Link Functions	28
4.1.2.13	Links	29

4.1.2.1 Reading, Homeworks, Projects, SemProjects

- Readings:
 - For today: ISLR5
 - For Thursday: ISLR6 (R4DS9-16)
 - Next is: Deep Learning with R (2nd Ed.)
- Laboratory Exercises:
 - LE2 is Due next Tuesday Feb. 14th
- Office Hours: (Class Canvas Calendar for Zoom Link)
 - Wednesdays @ 4:00 PM to 5:00 PM
 - Saturdays @ 3:00 PM to 4:00 PM
 - **Office Hours are on Zoom, and recorded**
- Semester Projects
 - Office Hours for SemProjs: Mondays at 4pm on Zoom
 - DSCI 453 Students Biweekly Updates Due
 - * Update # is Due ** **
 - DSCI 453 Students
 - * Next Report Out # is Due ** **
 - All DSCI 353/353M/453, E1453/2453 Students:
 - * Peer Grading of Report Out #1 is Due ** **
 - Exams
 - * MidTerm: **Thursday March 9th**, in class or remote, 11:30 - 12:45 PM
 - * Final: **Thursday May 4th**, 2023, 12:00PM - 3:00PM, Nord 356 or remote

4.1.2.2 Textbooks

4.1.2.2.1 Introduction to R and Data Science For students new to R, Coding, Inferential Statistics

- Peng: R Programming for Data Science
- Peng: Exploratory Data Analysis with R
- OIS = Diez, Barr, Çetinkaya-Runde: Open Intro Stat v4

4.1.2.2.2 Textbooks for this class

- R4DS = Wickham, Golemund: R for Data Science
- ISLR = James, Witten, Hastie, Tibshirani: Intro to Statistical Learning with R
- ESL = Trevor Hastie, Tibshirani, Friedman: Elements of Statistical Learning
- DLwR = Chollet, Allaire: Deep Learning with R

DL1 to DL6 are “Deep Learning” articles in 3-readings/2-articles/

4.1.2.3 Tidyverse Cheatsheets, Functions and Reading Your Code Look at the Tidyverse Cheat-sheet

- **Tidyverse For Beginners Cheatsheet**
 - In the Git/20s-dsci353-353m-453-prof/3-readings/3-CheatSheets/ folder
- **Data Wrangling with dplyr and tidyr Cheatsheet]**

Tidyverse Functions & Conventions

- The pipe operator `%>%`
- Use `dplyr::filter()` to subset data row-wise.
- Use `dplyr::arrange()` to sort the observations in a data frame
- Use `dplyr::mutate()` to update or create new columns of a data frame
- Use `dplyr::summarize()` to turn many observations into a single data point
- Use `dplyr::arrange()` to change the ordering of the rows of a data frame
- These can be combined using `dplyr::group_by()`

- which lets you perform operations “by group”.
- The `%in%` matches conditions provided by a vector using the `c()` function

Reading Your Code: Whenever you see

- The assignment operator `<-`, think “**gets**”
- The pipe operator, `%>%`, think “**then**”

4.1.2.4 Syllabus

4.1.2.5 Intro to Frequentist (Multilevel) Generalised Linear Models (GLM) in R with `glm` and `lme4`

- This will provide a basic introduction to generalized linear models (GLM)
 - using the frequentist approach.

Specifically, we’ll focus on the use of logistic regression

- in both binary-outcome and
 - count/proportion-outcome scenarios,
- and the respective approaches to model evaluation.

We’ll use the Thai Educational Data example

- from Chapter 6 of the book - Multilevel analysis: Techniques and applications [1]. - by Joop Hox, Mirjam Moerbeek, and Rens van de Schoot

Furthermore, we’ll demonstrate

- the multilevel extension of GLM models
 - with the `lme4` package in R.

Lastly, more distributions and link functions

- in the GLM framework are discussed.

4.1.2.5.1 More readings

- This is meant for beginners and we won’t
 - delve into technical details and complex models.

For a detailed introduction into frequentist multilevel models,

- see [this LME4 Tutorial](#).

For an extensive overview of GLM models,

- see the [Wikipedia GLM article here](#).

Also `glmnet` is a very nice “glm” package

- and has a [good glmnet site](#)

The Bayesian version of this tutorial

- can also be [found here](#).

4.1.2.5.2 The packages being used (all installed on Markov and ODS Desktop)

- We’ll be using
 - `lme4` for multilevel modelling (this tutorial uses version 1.1-31);
 - `tidyverse` for data manipulation and plotting with `ggplot2`;

Day:Date	Foundation	Practicum	Readings(optional)	Due(optional)
w01a:Tu:1/17/23 w01b:Th:1/19/23	Markov Cluster Stat. Learning, Approach	R, Rstudio IDE, Git Bash, Git, Class Repo	ISLR1,2 (R4DS-1-3)	(LE0)
w02a:Tu:1/24/23 w02b:Th:1/26/23	Lin. Regr. Bias-Var. Train/Test, Bias vs. Vari.	SemProjs; Regr. Ovrvw Tidyverse Review	ISLR3,(R4DS-4-6) DL01 DL02 (R4DS-7,8)	(LE0:Due) LE1
w02Pr:Fr:1/27/23	ADD DROP	DEADLINE		453 Update 1
w03a:Tu:1/31/23	Logistic Regr. Classif	Pred. Analytics, Regr.	DL03,ISLR4	
w03b:Th:2/2/23 w03Sa:2/4/23	LDA/QDA	ggPlot2, Code Expect.	DL04, DL05	LE1:Due, LE2 LE1:Due
w04a:Tu:2/7/23 w04b:Th:2/9/23	Resample Cross-Valid. DL, ML Overview	Multilevel Mod. ML with NNs	ISLR5 ISLR6 (R4DS9-16)	
w04Pr:Fr:2/10/23				453 Update 2
w05a:Tu:2/14/23	Bootstrap		DL2R1, DL06,07	LE2:Due, LE3
w05b:Th:2/16/23 w05Pr:Fr:2/17/23	Subset Selec., Shrink.	Mixed Effects	DLwR2	453 Rep. Out 1
w06a:Tu:2/21/23 w06b:Th:2/23/23	Mod. Selec. Beyond Linear Modls	Dim. Red. Feature Select., Caret	ISLR7 ISLR10, (R4DS22-25)	LE3:Due, LE4
w06Pr:Fr:2/24/23				453 Update 3
w07a:Tu:2/28/23	Dec. Trees, Rand. Forest.	Tidy Modeling	ISLR8, DL08,09	
w07b:Th:3/2/23	MidTerm Review, SVM	SVM, SVR, ROC	ISLR9 (R4DS26-30)	Peer Review 1
w08a:Tu:3/7/23	R-Keras/TensorFlow2		DLwR1	
w08b:Th:3/9/23 w08Pr:Fr:3/10/23	MIDTERM EXAM		DL10,11	LE4:Due LE5 453 Update 4
Tu:3/14/23 Th:3/16/23	SPRING SPRING	BREAK BREAK	ISLR10 DL12,13	
w09a:Tu:3/21/23 w09b:Th:3/23/23	Deep Learning Computer Vision, CNN	TF2 Keras Intro CNN w/TF2, Overfit	Pocket Perceptron DLR4	ISLR10, DLR3
w09Pr:Fr:3/24/23				453 Rep. Out 2
w10a:Tu:3/28/23	Deep Learn Intro	NN Types	DLR5	
w10b:Th:3/30/23 w10Pr:Fr:3/31/23	DL CNN,RNN ImageNet	NN Types, CNN w/TF2	Hinton ImageNet	453 Upd.5 & PrRev 2
Sa:4/1/23				LE5:Due LE6
w11a:Tu:4/4/23	Fitting NNs	AUC,Prec,Recall Fruit		
w11b:Th:4/6/23	NLP, Graphs & ML		LeCun DL Rev. 2015	
w12a:Tu:4/11/23	Graphs & ML	NLP with sequences	DLR6	
w12b:Th:4/13/23	NLP w attention	Graph Repr Proc Wrk-flw		LE6:Due LE7
w13a:Tu:4/18/23	DL Frameworks	Explaining DL w Lime		
w13b:Th:4/20/23 w13Pr:Fr:4/21/23	Linux Distros XGBoost	Explain Preds	Deep Dream	453 Rep. Out 3 Due
w14a:Tu:4/25/23 w14b:Th:4/27/23	Transformers Final Exam Review	Torch NN & DeepLearn		LE7:Due
w14Pr:Fr:4/28/23				Peer Rev 3 Due
	FINAL EXAM	Th. 5/4/23, 12-3pm	Nord 356 & Zoom	
	453 Final PDF Report	Fr. 4/29, 11:59pm		

Figure 1: IT Fundamentals: Applied Data Science with R, Syllabus

- `haven` for reading `sav` format data;
- `jtools` for handling of model summaries;
- `ggstance` for visualisation purposes;
- `ROCR` for calculating area under the curve (AUC);
- `performance` for calculating intra-class correlation (ICC).
- `effects` for plotting parameter effects;
- Basic knowledge of hypothesis testing and statistical inference;
- Basic knowledge of correlation and regression;
- Basic knowledge of coding in R;
- Basic knowledge of plotting and data manipulation with `tidyverse`.

```
# install.packages(c("lme4", "tidyverse", "haven", "jtools", "ggstance", "ROCR"))
```

4.1.2.6 Introduction to GLM

- If you are already familiar with generalized linear models (GLM),
 - Its section 4.6 of ISLR2
 - you can skip to the next section.

Otherwise, here is a short introduction to GLM

4.1.2.6.1 Recall that in a linear regression model,

- the object is to model the expected value of a continuous variable, Y ,
- as a linear function of the predictor, $\epsilon = X\beta$.

The model structure is thus: $E(Y) = X\beta + \epsilon$,

- where ϵ refers to the residual error term.

The linear regression model assumes that Y

- is continuous and comes from a normal distribution,
- that ϵ is normally distributed and
- that the relationship between the linear predictor η and
- the expected outcome $E(Y)$ is strictly linear.

However, these assumptions are easily violated in many real world data examples,

- such as those with binary or proportional outcome variables and
- those with non-linear relationships between
 - the predictors and the outcome variable.

In these scenarios

- where linear regression models are clearly inappropriate,
- generalised linear models (GLM) are needed.

4.1.2.6.2 The GLM is the generalised version of linear regression

- that allows for deviations from the assumptions underlying linear regression.

The GLM generalises linear regression

- by assuming the dependent variable Y
 - to be generated from any particular distribution in an exponential family
 - (a large class of probability distributions that includes
 - the normal, binomial, Poisson and gamma distributions, among others).

In this way, the distribution of Y

- does not necessarily have to be normal.

In addition, the GLM allows the linear predictor η

- to be connected to the expected value of the outcome variable, $E(Y)$,
- via a link function $g(\cdot)$.

The outcome variable, Y , therefore, depends on η

- through $E(Y) = g^{-1}(\eta) = g^{-1}(X\beta)$.

In this way, the model does not assume

- a linear relationship between $E(Y)$ and η ;
- instead, the model assumes a linear relationship
 - between $E(Y)$ and the transformed $g^{-1}(\eta)$.

This tutorial focuses on the probably most popular example of GLM: **logistic regression**.

Logistic regression has two variants,

- the well-known **binary logistic regression**
 - that is used to model binary outcomes (1 or 0; “yes” or “no”),
- and the less-known **binomial logistic regression**
 - suited to model count/proportion data.

Binary logistic regression assumes that Y

- comes from a Bernoulli distribution,
- where Y only takes a value of
 - 1 (target event) or 0 (non-target event).

Binary logistic regression connects $E(Y)$ and η

- via the logit link $\eta = \text{logit}(\pi) = \log(\pi/(1 - \pi))$,
 - where π refers to the probability of the target event ($Y = 1$).

Binomial logistic regression, in contrast,

- assumes a binomial distribution underlying Y ,
 - where Y is interpreted as the number of target events,
 - can take on any non-negative integer value
 - and is binomially distributed with regards to
 - * n number of trials and
 - * π probability of the target event.

The link function is the same as that of binary logistic regression.

The next section details the example data (**Thai Educational Data**)

- followed by the demonstration of the use of
- both binary and binomial logistic regression.

4.1.2.7 Thai Educational Data

- The data used in this tutorial is the Thai Educational Data
 - that is also used as an example in Chapter 6
 - of Multilevel analysis: Techniques and applications.

The data can [be downloaded from here](#).

And its also in the `2-class/data` folder

The data stems from a national survey of primary education in Thailand

- (Raudenbush & Bhumirat, 1992).

Each row in the data refers to a pupil.

- The outcome variable REPEAT is a dichotomous variable
 - indicating whether a pupil has repeated a grade during primary education.
- The SCHOOLID variable indicates the school of a pupil.
- The person-level predictors include: SEX (0 = female, 1 = male)
- and PPED (having had preschool education, 0 = no, 1 = yes).
- The school-level is MESC,
 - representing school mean SES (socioeconomic status) scores.

The main research questions that this tutorial seeks to answer

- using the Thai Educational Data are:

Ignoring the clustering structure of the data,

- what are the effects of gender and preschool education
 - on whether a pupil repeats a grade?

Ignoring the clustering structure of the data,

- what is the effect of school mean SES
 - on the proportion of pupil repeating a grade?

Considering the clustering structure of the data,

- what are the effects of gender, preschool education and school mean SES
 - on whether a pupil repeats a grade?

These three questions are answered

- by using these following models, respectively:
 - binary logistic regression;
 - binomial logistic regression;
 - multilevel binary logistic regression.

4.1.2.8 Data Preparation

```
# if you don't have these packages installed yet, please use the install.packages("package_name") command
library(lme4) # for multilevel models
```

4.1.2.8.1 Load necessary packages

```
## Loading required package: Matrix
```

```
library(tidyverse) # for data manipulation and plots
```

```
## -- Attaching packages ----- tidyverse 1.3.2 --
```

```
## v ggplot2 3.4.0      v purrr   1.0.0
```

```
## v tibble  3.1.8      v dplyr  1.0.10
```

```
## v tidyr   1.2.1      v stringr 1.5.0
```

```
## v readr   2.1.3      v forcats 0.5.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x tidyr::expand() masks Matrix::expand()
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()    masks stats::lag()
```

```
## x tidyr::pack()   masks Matrix::pack()
```

```
## x tidyr::unpack() masks Matrix::unpack()
```

```
library(haven) #for reading sav data
library(sjstats) # used to be to calc icc. but now use performance::icc
library(effects) #for plotting parameter effects
```

```
## Loading required package: carData
## lattice theme set by effectsTheme()
## See ?effectsTheme for details.
```

```
library(jtools) #for transforming model summaries
library(ROCR) #for calculating area under the curve (AUC) statistics
```

```
# ThaiEdu_Raw <- read_sav("https://github.com/MultiLevelAnalysis/Datasets-third-edition-Multilevel-book,
ThaiEdu_Raw <- haven::read_sav("./data/thaieduc.sav")
head(ThaiEdu_Raw)
```

4.1.2.8.2 Import Data

```
## # A tibble: 6 x 5
##   SCHOOLID SEX      PPED      REPEAT      MSESC
##   <dbl> <dbl+lbl> <dbl+lbl> <dbl+lbl> <dbl>
## 1 10101 0 [girl] 1 [yes] 0 [no] NA
## 2 10101 0 [girl] 1 [yes] 0 [no] NA
## 3 10101 0 [girl] 1 [yes] 0 [no] NA
## 4 10101 0 [girl] 1 [yes] 0 [no] NA
## 5 10101 0 [girl] 1 [yes] 0 [no] NA
## 6 10101 0 [girl] 1 [yes] 0 [no] NA
```

```
ThaiEdu_New <- ThaiEdu_Raw %>%
  mutate(
    SCHOOLID = factor(SCHOOLID),
    SEX = if_else(SEX == 0, "girl", "boy"),
    SEX = factor(SEX, levels = c("girl", "boy")),
    PPED = if_else(PPED == 0, "no", "yes"),
    PPED = factor(PPED, levels = c("no", "yes"))
  )
head(ThaiEdu_New)
```

4.1.2.8.3 Data Processing

```
## # A tibble: 6 x 5
##   SCHOOLID SEX PPED REPEAT MSESC
##   <fct> <fct> <fct> <dbl+lbl> <dbl>
## 1 10101 girl yes 0 [no] NA
## 2 10101 girl yes 0 [no] NA
## 3 10101 girl yes 0 [no] NA
## 4 10101 girl yes 0 [no] NA
## 5 10101 girl yes 0 [no] NA
## 6 10101 girl yes 0 [no] NA
```

```
ThaiEdu_New %>%
  summarise_each(list( ~ sum(is.na(.)))) %>%
```



```
gather()
```

4.1.2.8.4 Inspect Missing Data

```
## # A tibble: 5 x 2
##   key      value
##   <chr>   <int>
## 1 SCHOOLID     0
## 2 SEX          0
## 3 PPED         0
## 4 REPEAT       0
## 5 MSEC        1066
```

The data has 1066 observations missing for the MSEC variable.

The treatment of missing data is a complicated topic on its own.

For the sake of convenience,

- we simply list-wise delete the cases with missing data in this tutorial.

```
ThaiEdu_New <- ThaiEdu_New %>%
  filter(!is.na(MSEC))
```

4.1.2.9 Binomial Logistic Regression

```
ThaiEdu_New %>%
  group_by(SEX) %>%
  summarise(REPEAT = sum(REPEAT))
```

4.1.2.9.1 Explore Data: number of REPEAT by SEX and PPED

```
## # A tibble: 2 x 2
##   SEX REPEAT
##   <fct> <dbl>
## 1 girl   428
## 2 boy    639
```

```
ThaiEdu_New %>%
  group_by(PPED) %>%
  summarise(REPEAT = sum(REPEAT))
```

```
## # A tibble: 2 x 2
##   PPED REPEAT
##   <fct> <dbl>
## 1 no    673
## 2 yes   394
```

It seems that the number of pupils who repeated a grade

- differs quite a bit between the two genders,
- with more male pupils having to repeat a grade.

More pupils who did not have preschool education

- repeated a grade.

This observation suggests that SEX and PPED

- might be predictive of REPEAT.

4.1.2.9.2 Fit a Binary Logistic Regression Model

- R has the base package installed by default,
 - which includes the `glm` function that runs GLM.

The arguments for `glm` are similar to those for `lm`: formula and data.

However, `glm` requires an additional argument:

- `family`, which specifies the assumed distribution of the outcome variable;
 - within family we also need to specify the link function.
- The default of family is `gaussian(link = "identity")`,
 - which leads to a linear model
 - that is equivalent to a model specified by `lm`.
- In the case of binary logistic regression,
 - `glm` requires that we specify a binomial distribution
 - with the logit link, namely `family = binomial(link = "logit")`.

```
Model_Binary <- glm(
  formula = REPEAT ~ SEX + PPED,
  family = binomial(link = "logit"),
  data = ThaiEdu_New
)
summary(Model_Binary)

##
## Call:
## glm(formula = REPEAT ~ SEX + PPED, family = binomial(link = "logit"),
##      data = ThaiEdu_New)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.6844  -0.5630  -0.5170  -0.4218   2.2199
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.76195     0.05798 -30.387 < 2e-16 ***
## SEXboy       0.42983     0.06760   6.358 2.04e-10 ***
## PPEDyes     -0.61298     0.06833  -8.971 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6140.8  on 7515  degrees of freedom
## Residual deviance: 6016.2  on 7513  degrees of freedom
## AIC: 6022.2
##
## Number of Fisher Scoring iterations: 4
```

4.1.2.9.3 Interpretation

- From the summary output above, we can see that
 - SEX positively and significantly
 - * predicts a pupil's probability of repeating a grade,
 - while PPED negatively and significantly so.

Specifically, in comparison to being a girl,

- being a boy is more likely to repeat a grade.
- Having previous schooling is less likely to result in repeating a grade.

To interpret the value of the parameter estimates,

- we need to exponentiate the estimates.

The `summ` function from the `jtools` packages

- provides an easy way to do so for any model fitted by `glm`. See below.

```
summ(Model_Binary, exp = T) # set "exp = T" to show exponentiated estimates; if you need standardised e
```

Observations	7516
Dependent variable	REPEAT
Type	Generalized linear model
Family	binomial
Link	logit

$\chi^2(2)$	124.55
Pseudo-R ² (Cragg-Uhler)	0.03
Pseudo-R ² (McFadden)	0.02
AIC	6022.21
BIC	6042.98

	exp(Est.)	2.5%	97.5%	z val.	p
(Intercept)	0.17	0.15	0.19	-30.39	0.00
SEXboy	1.54	1.35	1.75	6.36	0.00
PPEDyes	0.54	0.47	0.62	-8.97	0.00

Standard errors: MLE

Note that the interpretation of the parameter estimates

- is linked to the odds rather than probabilities.

The definition of odds is:

- $P(\text{event occurring})/P(\text{event not occurring})$.

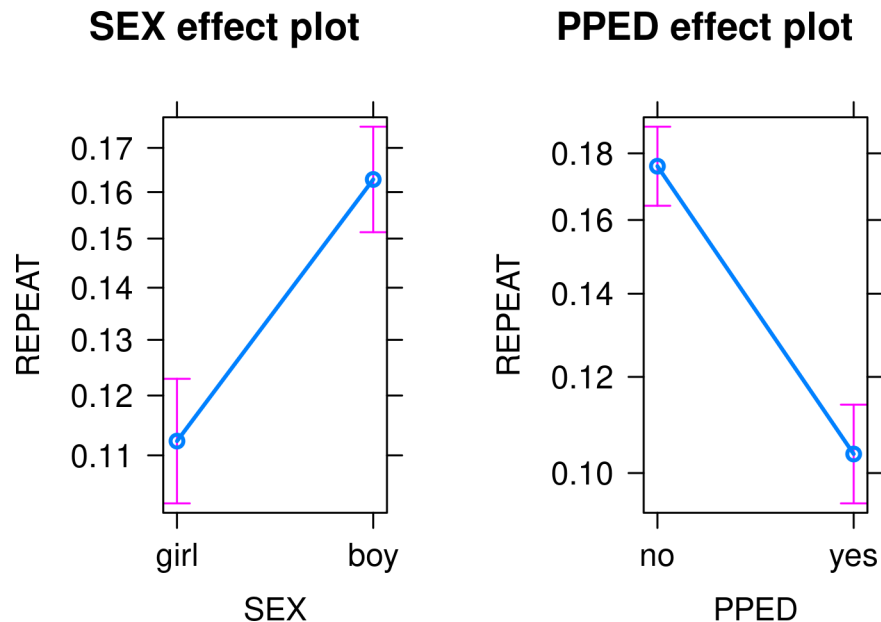
In this analysis, assuming everything else stays the same,

- being a boy increases the odds of repeating a grade by 54%,
 - in comparison to being a girl;
- having preschool education lowers the odds of repeating a grade
 - by $(1 - 0.54)\% = 46\%$,
 - in comparison to not having preschool education,
 - assuming everything else stays constant.

4.1.2.9.4 Visualisation of Parameter Effects

- To make the interpretation of the parameter effects even easier,
 - we can use the `allEffects` function from the `effects` package
 - to visualize the parameter effects. See below.

```
plot(allEffects(Model_Binary))
```



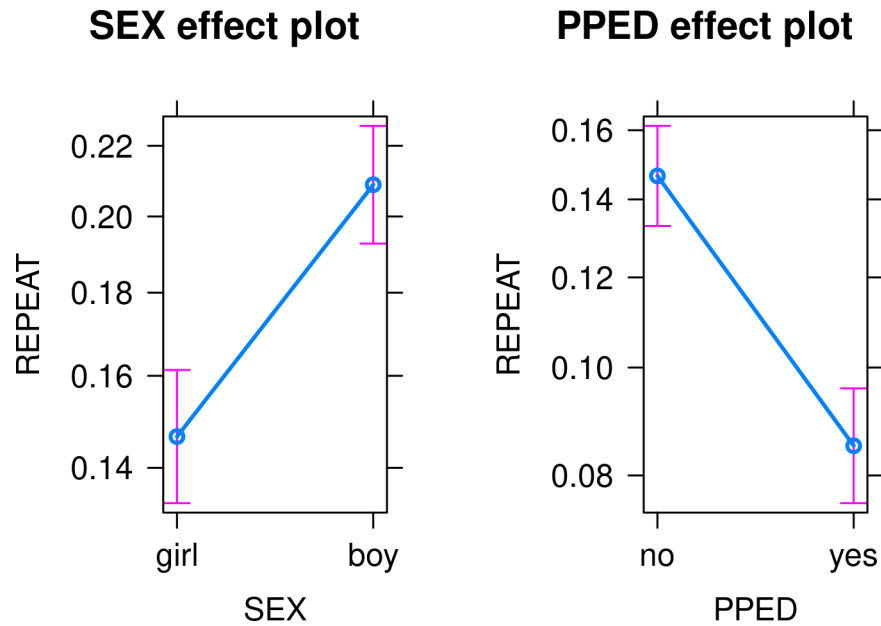
Note that in both plots, the y scale refers to

- the probability of repeating a grade rather than the odds.
 - Probabilities are more interpretable than odds.
 - The probability scores for each variable are calculated
 - by assuming that the other variables in the model are constant
 - and take on their average values.
 - As we can see, assuming that a pupil has an average preschool education,
 - being a boy has a higher probability (~0.16) of repeating a grade
 - than being a girl (~0.11).
 - Likewise, assuming that a pupil has an average gender,
 - having preschool education has a lower probability (~0.11)
 - of repeating a grade than not having preschool education (~0.18).
- Note that in both plots the confidence intervals for the estimates are also included to give us some idea of the uncertainties of the estimates.

Note that the notion of average preschool education and gender

- may sound strange, given they are categorical variables (i.e. factors).
- If you are not comfortable with the idea of assuming an average factor,
 - you can specify your intended factor level as the reference point,
 - by using the `fixed.predictors = list(given.values = ...)` argument
 - in the `allEffects` function. See below:

```
plot(allEffects(Model_Binary, fixed.predictors =
  list(given.values = c(
    SEXboy = 0, PPEDyes = 0
  ))))
```



Setting `SEXboy = 0` means that for the PPED effect plot,

- the reference level of the SEX variable is set to 0;
- `PPEDyes = 0` results in the 0 being the reference level
 - of the PPED variable in the SEX effect plot.

Therefore, as the two plots above show,

- assuming that a pupil has no preschool education,
 - being a boy has a higher probability (~0.20) of repeating a grade
 - than being a girl (~0.14);
- assuming that a pupil is female,
 - having preschool education has a lower probability (~0.09)
 - of repeating a grade than not having preschool education (~0.15).

4.1.2.9.5 Model Evaluation: Goodness of Fit

- There are different ways to evaluate the goodness of fit
 - of a logistic regression model.

Likelihood ratio test

- A logistic regression model has a better fit to the data if the model,
 - compared with a model with fewer predictors,
 - demonstrates an improvement in the fit.

This is performed using the likelihood ratio test,

- which compares the likelihood of the data under the full model
- against the likelihood of the data under a model with fewer predictors.

Removing predictor variables from a model

- will almost always make the model fit less well
 - (i.e. a model will have a lower log likelihood),
- but it is useful to test whether the observed difference in model fit
 - is statistically significant.

```

#specify a model with only the `SEX` variable
Model_Binary_Test <- glm(
  formula = REPEAT ~ SEX,
  family = binomial(link = "logit"),
  data = ThaiEdu_New
)

#use the `anova()`` function to run the likelihood ratio test
anova(Model_Binary_Test, Model_Binary, test = "Chisq")

## Analysis of Deviance Table
##
## Model 1: REPEAT ~ SEX
## Model 2: REPEAT ~ SEX + PPED
##   Resid. Df Resid. Dev Df Deviance  Pr(>Chi)
## 1      7514      6099.1
## 2      7513      6016.2  1    82.941 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

As we can see, the model with both SEX and PPED predictors

- provide a significantly better fit to the data
 - than does the model with only the SEX variable.
- Note that this method can also be used to determine
 - whether it is necessary to include one or a group of variables.

AIC

- The Akaike information criterion (AIC) is another measure for model selection.

Different from the likelihood ratio test,

- the calculation of AIC not only regards the goodness of fit of a model,
 - but also takes into account the simplicity of the model.
- In this way, AIC deals with the trade-off
 - between goodness of fit and complexity of the model,
 - and as a result, discourages overfitting.
- A smaller AIC is preferred.

```
Model_Binary_Test$aic
```

```
## [1] 6103.148
```

```
Model_Binary$aic
```

```
## [1] 6022.207
```

With a smaller AIC value, the model with both SEX and PPED predictors

- is preferred to the one with just the SEX predictor.

Correct Classification Rate

- The percentage of correct classification is another useful measure
 - to see how well the model fits the data.

```

# use the `predict()`` function to calculate the predicted probabilities of pupils in the original data .
Pred <- predict(Model_Binary, type = "response")
Pred <- if_else(Pred > 0.5, 1, 0)

```

```
ConfusionMatrix <-
  table(Pred, pull(ThaiEdu_New, REPEAT)) #`pull` results in a vector
#correct classification rate
sum(diag(ConfusionMatrix)) / sum(ConfusionMatrix)
```

```
## [1] 0.8580362
```

```
ConfusionMatrix
```

```
##
```

```
## Pred    0    1
```

```
##      0 6449 1067
```

We can see that the model correctly classifies 85.8% of all the observations.

However, a closer look reveals that the model

- predicts all of the observations to belong to class “0”,
 - meaning that all pupils are predicted not to repeat a grade.
- Given that the majority category of the REPEAT variable is 0 (No),
 - the model does not perform better in classification
 - than simply assigning all observations to the majority class 0 (No).

4.1.2.9.6 AUC (area under the curve).

- An alternative to using correct classification rate
 - is the Area under the Curve (AUC) measure.

The AUC measures discrimination, that is,

- the ability of the test to correctly classify
 - those with and without the target response.

In the current data, the target response is repeating a grade.

- We randomly pick one pupil from the “repeating a grade” group
 - and one from the “not repeating a grade” group.
- The pupil with the higher predicted probability
 - should be the one from the “repeating a grade” group.
- The AUC is the percentage of randomly drawn pairs for which this is true.
- This procedure sets AUC apart from the correct classification rate
 - because the AUC is not dependent
 - on the imbalance of the proportions of classes in the outcome variable.
- A value of 0.50 means that the model does not classify better than chance.
- A good model should have an AUC score much higher than 0.50
 - (preferably higher than 0.80).

```
# Compute AUC for predicting Class with the model
Prob <- predict(Model_Binary, type = "response")
Pred <- prediction(Prob, as.vector(pull(ThaiEdu_New, REPEAT)))
AUC <- performance(Pred, measure = "auc")
AUC <- AUC@y.values[[1]]
AUC
```

```
## [1] 0.6013622
```

With an AUC score of 0.60,

- the model does not discriminate well.

4.1.2.10 Binomial Logistic Regression

- As mentioned in the beginning,
 - logistic regression can also be used to model count or proportion data.

Binary logistic regression assumes that the outcome variable

- comes from a Bernoulli distribution
 - (which is a special case of binomial distributions)
 - where the number of trial n is 1 and
 - thus the outcome variable can only be 1 or 0.
- In contrast, binomial logistic regression
 - assumes that the number of the target events
 - follows a binomial distribution with n trials and probability q .
- In this way, binomial logistic regression allows the outcome variable
 - to take any non-negative integer value
 - and thus is capable of handling count data.

The Thai Educational Data records information about individual pupils

- that are clustered within schools.

By aggregating the number of pupils who repeated a grade by school,

- we obtain a new data set where each row represents a school,
- with information about the proportion of pupils
 - repeating a grade in that school.
- The MSEC (mean SES score) is also on the school level;
- therefore, it can be used to predict proportion or count of pupils
 - who repeat a grade in a particular school. See below.

```
ThaiEdu_Prop <- ThaiEdu_New %>%
  group_by(SCHOOLID, MSEC) %>%
  summarise(REPEAT = sum(REPEAT),
            TOTAL = n()) %>%
  ungroup()
```

4.1.2.10.1 Transform Data

`summarise()` has grouped output by 'SCHOOLID'. You can override using the
`.groups` argument.

```
head(ThaiEdu_Prop)
```

```
## # A tibble: 6 x 4
##   SCHOOLID MSEC REPEAT TOTAL
##   <fct>    <dbl> <dbl> <int>
## 1 10103     0.88     1     17
## 2 10104     0.2      0     29
## 3 10105    -0.07     5     18
## 4 10106     0.47     0      5
## 5 10108     0.76     3     19
## 6 10109     1.06     9     21
```

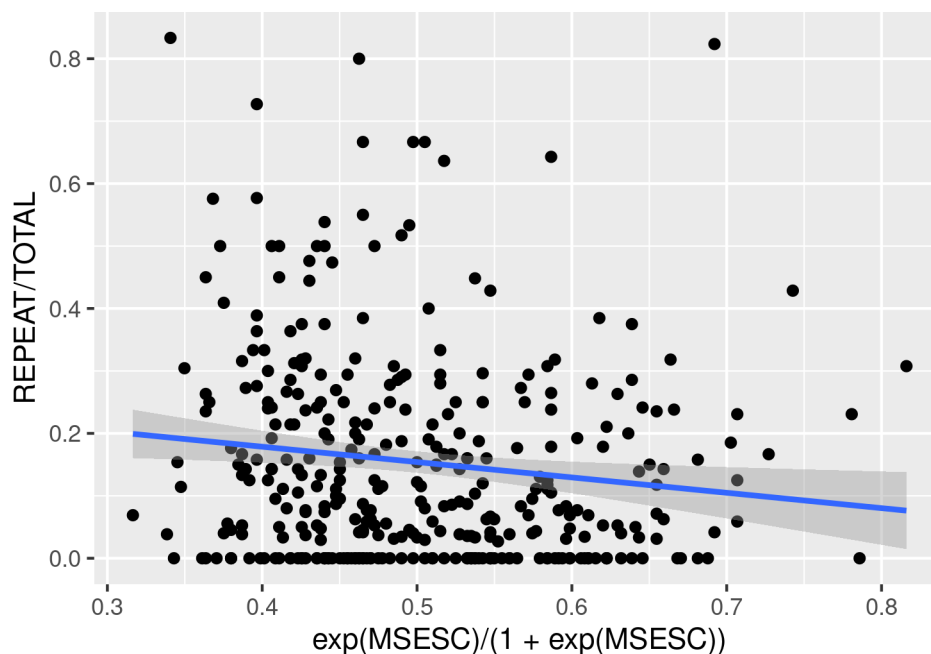
In this new data set,

- REPEAT refers to the number of pupils who repeated a grade;
- TOTAL refers to the total number of students in a particular school.


```
ThaiEdu_Prop %>%
  ggplot(aes(x = exp(MSESC) / (1 + exp(MSESC)), y = REPEAT / TOTAL)) +
  geom_point() +
  geom_smooth(method = "lm")
```

4.1.2.10.2 Explore Data

```
## `geom_smooth()` using formula = 'y ~ x'
```



We can see that the proportion of students who repeated a grade

- is negatively related to the inverse-logit of MSESC.

Note that we model the variable MSESC as its inverse-logit

- because in a binomial regression model,
- we assume a linear relationship between
 - the inverse-logit of the linear predictor
 - and the outcome (i.e. proportion of events),
- not linearity between the predictor itself and the outcome.

4.1.2.10.3 Fit a Binomial Logistic Regression Model

- To fit a binomial logistic regression model, we also use the `glm` function.

The only difference is in the specification

- of the outcome variable in the formula.

We need to specify both

- the number of target events (REPEAT)
- and the number of non-events (TOTAL-REPEAT)
- and wrap them in `cbind()`.

```
Model_Prop <- glm(
  formula = cbind(REPEAT, TOTAL - REPEAT) ~ MSESC,
```

```

family = binomial(logit),
data = ThaiEdu_Prop
)

summary(Model_Prop)

##
## Call:
## glm(formula = cbind(REPEAT, TOTAL - REPEAT) ~ MSEC, family = binomial(logit),
##      data = ThaiEdu_Prop)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.3629  -1.8935  -0.5083   1.1674   6.9494
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.80434    0.03324 -54.280 < 2e-16 ***
## MSEC        -0.43644    0.09164  -4.763 1.91e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1480.7  on 355  degrees of freedom
## Residual deviance: 1457.3  on 354  degrees of freedom
## AIC: 2192
##
## Number of Fisher Scoring iterations: 5

```

4.1.2.10.4 Interpretation

- The parameter interpretation in a binomial regression model
 - is the same as that in a binary logistic regression model.

We know from the model summary above

- that the mean SES score of a school is negatively related to
- the odds of students repeating a grade in that school.

To enhance interpretability, we use the `summ()` function again

- to calculate the exponentiated coefficient estimate of MSEC.

Since MSEC is a continuous variable,

- we can standardize the exponentiated MSEC estimate
- (by multiplying the original estimate with the SD of the variable,
 - and then then exponentiating the resulting number).

```

# Note that to use the summ() function for a binomial regression model, we need to make the outcome var
REPEAT <- pull(filter(ThaiEdu_Prop, !is.na(MSEC)), REPEAT)
TOTAL <- pull(filter(ThaiEdu_Prop, !is.na(MSEC)), TOTAL)
summ(Model_Prop, exp = T, scale = T)

```

We can see that with a SD increase in MSEC,

- the odds of students repeating a grade is lowered by $1 - 85\% = 15\%$.

Observations	356
Dependent variable	cbind(REPEAT, TOTAL - REPEAT)
Type	Generalized linear model
Family	binomial
Link	logit

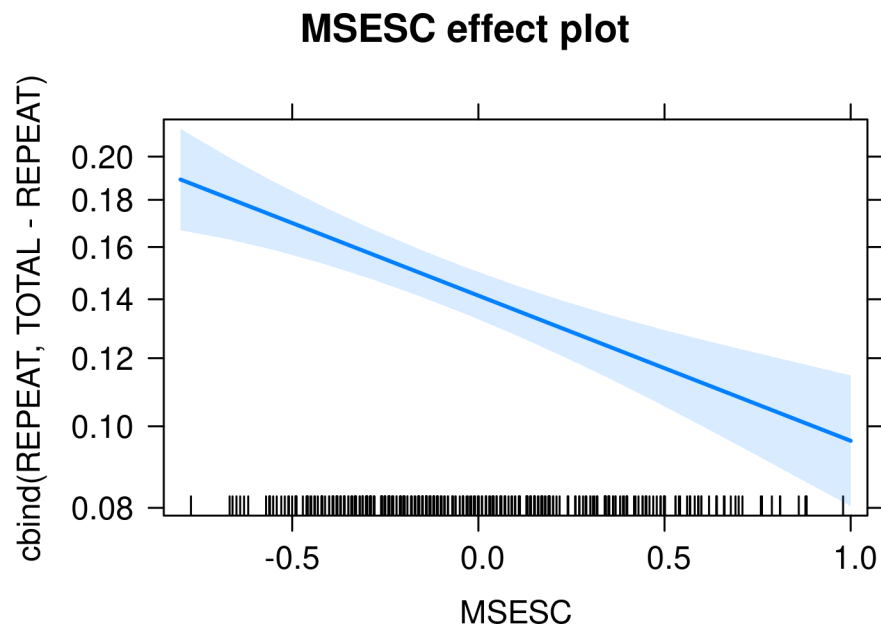
$\chi^2(1)$	23.36
Pseudo-R ² (Cragg-Uhler)	0.06
Pseudo-R ² (McFadden)	0.01
AIC	2191.96
BIC	2199.71

	exp(Est.)	2.5%	97.5%	z val.	p
(Intercept)	0.16	0.15	0.18	-54.26	0.00
MSESC	0.85	0.79	0.91	-4.76	0.00

Standard errors: MLE; Continuous predictors are mean-centered and scaled by 1 s.d.

We can visualize the effect of MSESC.

```
plot(allEffects(Model_Prop))
```



The plot above shows the expected influence of MSESC

- on the probability of a pupil repeating a grade.

Holding everything else constant, as MSESC increases,

- the probability of a pupil repeating a grade lowers
 - (from 0.19 to 0.10).
- The blue shaded areas indicate the 95% confidence intervals
 - of the predicted values at each value of MSESC.

4.1.2.11 Multilevel Binary Logistic Regression

- The binary logistic regression model introduced earlier
 - is limited to modelling the effects of pupil-level predictors;

The binomial logistic regression

- is limited to modelling the effects of school-level predictors.

To incorporate both pupil-level and school-level predictors,

- we can use multilevel models,
- specifically, multilevel binary logistic regression.

If you are unfamiliar with multilevel models,

- you can use [Multilevel analysis: Techniques and Applications](#) for reference
- and this tutorial for a good introduction to multilevel models
 - with the `lme4` package in R.

In addition to the motivation above,

- there are more reasons to use multilevel models.

For instance, as the data are clustered within schools,

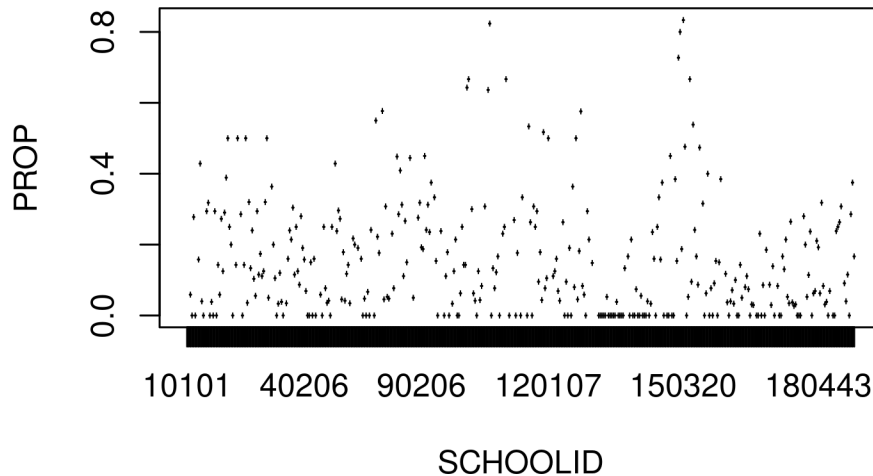
- it is likely that pupils from the same school
 - are more similar to each other than those from other schools.
- Because of this, in one school,
 - the probability of a pupil repeating a grade may be high,
 - while in another school, low.
- Furthermore, even the relationship between the outcome
 - (i.e. repeating a grade)
- and the predictor variables (e.g. gender, preschool education, SES)
 - may be different across schools.
- Also note that there are missing values in the MSES variable.
- Using multilevel models can appropriately address these issues.

See the following plot as an example.

The plot shows the proportions of students repeating a grade across schools.

- We can see vast differences across schools.
- Therefore, we may need multilevel models.

```
ThaiEdu_New %>%
  group_by(SCHOOLID) %>%
  summarise(PROP = sum(REPEAT) / n()) %>%
  plot()
```



4.1.2.11.1 Center Variables

- Prior to fitting a multilevel model,
 - it is necessary to center the predictors
 - * by using an appropriately chosen centering method
 - * (i.e. grand-mean centering or within-cluster centering),
 - because the centering approach matters
 - * for the interpretation of the model estimates.
 - Following the advice of [Enders and Tofghi \(2007\)](#),
 - * we should use within-cluster centering
 - * for the first-level predictors SEX and PPED,
 - and grand-mean centering
 - * for the second-level predictor MDESC.

```
ThaiEdu_Center <- ThaiEdu_New %>%
  mutate(SEX = if_else(SEX == "girl", 0, 1),
         PPED = if_else(PPED == "yes", 1, 0)) %>%
  group_by(SCHOOLID) %>%
  mutate(SEX = SEX - mean(SEX),
         PPED = PPED - mean(PPED)) %>%
  ungroup() %>%
  mutate(MDESC = MDESC - mean(MDESC, na.rm = T))

head(ThaiEdu_Center)
```

```
## # A tibble: 6 x 5
##   SCHOOLID  SEX  PPED REPEAT  MDESC
##   <fct>    <dbl> <dbl> <dbl+lbl> <dbl>
## 1 10103    -0.647 -0.882 0 [no]    0.870
## 2 10103    -0.647 -0.882 0 [no]    0.870
## 3 10103    -0.647  0.118 0 [no]    0.870
## 4 10103    -0.647  0.118 0 [no]    0.870
## 5 10103    -0.647  0.118 0 [no]    0.870
## 6 10103    -0.647  0.118 0 [no]    0.870
```

4.1.2.11.2 Intercept Only Model

- To specify a multilevel model,
 - we use the `glmer` function from the `lme4` package.

Note that the random effect term

- should be included in parentheses.

In addition, within the parentheses,

- the random slope term(s) and
- the cluster terms should be separated by | (the “pipe” symbol).

Note that we use an additional argument

- `control = glmerControl(optimizer = "bobyqa", optCtrl = list(maxfun=2e5))`
- in the `glmer` function
 - to specify a higher number of maximum iterations than default (10000).
- This might be necessary because a multilevel model
 - may require a large number of iterations to converge.

We start by specifying an intercept-only model,

- in order to assess the impact of the clustering structure of the data.

```
Model_Multi_Intercept <- glmer(
  formula = REPEAT ~ 1 + (1 | SCHOOLID),
  family = binomial(link = "logit"),
  data = ThaiEdu_Center,
  control = glmerControl(optimizer = "bobyqa",
                        optCtrl = list(maxfun = 2e5))
)

summary(Model_Multi_Intercept)

## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula: REPEAT ~ 1 + (1 | SCHOOLID)
## Data: ThaiEdu_Center
## Control: glmerControl(optimizer = "bobyqa", optCtrl = list(maxfun = 2e+05))
##
##      AIC      BIC   logLik deviance df.resid
## 5547.1  5560.9 -2771.5  5543.1     7514
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.6254 -0.4174 -0.2487 -0.1765  4.7824
##
## Random effects:
## Groups   Name      Variance Std.Dev.
## SCHOOLID (Intercept) 1.646    1.283
## Number of obs: 7516, groups: SCHOOLID, 356
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -2.22481    0.08391  -26.52   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Below we calculate

- the ICC (intra-class correlation)

- of the intercept-only model.

```
performance::icc(Model_Multi_Intercept)
```

```
## # Intraclass Correlation Coefficient
##
##      Adjusted ICC: 0.333
##      Unadjusted ICC: 0.333
```

An ICC of 0.33 means

- that 33% of the variation in the outcome variable
 - can be accounted for by the clustering structure of the data.
- This provides evidence that a multilevel model
 - may make a difference to the model estimates,
 - in comparison with a non-multilevel model.
- Therefore, the use of multilevel models is necessary and warranted.

4.1.2.11.3 Full Model

- It is good practice to build a multilevel model step by step.

However, as this tutorial's focus is not on multilevel modelling,

- we go directly from the intercept-only model
 - to the full-model that we are ultimately interested in.

In the full model,

- we include not only fixed effect terms of SEX, PPED and MSEC
 - and a random intercept term,
- but also random slope terms for SEX and PPED.

Note that we specify `family = binomial(link = "logit")`,

- as this model is essentially a binary logistic regression model.

```
Model_Multi_Full <-
  glmer(
    REPEAT ~ SEX + PPED + MSEC + (1 + SEX + PPED | SCHOOLID),
    family = binomial(link = "logit"),
    data = ThaiEdu_Center,
    control = glmerControl(optimizer = "bobyqa", optCtrl = list(maxfun = 2e5))
  )
```

```
## boundary (singular) fit: see help('isSingular')
```

```
?isSingular
```

boundary (singular) fit: see ?isSingular

```
summary(Model_Multi_Full)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula: REPEAT ~ SEX + PPED + MSEC + (1 + SEX + PPED | SCHOOLID)
## Data: ThaiEdu_Center
## Control: glmerControl(optimizer = "bobyqa", optCtrl = list(maxfun = 2e+05))
##
##      AIC      BIC    logLik deviance df.resid
```

```
##    5468.0    5537.2   -2724.0    5448.0      7506
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.4509 -0.4003 -0.2451 -0.1732  5.6938
##
## Random effects:
##   Groups   Name                Variance Std.Dev. Corr
## SCHOOLID (Intercept) 1.66585   1.2907
##           SEX         0.15439   0.3929    0.56
##           PPED         0.04748   0.2179   -0.61  0.32
## Number of obs: 7516, groups: SCHOOLID, 356
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -2.2727     0.0886 -25.652  < 2e-16 ***
## SEX           0.4093     0.1105   3.703  0.000213 ***
## PPED          -0.5555     0.1534  -3.621  0.000293 ***
## MSESC         -0.5054     0.2173  -2.326  0.020020 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##      (Intr) SEX    PPED
## SEX    0.017
## PPED   0.024  0.064
## MSESC  0.048  0.064  0.077
## optimizer (bobyqa) convergence code: 0 (OK)
## boundary (singular) fit: see help('isSingular')
```

The results (pertaining to the fixed effects) are similar

- to the results of the previous
 - binary logistic regression and binomial logistic regression models.
- On the pupil-level,
 - SEX has a significant and positive influence
 - on the odds of a pupil repeating a grade,
 - while PPED has a significant and negative influence.
- On the school-level, MSESC has a significant and negative effect
 - on the outcome variable.
- Let's also look at the variance of the random effect terms.

Again, we can use the `summ()` function

- to retrieve the exponentiated coefficient estimates for easier interpretation.

```
summ(Model_Multi_Full, exp = T)
```

Observations	7516
Dependent variable	REPEAT
Type	Mixed effects generalized linear model
Family	binomial
Link	logit

We can also use the `allEffects` function

AIC	5467.99
BIC	5537.23
Pseudo-R ² (fixed effects)	0.02
Pseudo-R ² (total)	0.36

Fixed Effects				
	exp(Est.)	S.E.	z val.	p
(Intercept)	0.10	0.09	-25.65	0.00
SEX	1.51	0.11	3.70	0.00
PPED	0.57	0.15	-3.62	0.00
MSESC	0.60	0.22	-2.33	0.02

Random Effects		
Group	Parameter	Std. Dev.
SCHOOLID	(Intercept)	1.29
SCHOOLID	SEX	0.39
SCHOOLID	PPED	0.22

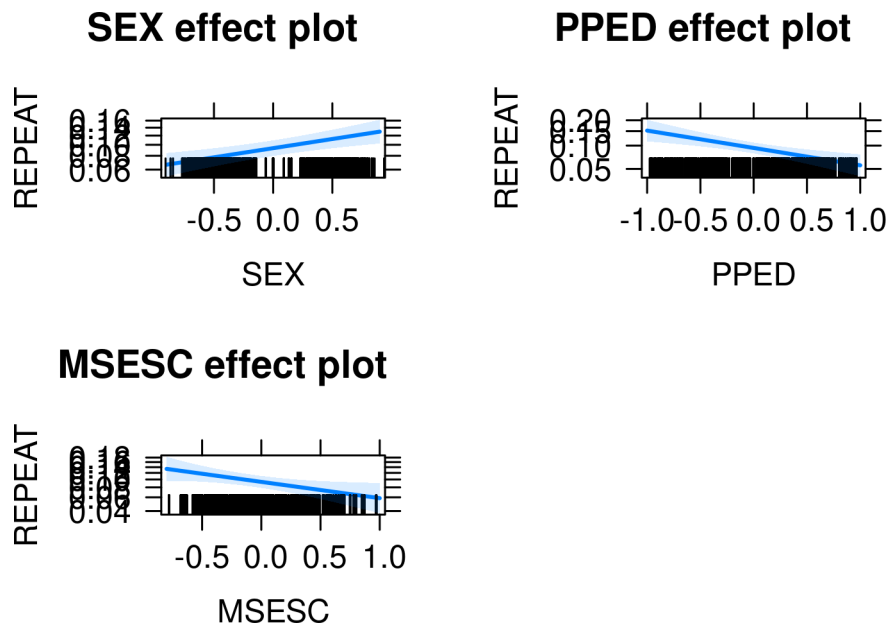
Grouping Variables		
Group	# groups	ICC
SCHOOLID	356	0.34

- to visualize the effects of the parameter estimates.

Note that because the first-level categorical variables

- (SEX and PPED) are centered,
- they are treated as continuous variables in the model
- and as well in the following effect plots.

```
plot(allEffects(Model_Multi_Full))
```



In addition to the fixed-effect terms,

- let's also look at the random effect terms.

From the ICC value before,

- we know that it's necessary to include a random intercept.

However, the necessity of including

- random slopes for SEX and PPED is less clear.

To find this out, we can use

- the likelihood ratio test
 - and AIC
- to judge whether the inclusion of the random slope(s)
 - improves model fit.

```
# let's fit a less-than-full model that leaves out the random slope term of `SEX`
Model_Multi_Full_No_SEX <-
  glmer(
    REPEAT ~ SEX + PPED + MDESC + (1 + PPED | SCHOOLID),
    family = binomial(link = "logit"),
    data = ThaiEdu_Center,
    control = glmerControl(optimizer = "bobyqa",
                           optCtrl = list(maxfun = 2e5))
  )
```

```
## boundary (singular) fit: see help('isSingular')
```

```
boundary (singular) fit: see ?isSingular
```

```
# let's fit a less-than-full model that leaves out the random slope term of `PPED`
Model_Multi_Full_No_PPED <-
  glmer(
    REPEAT ~ SEX + PPED + MDESC + (1 + SEX | SCHOOLID),
    family = binomial(link = "logit"),
    data = ThaiEdu_Center,
```

```

    control = glmerControl(optimizer = "bobyqa", optCtrl =
                          list(maxfun = 2e5))
  )

# let's fit a less-than-full model that leaves out the random slope terms of both `SEX` and `PPED`
Model_Multi_Full_No_Random_Slope <-
  glmer(
    REPEAT ~ SEX + PPED + MSESC +
      (1 | SCHOOLID),
    family = binomial(link = "logit"),
    data = ThaiEdu_Center,
    control = glmerControl(optimizer = "bobyqa",
                          optCtrl = list(maxfun = 2e5))
  )

```

Likelihood ratio test:

```

# compare the full model with that model that excludes `SEX`
anova(Model_Multi_Full_No_SEX, Model_Multi_Full, test = "Chisq")

## Data: ThaiEdu_Center
## Models:
## Model_Multi_Full_No_SEX: REPEAT ~ SEX + PPED + MSESC + (1 + PPED | SCHOOLID)
## Model_Multi_Full: REPEAT ~ SEX + PPED + MSESC + (1 + SEX + PPED | SCHOOLID)
##
##          npar    AIC    BIC logLik deviance Chisq Df
## Model_Multi_Full_No_SEX      7 5466.6 5515.1 -2726.3   5452.6
## Model_Multi_Full           10 5468.0 5537.2 -2724.0   5448.0 4.6054  3
##
##          Pr(>Chisq)
## Model_Multi_Full_No_SEX
## Model_Multi_Full           0.2031

# compare the full model with that model that excludes `PPED`
anova(Model_Multi_Full_No_PPED, Model_Multi_Full, test = "Chisq")

```

```

## Data: ThaiEdu_Center
## Models:
## Model_Multi_Full_No_PPED: REPEAT ~ SEX + PPED + MSESC + (1 + SEX | SCHOOLID)
## Model_Multi_Full: REPEAT ~ SEX + PPED + MSESC + (1 + SEX + PPED | SCHOOLID)
##
##          npar    AIC    BIC logLik deviance Chisq Df
## Model_Multi_Full_No_PPED      7 5462.9 5511.4 -2724.4   5448.9
## Model_Multi_Full           10 5468.0 5537.2 -2724.0   5448.0 0.9052  3
##
##          Pr(>Chisq)
## Model_Multi_Full_No_PPED
## Model_Multi_Full           0.8242

anova(Model_Multi_Full_No_Random_Slope, Model_Multi_Full, test = "Chisq")

## Data: ThaiEdu_Center
## Models:
## Model_Multi_Full_No_Random_Slope: REPEAT ~ SEX + PPED + MSESC + (1 | SCHOOLID)
## Model_Multi_Full: REPEAT ~ SEX + PPED + MSESC + (1 + SEX + PPED | SCHOOLID)
##
##          npar    AIC    BIC logLik deviance Chisq Df
## Model_Multi_Full_No_Random_Slope      5 5463.2 5497.8 -2726.6   5453.2
## Model_Multi_Full           10 5468.0 5537.2 -2724.0   5448.0 5.2249  5
##
##          Pr(>Chisq)
## Model_Multi_Full_No_Random_Slope

```

```
## Model_Multi_Full 0.3891
```

From the all insignificant likelihood ratio test results

- $(\text{Pr}(>\text{Chisq}) > 0.05)$,
- we can conclude that there is no significant improvement in model fit
 - by adding any random slope terms.

AIC:

```
AIC(logLik(Model_Multi_Full)) #full model
```

```
## [1] 5467.985
```

```
AIC(logLik(Model_Multi_Full_No_SEX)) #model without SEX
```

```
## [1] 5466.591
```

```
AIC(logLik(Model_Multi_Full_No_PPED)) #model without PPED
```

```
## [1] 5462.89
```

```
AIC(logLik(Model_Multi_Full_No_Random_Slope)) #model without random slopes
```

```
## [1] 5463.21
```

From the AIC results, we see that

- including random slope terms
 - either does not substantially improve AIC (indicated by lower AIC value)
 - or leads to worse AIC (i.e. higher).
- Therefore, we also conclude
 - there is no need to include the random effect term(s).

4.1.2.12 Other Family (Distribution) and Link Functions

- So far, we have introduced binary and binomial logistic regression,
 - both of which come from the binomial family with the logit link.

However, there are many more distribution families and link functions

- that we can use in `glm` analysis.

For instance, to model binary outcomes,

- we can also use the `probit` link
 - or the complementary log-log (`cloglog`)
- instead of the logit link.

To model count data,

- we can also use Poisson regression,
- which assumes that the outcome variable
 - comes from a Poisson distribution
 - and uses the logarithm as the link function.

For an overview of possible `glm` models,

- see the [Wikipedia page for GLM](#).

4.1.2.13 Links

- [1] Joop Hox, Mirjam Moerbeek, and Rens van de Schoot, [Multilevel Analysis: Techniques and Applications, Third Edition](#).
- Qixiang Fang and Rens van de Schoot, Uni Utrecht, [Intro to Frequentist \(Multilevel\) Generalised Linear Models \(GLM\) in R with glm and lme4](#)

Bates, D., Maechler, M., Bolker, B., & Walker, S. (2015). Fitting Linear Mixed-Effects Models Using lme4. *Journal of Statistical Software*, 67(1), 1-48. doi:[%5B10.18637/jss.v067.i01](#)](<http://dx.doi.org/10.18637/jss.v067.i01>)

Enders, C. K., & Tofighi, D. (2007). Centering predictor variables in cross-sectional multilevel models: A new look at an old issue. *Psychological Methods*, 12(2), 121-138. doi:[10.1037/1082-989X.12.2.121](#)

Fox, J. (2003). Effect Displays in R for Generalised Linear Models. *Journal of Statistical Software*, 8(15), 1-27. <http://www.jstatsoft.org/v08/i15/>

Long, JA. (2019). jtools: Analysis and Presentation of Social Scientific Data. R package version 2.0.1, <https://cran.r-project.org/package=jtools>

Lüdtke, D. (2019). sjstats: Statistical Functions for Regression Models (Version 0.17.5). doi: [10.5281/zenodo.1284472](#)

Raudenbush, S. W., & Bhunirat, C. (1992). The distribution of resources for primary education and its consequences for educational achievement in Thailand. *International Journal of Educational Research*, 17(2), 143-164. doi:[10.1016/0883-0355\(92\)90005-Q](#)

Sing, T., Sander, O., Beerenwinkel, N. & Lengauer, T. (2005). ROCR: visualizing classifier performance in R. *Bioinformatics*, 21(20), pp. 7881. <http://rocr.bioinf.mpi-sb.mpg.de>

Wickham, H. (2017). tidyverse: Easily Install and Load the ‘Tidyverse’. R package version 1.2.1. <https://CRAN.R-project.org/package=tidyverse>