

DSCI353-353m-453: Class 04a GGplot Graphics in Action

Profs: R. H. French, L. S. Bruckman, P. Leu, K. Davis, S. Cirlos

TAs: W. Oltjen, K. Hernandez, M. Li, M. Li, D. Colvin

07 February, 2023

Contents

4.1.3.1	ggplot2 Graphics	1
4.1.3.2	Creating a graph with ggplot	1
4.1.3.2.1	2
4.1.3.2.2	Geoms	2
4.1.3.2.3	Grouping	7
4.1.3.2.4	Scales	8
4.1.3.2.5	Facets	12
4.1.3.2.6	Labels	14
4.1.3.2.7	Themes	15
4.1.3.3	ggplot2 details	17
4.1.3.3.1	Placing the data and mapping options	17
4.1.3.3.2	Graphs as objects	18
4.1.3.3.3	Exporting graphs	20
4.1.3.3.4	Common Mistakes	22
4.1.3.4	ggplot2 Summary	23
4.1.3.5	Links	24

4.1.3.1 **ggplot2** Graphics

- A good book on **ggplot2** is [The Grammar of Graphics](#)

4.1.3.2 Creating a graph with ggplot

- The **ggplot2** package uses a series of functions
 - to build up a graph in layers.
- We'll build a complex graph
 - by starting with a simple graph
 - and adding additional elements, one at a time.
- By default, **ggplot2** graphs
 - appear on a grey background
 - with white reference lines.
- We'll start by setting the default theme
 - to a white background
 - with light grey reference lines.
- This looks better when printed in black and white.

Let's load the **ggplot2** package and set this default theme.

```
library(ggplot2)
theme_set(theme_bw())
```


4.1.3.2.1

- The first function in building a graph is the `ggplot()` function.
- It specifies the
 - data frame containing the data to be plotted
 - the mapping of the variables to visual properties of the graph.

The mappings are placed in an `aes()` function

- (which stands for aesthetics or “something you can see”).

```
library(ggplot2)
library(mosaicData)
ggplot(data = CPS85, mapping = aes(x = exper, y = wage))
```



Why is the graph empty?

- We specified that the `exper` variable
 - should be mapped to the x-axis
- and that the `wage` variable
 - should be mapped to the y-axis,
- but we haven't yet specified what we wanted placed on the graph.

In this case,

- we'll want points to represent each participant.

4.1.3.2.2 Geoms

Geoms are the geometric objects

- (points, lines, bars, and shaded regions)
 - that can be placed on a graph.
- They are added using functions that start with the phrase `geom_`.
- Currently, 37 different `geoms` are available and the list is growing.

Table 4.1 describes the more common `geoms`,

- along with frequently used options for each.

We'll add points using the `geom_point()` function,

Table 4.1 Geom functions

Function	Adds	Options
<code>geom_bar()</code>	Bar chart	color, fill, alpha
<code>geom_boxplot()</code>	Box plot	color, fill, alpha, notch, width
<code>geom_density()</code>	Density plot	color, fill, alpha, linetype
<code>geom_histogram()</code>	Histogram	color, fill, alpha, linetype, binwidth
<code>geom_hline()</code>	Horizontal lines	color, alpha, linetype, size
<code>geom_jitter()</code>	Jittered points	color, size, alpha, shape
<code>geom_line()</code>	Line graph	color, alpha, linetype, size
<code>geom_point()</code>	Scatterplot	color, alpha, shape, size
<code>geom_rug()</code>	Rug plot	color, side
<code>geom_smooth()</code>	Fitted line	method, formula, color, fill, linetype, size
<code>geom_text()</code>	Text annotations	Many; see the help for this function
<code>geom_violin()</code>	Violin plot	color, fill, alpha, linetype
<code>geom_vline()</code>	Vertical lines	color, alpha, linetype, size


Figure 1: fig

- creating a scatterplot.

In ggplot2 graphs, functions are chained together

- using the + sign to build a final plot.

```
library(ggplot2)
library(mosaicData)
ggplot(data = CPS85, mapping = aes(x = exper, y = wage)) + geom_point()
```




It appears that as experience goes up,

- wages go up,
- but the relationship is weak.

The graph also indicates

- that there is an outlier.
- One individual has a wage
 - much higher than the rest.
- We'll delete this case and reproduce the plot.

```
CPS85 <- CPS85[CPS85$wage < 40, ]
ggplot(data = CPS85, mapping = aes(x = exper, y = wage)) +
  geom_point()
```




A number of options can be specified

- in a `geom_` function (see table 4.1).
- Options for `geom_point()`
 - include `color`, `size`, `shape`, and `alpha`.
- These control
 - the point color, size, shape, and transparency, respectively.
- Colors can be specified by name or hexadecimal code.
- Shape and linetype can be specified by
 - the name or number representing the pattern or symbol respectively.
- Point size is specified with positive real numbers starting at zero.
- Large numbers produce larger point sizes.
- Transparency (`alpha`) ranges from
 - 0 (completely transparent)
 - to 1 (completely opaque).
- Adding a degree of transparency
 - can help visualize overlapping points.

We'll also change the gray background to white using theme

```
ggplot(data = CPS85, mapping = aes(x = exper, y = wage)) +
  geom_point(color = "cornflowerblue",
            alpha = .7,
            size = 3)
```



It would be helpful if the graph

- had a line summarizing the trend between experience and wages.


We can add this line with the `geom_smooth()` function.

- Options control the type of line
 - (linear, quadratic, nonparametric),
- the thickness of the line,
- the line's color,
- and the presence or absence of a confidence interval.

Here we request a linear regression (`method = lm`) line.

```
ggplot(data = CPS85, mapping = aes(x = exper, y = wage)) +
  geom_point(color = "cornflowerblue",
            alpha = .7,
            size = 3) +
  geom_smooth(method = "lm")
```

```
## `geom_smooth()` using formula = 'y ~ x'
```



We can see from this line that on average,

- wages appear to increase to a moderate degree with experience.

4.1.3.2.3 Grouping


- In the previous section, we set graph characteristics
 - such as color and transparency to a constant value.
- However, we can also map variables values
 - to the color, shape, size, transparency, line style,
 - and other visual characteristics of geometric objects.
- This allows groups of observations
 - to be superimposed in a single graph
 - (a process called grouping).

Let's add sex to the plot

- and represent it by color, shape, and linetype.

```
ggplot2::ggplot(
  data = CPS85,
  mapping = aes(
    x = exper,
    y = wage,
    color = sex,
    shape = sex,
    linetype = sex
  )
) +
  geom_point(alpha = .7, size = 3) +
  geom_smooth(method = "lm", se = FALSE, size = 1.5)
```

```
## `geom_smooth()` using formula = 'y ~ x'
```



By default, the first group (female) is represented

- by pink filled circles
 - and a solid pink line,
- while the second group (male) is represented
 - by teal filled triangles
 - and a dashed teal line.

Note that the `color = sex`, `shape = sex`, and `linetype = sex`, options

- are placed in the `aes()` function
 - because we are mapping a variable to an aesthetic.
- The `geom_smooth` option (`se = FALSE`)
 - was added to suppresses the confidence intervals,
 - making the graph less busy and easier to read.
- The `size = 1.5` option
 - makes the line a bit thicker.

4.1.3.2.4 Scales As we've seen, the `aes()` function is used

- to map variables to the visual characteristics of a plot.
- Scales specify how each of these mappings occurs.
- For example, `ggplot2` automatically creates plot axes
 - with tick marks, tick mark labels, and axis labels.
- Often they look fine, but occasionally
 - you'll want to take greater control over their appearance.
- Colors that represent groups are chosen automatically,
 - but you may want to select a different set of colors
 - based on your tastes or a publication's requirements.

Scale functions (which start with `scale_`)

- allow you to modify these default scaling.

In the next plot,

- we'll change the x- and y-axis scaling,

Function	Description
<code>scale_x_continuous()</code> , <code>scale_y_continuous()</code>	Scales the x and y axes for quantitative variables. Options include <code>breaks</code> for specifying tick marks, <code>labels</code> for specifying tick mark labels, and <code>limits</code> to control the range of the values displayed.
<code>scale_x_discrete()</code> , <code>scale_y_discrete()</code>	Same as above for axes representing categorical variables.
<code>scale_color_manual()</code>	Specifies the colors used to represent the levels of a categorical variable. The <code>values</code> option specifies the colors. A table of colors can be found at http://research.stowers.org/mcm/efg/R/Color/Chart/ColorChart.pdf


Figure 2: Scales

- and the colors representing males and females.
- The x-axis representing `exper`
 - will range from 0 to 60 by 10,
- and the y-axis representing `wage`
 - will range from 0 to 30 by 5.
- Females will be coded with an off-red color
 - and males will be coded with an off-blue color.

The code

```
ggplot(data = CPS85,
       mapping = aes(x = exper, y = wage,
                     color = sex, shape=sex, linetype=sex)) +
  geom_point(alpha = .7, size = 3) +
  geom_smooth(method = "lm", se = FALSE, size = 1.5) +
  scale_x_continuous(breaks = seq(0, 60, 10)) +
  scale_y_continuous(breaks = seq(0, 30, 5)) +
  scale_color_manual(values = c("indianred3", "cornflowerblue"))


## `geom_smooth()` using formula = 'y ~ x'
```



Code Styling, using Rstudio's Cntrl-Shift-A to reformat the code.

```
ggplot(
  data = CPS85,
  mapping = aes(
    x = exper,
    y = wage,
    color = sex,
    shape = sex,
    linetype = sex
  )
) +
  geom_point(alpha = .7, size = 3) +
  geom_smooth(method = "lm", se = FALSE, size = 1.5) +
  scale_x_continuous(breaks = seq(0, 60, 10)) +
  scale_y_continuous(breaks = seq(0, 30, 5)) +
  scale_color_manual(values = c("indianred3", "cornflowerblue"))

## `geom_smooth()` using formula = 'y ~ x'
```




The numbers on the x- and y-axes are better,

- and the colors are more attractive.
- However,
 - wages are in dollars.
- We can change the labels on the y-axis
 - to represent dollars using the `scales` package.
- The `scales` package provides
 - label formatting for dollars, euros, percents, and more.

```
ggplot(
  data = CPS85,
  mapping = aes(
    x = exper,
    y = wage,
    color = sex,
    shape = sex,
    linetype = sex
  )
) +
  geom_point(alpha = .7, size = 3) +
  geom_smooth(method = "lm", se = FALSE, size = 1.5) +
  scale_x_continuous(breaks = seq(0, 60, 10)) +
  scale_y_continuous(breaks = seq(0, 30, 5),
                     label = scales::dollar) +
  scale_color_manual(values = c("indianred3", "cornflowerblue"))

## `geom_smooth()` using formula = 'y ~ x'
```



We are definitely getting there.

4.1.3.2.5 Facets

- Here is the next question.
- Is the relationship between
 - `experience`, `wages` and `sex`
 - the same for each job sector?
- Let's repeat this graph once
 - for each job sector in order to explore this.

Sometimes relationships are clearer

- if groups appear in side-by-side graphs
 - rather than overlapping in a single graph.
- Facets reproduce a graph
 - for each level of a given variable
 - * (or combination of variables).
- You can create faceted graphs
 - using the `facet_wrap()` and `facet_grid()` functions.

The syntax is given in table 14.3,

- where `var`, `rowvar`, and `colvar` are factors.

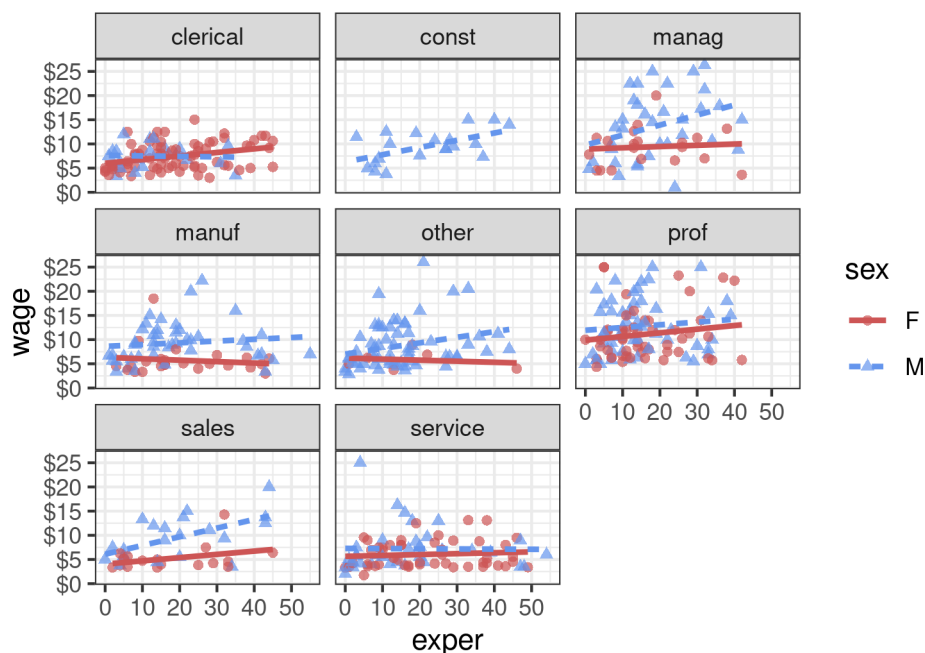
Syntax	Results
<code>facet_wrap(~var, ncol=n)</code>	Separate plots for each level of <code>var</code> arranged into <code>n</code> columns
<code>facet_wrap(~var, nrow=n)</code>	Separate plots for each level of <code>var</code> arranged into <code>n</code> rows
<code>facet_grid(rowvar~colvar)</code>	Separate plots for each combination of <code>rowvar</code> and <code>colvar</code> , where <code>rowvar</code> represents rows and <code>colvar</code> represents columns
<code>facet_grid(rowvar~.)</code>	Separate plots for each level of <code>rowvar</code> , arranged as a single column
<code>facet_grid(.~colvar)</code>	Separate plots for each level of <code>colvar</code> , arranged as a single row

Here, facets will be defined

- by the eight levels of the sector variable.
- Since each facet
 - will be smaller than a one panel graph alone,
- we'll omit `size = 3` from `geom_point()`
 - and `size = 1.5` from `geom_smooth()`.
- This will reduce the point and line sizes
 - compared with the previous graphs
 - and looks better in a faceted graph.

```
ggplot(
  data = CPS85,
  mapping = aes(
    x = exper,
    y = wage,
    color = sex,
    shape = sex,
    linetype = sex
  )
) +
  geom_point(alpha = .7) +
  geom_smooth(method = "lm", se = FALSE) +
  scale_x_continuous(breaks = seq(0, 60, 10)) +
  scale_y_continuous(breaks = seq(0, 30, 5),
                     label = scales::dollar) +
  scale_color_manual(values = c("indianred3", "cornflowerblue")) +
  facet_wrap(~ sector)
```

```
## `geom_smooth()` using formula = 'y ~ x'
```



It appears that the differences between men and women

- depend on the job sector under consideration.
- For example, there is a strong positive relationship
 - between experience and wages for male managers,
 - but not for female managers.
- To a lesser extent,
 - this is also true for sales workers.
- There appears to be no relationship
 - between experience and wages
 - for both male and female service workers.
 - In either case, males make slightly more.
- Wages go up with experience for female clerical workers,
 - but may go down for male clerical workers
 - (the relationship may not be significant here).
- We have gained a great deal of insight into
 - the relationship of wages and experience at this point.

4.1.3.2.6 Labels

- Graphs should be easy to interpret
 - and informative labels are a key element
 - in achieving this goal.
- The `labs()` function provides customized labels
 - for the axes and legends.
- Additionally, a custom title, subtitle, and caption can be added.


```
ggplot(
  data = CPS85,
  mapping = aes(
    x = exper,
    y = wage,
    color = sex,
    shape = sex,
```

```

    linetype = sex
  )
) +
  geom_point(alpha = .7) +
  geom_smooth(method = "lm", se = FALSE) +
  scale_x_continuous(breaks = seq(0, 60, 10)) +
  scale_y_continuous(breaks = seq(0, 30, 5),
                    label = scales::dollar) +
  scale_color_manual(values = c("indianred3",
                                "cornflowerblue")) +
  facet_wrap(~ sector) +
  labs(
    title = "Relationship between wages and experience",
    subtitle = "Current Population Survey",
    caption = "source: http://mosaic-web.org/",
    x = "Years of Experience",
    y = "Hourly Wage",
    color = "Gender",
    shape = "Gender",
    linetype = "Gender"
  )

```

```
## `geom_smooth()` using formula = 'y ~ x'
```



Now a viewer doesn't need to guess

- what the labels `expr` and `wage` mean,
- or where the data come from.


4.1.3.2.7 Themes

- Finally, we can fine tune the appearance of the graph
 - using themes.
- Theme functions (which start with `theme_`)

- control background colors, fonts, grid-lines, legend placement,
- and other non-data related features of the graph.
- Let's use a cleaner theme.
- We used themes at the beginning in order
 - to give each plot a white background.
- Let's try a different theme
 - one that is more minimalistic.

```
ggplot(data = CPS85,
       mapping = aes(x = exper, y = wage, color = sex)) +
  geom_point(alpha = .6) +
  geom_smooth(method = "lm", se = FALSE) +
  scale_x_continuous(breaks = seq(0, 60, 10)) +
  scale_y_continuous(breaks = seq(0, 30, 5),
                    label = scales::dollar) +
  scale_color_manual(values = c("indianred3", "cornflowerblue")) +
  facet_wrap(~ sector) +
  labs(
    title = "Relationship between wages and experience",
    subtitle = "Current Population Survey",
    caption = "source: http://mosaic-web.org/",
    x = "Years of Experience",
    y = "Hourly Wage",
    color = "Gender"
  ) +
  theme_minimal()
```

`geom_smooth()` using formula = 'y ~ x'



This is our finished graph,

- ready for publication.
- Of course, these findings are tentative.
- They are based on a limited sample size and
 - don't involve statistical testing

- to assess whether differences may be due to chance variation.

4.1.3.3 ggplot2 details

- Before we finish, there are three important topics to consider:
 - the placement of the `aes()` function,
 - the treatment of `ggplot2` graphs as R objects,
 - and various methods to save your graphs
 - * for use in reports and webpages.

4.1.3.3.1 Placing the data and mapping options

- Plots created with `ggplot2`
 - always start with the `ggplot` function.
- In the previous examples,
 - the `data=` and `mapping=` options were placed in this function.
- In this case they apply to each geom function that follows.


You can also place these options

- directly within a geom.
- In that case, they only apply to that specific geom.

Consider the following graph.

```
ggplot(CPS85,
       mapping = aes(x = exper, y = wage, color = sex)) +
  geom_point(alpha = .7, size = 3) +
  geom_smooth(method = "lm", se = FALSE, size = 1.5)
```

`geom_smooth()` using formula = 'y ~ x'



Since the mapping of sex to color


- appears in the `ggplot()` function,
 - it applies to both `geom_point` and `geom_smooth`.
- The color of the point indicates the sex,

- and a separate colored trend line is produced for men and women.

Compare this to

```
ggplot(CPS85, aes(x = exper, y = wage)) +
  geom_point(aes(color = sex), alpha = .7, size = 3) +
  geom_smooth(method = "lm", se = FALSE, size = 1.5)
```

```
## `geom_smooth()` using formula = 'y ~ x'
```



Since the sex to color mapping

- only appears in the `geom_point()` function,
 - it is only used there.
- A single trend line is created for all observations.

Most examples place the data and mapping options

- in the `ggplot` function.
- Additionally, the phrases `data =` and `mapping =` are omitted
 - since the first option always refers to data
 - and the second option always refers to mapping.

But your code will be more readable and understandable

- If you put the `data =` and `mapping =` in your code explicitly.


4.1.3.3.2 Graphs as objects

- A `ggplot2` graph can be saved as
 - a named R object (a list),
- or manipulated further,
 - and then printed or saved to disk.


```
data(CPS85, package = "mosaicData")      #A
CPS85 <- CPS85[CPS85$wage < 40,]          #A

myplot <- ggplot(data = CPS85, aes(x = exper, y = wage)) + geom_point()      #B
```

```
myplot #C
```




```
myplot2 <- myplot + geom_point(size = 3, color = "blue") #D  
myplot2 #D
```



```
myplot + geom_smooth(method = "lm") + #E  
  labs(title = "Mildly interesting graph") #E
```

```
## `geom_smooth()` using formula = 'y ~ x'
```

Mildly interesting graph



And here is what the codes doing at each letter step

- #A Prepare data
- #B Create a scatterplot and save it as myplot
- #C Display myplot
- #D Make the points larger and blue,
 - save it as myplot2 and display the graph
- #E Display myplot with a best fit line and a title

First the data are imported and outliers are removed #A.

- Then a simple scatter plot of experience vs. wages
 - is created and saved as myplot #B.
- Next, the plot is printed #C.
- The plot is then modified
 - by changing the point size and color,
 - saved as myplot2 and printed #D.
- Finally, the original plot
 - is given a line of best fit and title, and printed #E.

Note that these changes are not saved.

The ability to save graphs as objects

- allows you to continue to work with and modify them.
- This can be a real time saver (and help you avoid carpal tunnel syndrome).
- It is also handy when saving graphs programmatically.

4.1.3.3.3 Exporting graphs

- You can export graphs created by `ggplot2`
 - in a variety of image formats
 - using the RStudio GUI
 - or through your code.
- To export a graph using the RStudio menus,
 - go to the Plots tab and choose Export

To export a graph via code

- use the `ggsave()` function.
- You can specify the plot to save,
 - its size and format,
 - and where to save it. For example,
 - `ggsave(file = "mygraph.png", plot = myplot, width = 5, height = 4)`
 - saves myplot as a 5-inch by 4-inch PNG file
 - named mygraph.png
 - in the current working directory.

You can save the graph in a different format

- by changing the file extension.

Extension	Format
pdf	Portable Document Format
jpeg	JPEG
tiff	Tagged Image File Format
png	Portable Network Graphics
svg	Scalable Vector Graphics
wmf	Windows Metafile

Figure 3: Image File Formats

The pdf, svg, and wmf formats are lossless

- they resize without fuzziness or pixilation.
- The other formats are lossy
 - they will pixelate when resized.
 - This is especially noticeable when small images are enlarged.
- The png format is popular for images destined for webpages.
- The jpeg and tif formats are usually reserved for photographs.

The wmf format is usually recommended for graphs

- that will appear in Microsoft Word or PowerPoint documents.
- MS Office does not support pdf or svg files,
 - and the wmf format will rescale well.
- However, note that wmf files
 - will lose any transparency settings that have been set.

If you omit the `plot =` option,

- the most recently created graph is saved.


The code is valid and

- saves the graph to disk as a PDF document.

- See `help(ggsave)` for additional details.

```
ggplot(data = mtcars, aes(x = mpg)) + geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
ggsave(file = "mygraph.pdf")
```

```
## Saving 5 x 3.5 in image
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

4.1.3.3.4 Common Mistakes After working with ggplot2 for years,

- I've found that there are two mistakes that are frequently made.
- The first is omitting or misplacing a closing parentheses.
 - This happens most often following the `aes()` function.

Consider the following code.

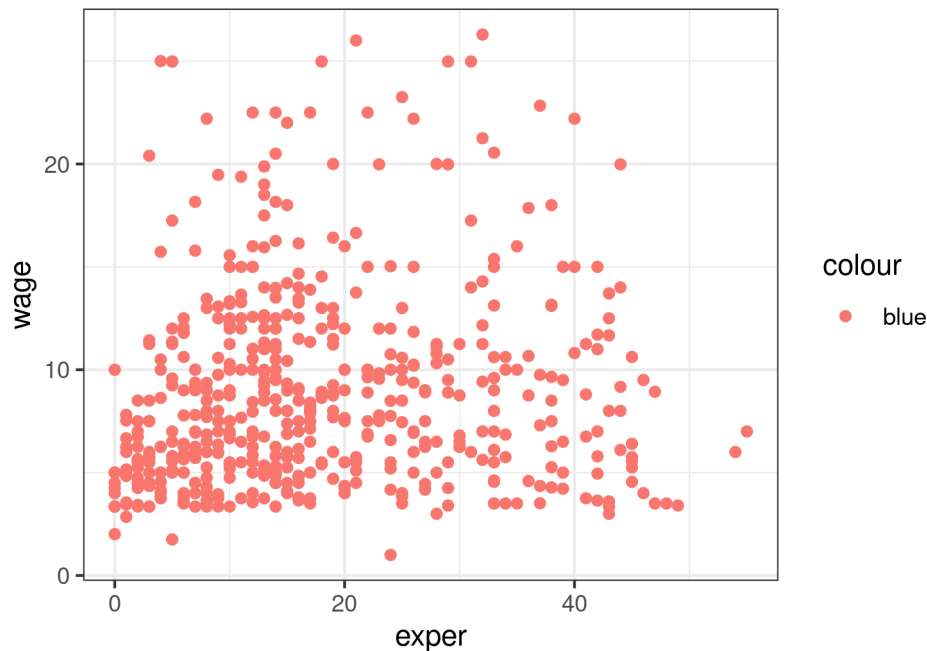
```
ggplot(CPS85, aes(x = exper, y = wage, color = sex) +  
  geom_point()
```

- Note the lack of a closing parentheses at the end of the first line.
 - I can't tell you how many times I've done this.

The second error is confusing an assignment for a mapping.

This code produces the next graph.

```
ggplot(CPS85, aes(x = exper, y = wage, color = "blue")) +  
  geom_point()
```



The points are red (not blue) and there is a strange legend.

- What happened?

The `aes()` function is used

- to map variables to the visual characteristics of the graph.
- Assigning constant values is done outside the `aes()` function.

4.1.3.4 ggplot2 Summary

- The `ggplot2` package provides a powerful platform
 - for creating both simple and complex graphs.
 - Graphs are built up in layers
 - * using functions chained together with the plus (+) symbol.
- The `ggplot()` function specifies a data frame
 - containing plot data
 - and an `aes()` function that maps variables
 - * to visual aspects of the graph.
 - `geom_` functions specify the geometric objects
 - * (bars, lines, points, etc.)
 - * to be placed on the graph.
- Optional `scale_` functions allow you to
 - customize how a variable's values
 - * will be translated into their visual representations on the graph
 - * (e.g., the x- and y-axis scales and labels to used,
 - * and what colors, shapes, and line-types
 - * will be mapped to a variable's values).
- Data from two or more groups
 - can be represented by grouping
 - (superimposing plots distinguished by visual aspects such as color)
 - or faceting (placing several small plots in a matrix-like array).
- Two common errors of `ggplot` are
 - missing/misplaced parentheses
 - nd confusing an assignment for a mapping.

- Graphs can be exported in a wide variety of image formats
 - (such as tiff, pdf, jpg, png, svg, and wmf)
 - using the RStudio GUI
 - or with the `ggsave()` function.

4.1.3.5 Links

- Robert I. Kabacoff, R in Action, 3rd Edition, Manning Publications 2020