

DSCI353-353m-453: Class w06a-p2-ggplot2-graphs-2

Profs: R. H. French, L. S. Bruckman, P. Leu, K. Davis, S. Cirlos

TAs: W. Oltjen, K. Hernandez, M. Li, M. Li, D. Colvin

21 February, 2023

Contents

6.1.2.1	ggplot2 Graphics	1
6.1.2.2	Bar Charts	2
6.1.2.2.1	Stacked, grouped and filled bar charts	3
6.1.2.2.2	Mean bar charts	6
6.1.2.2.3	Tweaking Bar Charts	8
6.1.2.2.4	Bar Chart Labels	9
6.1.2.3	Pie Charts (Something I don't use)	11
6.1.2.4	Tree Maps	11
6.1.2.5	Histograms	13
6.1.2.6	Kernel density plots	16
6.1.2.7	Box Plots	20
6.1.2.7.1	Using parallel box plots to compare groups	23
6.1.2.7.2	Violin Plots	25
6.1.2.8	Dot Plots	26
6.1.2.9	Summary	29
6.1.2.10	Links	29

6.1.2.1 ggplot2 Graphics

- Whenever we analyze data, the first thing we should do is look at it.
 - For each variable,
 - * what are the most common values?
 - * How much variability is present?
 - * Are there any unusual observations?
 - R provides a wealth of functions for visualizing data. Lets look at graphs that help is understand
 - * a single categorical or continuous variable.

This topic includes

- Visualizing the distribution of a variable
- Comparing the distribution of a variable across two or more groups

In both cases, the variable

- can be continuous (for example, car mileage as miles per gallon)
 - or categorical (for example, treatment outcome as none, some, or marked).
- Later we'll explore graphs that display
 - more complex relationships among variables.

Here we'll explore

- bar charts,
- pie charts,
- tree maps,
- histograms,
- kernel density plots,
- box plots,
- violin plots,
- and dot plots.

Some of these may be familiar to you,

- whereas others (such as tree charts or violin plots) may be new to you.

The goal, as always, is

- to understand your data better
- and to communicate this understanding to others.

6.1.2.2 Bar Charts

- In these examples, you'll plot the outcome of a study
 - investigating a new treatment for rheumatoid arthritis.
- The data are contained in the Arthritis data frame
 - distributed with the `vcd` package.
- Note that the `vcd` package isn't needed to create bar charts.
 - You're installing it to gain access to the Arthritis dataset.

In the Arthritis study, the variable `Improved`

- records the patient outcomes for individuals receiving a placebo or drug:

```
data(Arthritis, package = "vcd")
table(Arthritis$Improved)
```

```
##
##   None   Some Marked
##    42    14    28
```


Here, you see that

- 28 patients showed marked improvement,
- 14 showed some improvement,
- and 42 showed no improvement.

You can graph these counts using a vertical or horizontal bar chart.


```
library(ggplot2)
data(Arthritis, package = "vcd")
ggplot(Arthritis, aes(x = Improved)) + geom_bar() +
  labs(title = "Simple Bar chart",
       x = "Improvement",
       y = "Frequency")
```

Simple Bar chart



```
ggplot(Arthritis, aes(x = Improved)) + geom_bar() +
  labs(title = "Horizontal Bar chart",
        x = "Improvement",
        y = "Frequency") +
  coord_flip()
```

Horizontal Bar chart



6.1.2.2.1 Stacked, grouped and filled bar charts

- The central question in the Arthritis study is
 - “How does the level of improvement
 - * vary between the placebo and treated conditions?”.

The `table()` function can be used

- to generate a cross-tabulation of the variables.

```
table(Arthritis$Improved, Arthritis$Treatment)
```

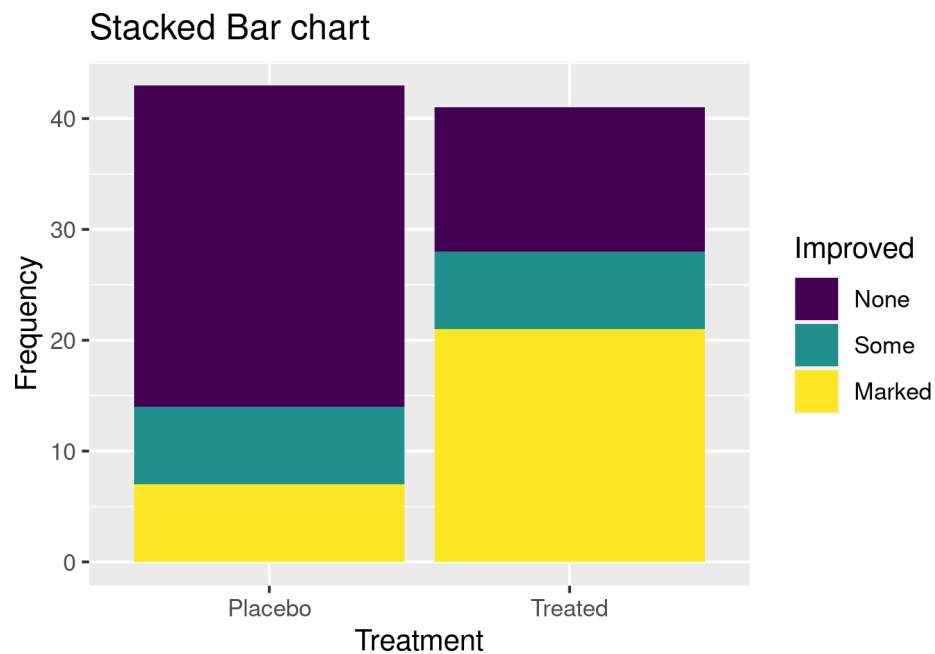
```
##
##           Placebo Treated
## None           29      13
## Some            7       7
## Marked          7      21
```

While the tabulation is helpful,

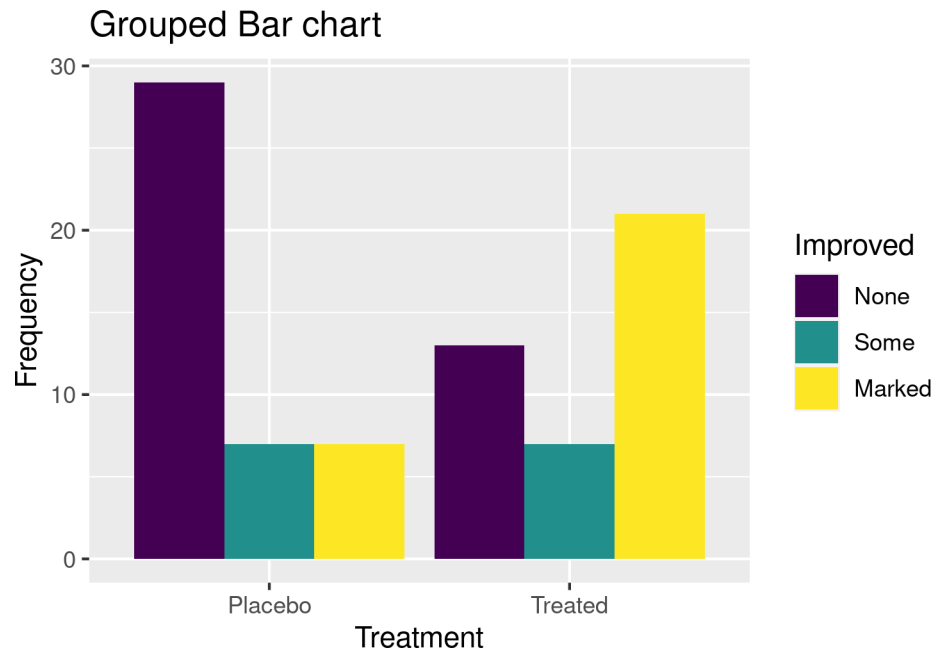
- the results are easier to grasp with a bar chart.
- The relationship between two categorical variables can be plotted
 - using stacked, grouped, or filled bar charts.

#A Stacked bar chart #B Grouped bar chart #C Filled bar chart

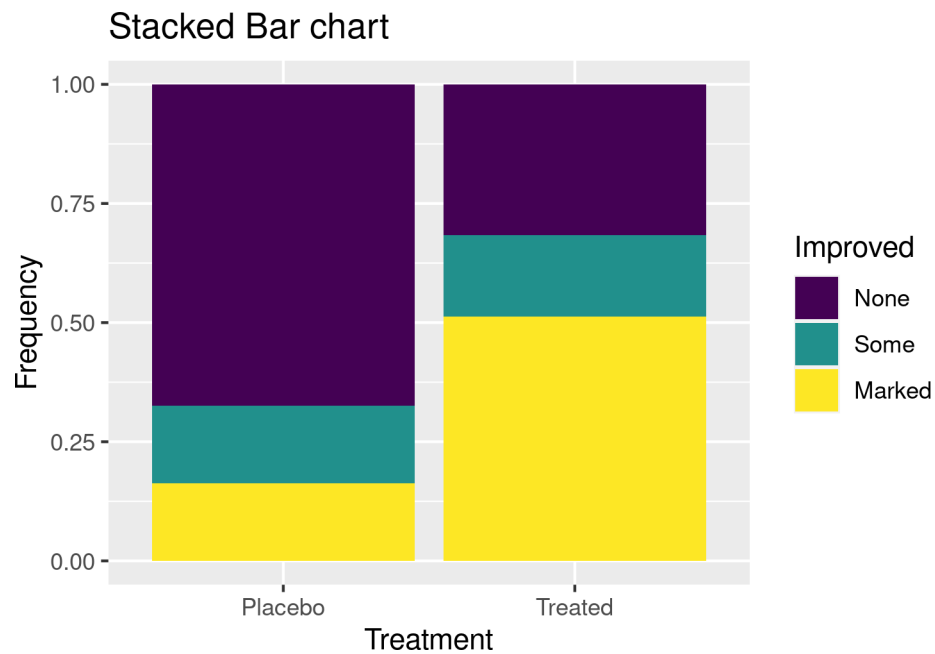
```
ggplot(Arthritis, aes(x = Treatment, fill = Improved)) +      #A
  geom_bar(position = "stack") +
  labs(title = "Stacked Bar chart",
        x = "Treatment",
        y = "Frequency")
```



```
ggplot(Arthritis, aes(x = Treatment, fill = Improved)) +      #B
  geom_bar(position = "dodge") +
  labs(title = "Grouped Bar chart",
        x = "Treatment",
        y = "Frequency")
```



```
ggplot(Arthritis, aes(x = Treatment, fill = Improved)) + #C
  geom_bar(position = "fill") +
  labs(title = "Stacked Bar chart",
       x = "Treatment",
       y = "Frequency")
```



In the stacked bar chart,

- each segment represents the frequency or proportion of cases
 - within in a given Treatment (Placebo, Treated)
 - and Improvement (None, Some, Marked) level combination.
- The segments are stacked separately for each Treatment level.
- The grouped bar chart

- places the segments representing Improvement
 - * side by side within each Treatment level.
- The filled bar chart is a stacked bar chart
 - rescaled so that the height of each bar is 1
 - and the segment heights represent proportions.
- Filled bar charts are particularly useful
 - for comparing the proportions of one categorical variable
 - over the levels of another categorical variable.
 - For example, the filled bar chart
 - * clearly displays the larger percentage of treated patients
 - * with marked improvement
 - * compared with patients receiving a placebo.

6.1.2.2.2 Mean bar charts

- Bar plots needn't be based on counts or frequencies.
 - You can create bar charts that represent
 - * means, medians, percents, standard deviations, and so forth
 - * by summarizing the data with an appropriate statistic
 - * and passing the results to ggplot2.

In the following graph, we'll plot

- the mean illiteracy rate for regions of the United States in 1970.
- The built-in R dataset `state.x77` has the illiteracy rates by state,
 - and the dataset `state.region` has the region names for each state.

#1 Generate means by region #2 Plot means in a sorted bar chart

```
states <- data.frame(state.region, state.x77)
```

```
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
## filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
## intersect, setdiff, setequal, union
```

```
plotdata <- states %>% #1
```

```
  group_by(state.region) %>%
```

```
  summarize(mean = mean(Illiteracy))
```

```
plotdata
```

```
## # A tibble: 4 x 2
```

```
##   state.region mean
```

```
##   <fct>      <dbl>
```

```
## 1 Northeast      1
```

```
## 2 South          1.74
```

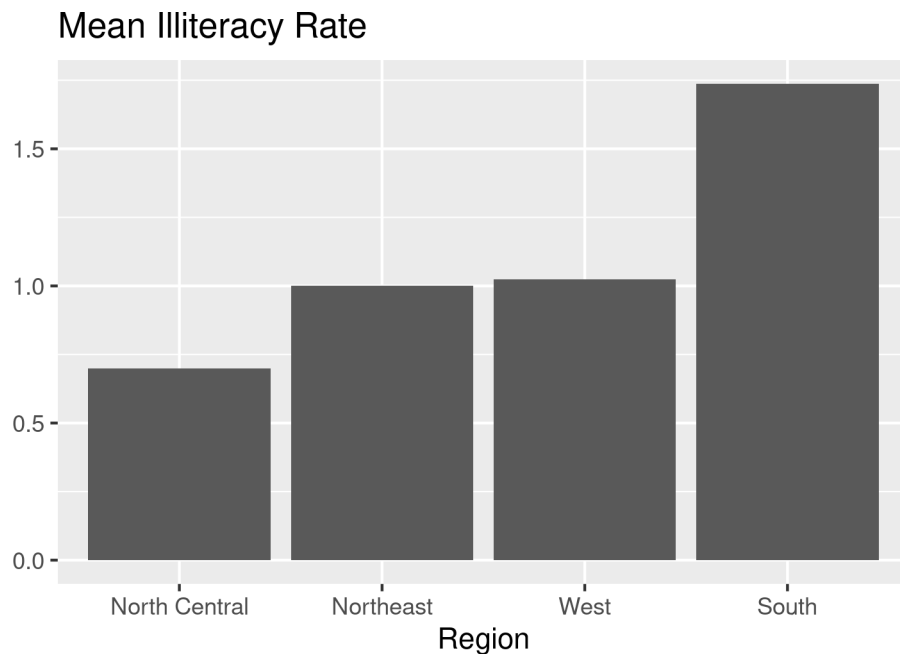
```
## 3 North Central  0.7
```

```
## 4 West           1.02
```

```
ggplot(plotdata, aes(x = reorder(state.region, mean), y = mean)) + #2
```

```
  geom_bar(stat = "identity") +
```

```
labs(x = "Region",
     y = "",
     title = "Mean Illiteracy Rate")
```



When plotting summary statistics such as means,

- it's good practice to indicate the variability of the estimates involved.
- One measure of variability
 - is the standard error of the statistic
- an estimate of the expected variation of the statistic
 - across hypothetical repeated samples.

#1 Generate means and standard errors by region #2 Plot means in a sorted bar chart #3 Add error bars


```
plotdata <- states %>%      #1
  group_by(state.region) %>%
  summarize(
    n = n(),
    mean = mean(Illiteracy),
    se = sd(Illiteracy) / sqrt(n)
  )
```

plotdata

```
## # A tibble: 4 x 4
##   state.region      n mean    se
##   <fct>          <int> <dbl> <dbl>
## 1 Northeast         9  1  0.0928
## 2 South            16  1.74 0.138
## 3 North Central    12  0.7  0.0408
## 4 West             13  1.02 0.169
```

```
ggplot(plotdata, aes(x = reorder(state.region, mean), y = mean)) +      #2
  geom_bar(stat = "identity", fill = "skyblue") +
  geom_errorbar(aes(ymin = mean - se, ymax = mean + se), width = 0.2) +  #3
```

```
labs(
  x = "Region",
  y = "",
  title = "Mean Illiteracy Rate",
  subtitle = "with standard error bars"
)
```



The means and standard errors are calculated for each region #1.

- The bars are then plotted in order of increasing illiteracy.
- The color is changed from a default dark grey
 - to a lighter shade (sky blue)
 - so that error bars to be added in the next step will stand out #2.
- Finally, the error bars are plotted #3.
 - The width option in the `geom_errorbar()` function
 - controls the horizontal width of the error bars and is purely aesthetic
 - * it has no statistical meaning.

In addition to displaying the mean illiteracy rates,

- we can see that the mean for the North Central region
 - is the most reliable (least variability)
- and the West region is least reliable
 - (largest variability).

6.1.2.2.3 Tweaking Bar Charts

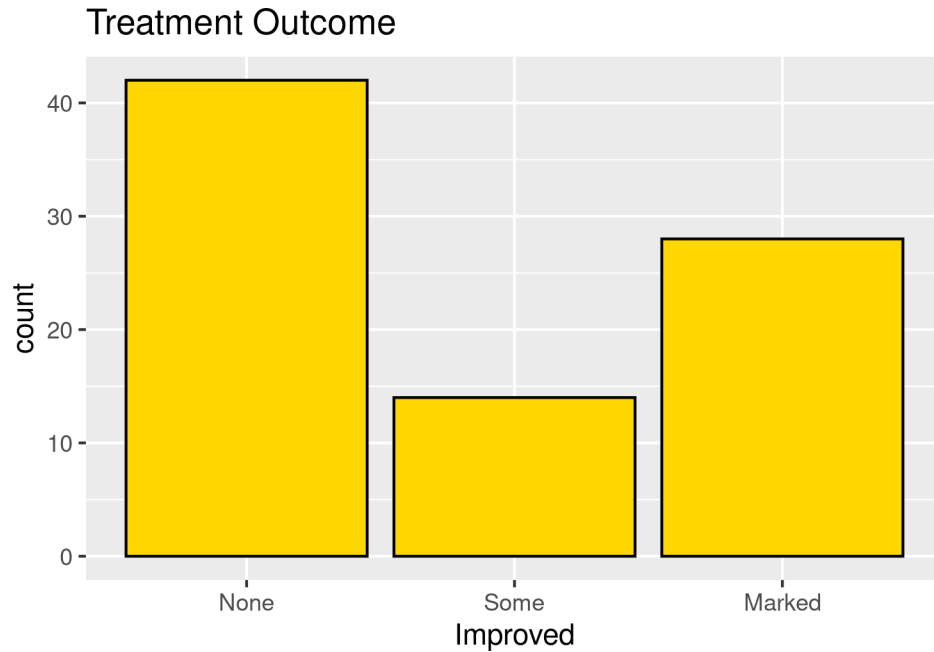
- There are several ways to tweak the appearance of a bar chart.
 - The most common are customizing the bar colors and labels.

Fill vs. Color

- In general, ggplot2 uses **fill**
 - to specify the color of geometric objects that have area
 - (such as bars, pie slices, boxes),
- and the term **color** when referring to

- the color of geometric objects without area
- (such as lines, points, and borders).

```
data(Arthritis, package = "vcd")
ggplot(Arthritis, aes(x = Improved)) +
  geom_bar(fill = "gold", color = "black") +
  labs(title = "Treatment Outcome")
```



Here, bar fill colors are mapped to the levels of the variable Improved.


- The `scale_fill_manual()` function specifies
 - red for None,
 - grey for Some,
 - and gold for Marked improvement.

6.1.2.2.4 Bar Chart Labels

- When there are many bars or long labels,
 - bar chart labels tend to overlap and become unreadable.
 - Consider the following example.
 - The dataset `mpg` in the `ggplot2` package
 - * describes fuel economy data from for 38 popular car models
 - in 1999 and 2008.
 - * Each model has several configurations
 - (transmission type, number of cylinders, etc.).

```
ggplot(mpg, aes(x = model)) +
  geom_bar() +
  labs(title = "Car models in the mpg dataset",
       y = "Frequency", x = "")
```

Car models in the mpg dataset




Two simple tweaks will make the labels readable.

- First, we can plot the data as a horizontal bar chart.
- Second, we can angle the label text and use a smaller font.

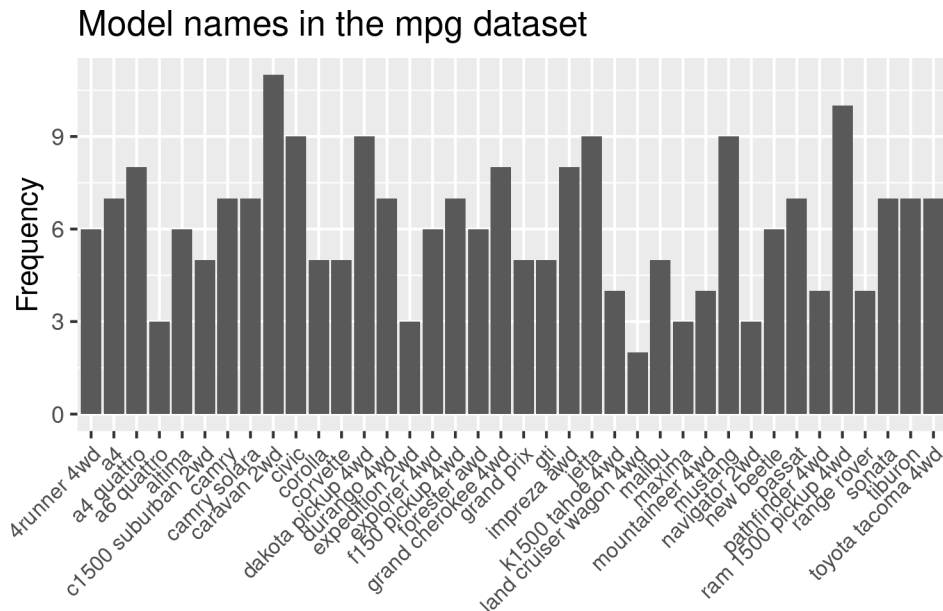
```
ggplot(mpg, aes(x = model)) +
  geom_bar() +
  labs(title = "Car models in the mpg dataset",
        y = "Frequency", x = "") +
  coord_flip()
```

Car models in the mpg dataset



```
ggplot(mpg, aes(x = model)) +  
  geom_bar() +  
  labs(title = "Model names in the mpg dataset",  
        y = "Frequency", x = "") +
```

```
theme(axis.text.x = element_text(
  angle = 45,
  hjust = 1,
  size = 8
))
```



6.1.2.3 Pie Charts (Something I don't use)

- Pie charts are ubiquitous in the business world,
 - but they're denigrated by most statisticians,
 - * including the authors of the R documentation.
 - They recommend bar or dot plots over pie charts
 - * because people are able to judge length more accurately than volume.
 - Perhaps for this reason, the pie chart options in R
 - * are severely limited
 - * when compared with other statistical platforms.

The issue with Pie Charts

6.1.2.4 Tree Maps

- An alternative to a pie chart is a tree map.
 - A tree map displays the distribution of a categorical variable
 - * using rectangles that are proportional to variable levels.
 - Unlike pie charts, tree maps
 - * can handle categorical variables with many levels.
 - We'll use the **treemapify** package

We'll start by creating a tree map

- displaying the distribution of car manufacturers in the mpg data frame.

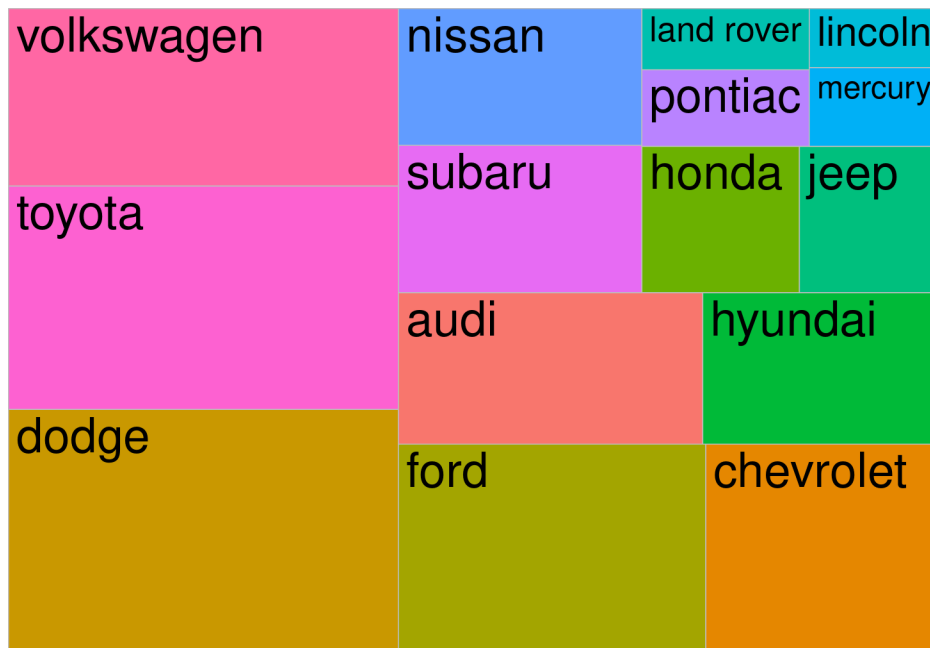
#1 Summarize the data #2 Create the tree map

```
library(ggplot2)
library(dplyr)
```

```
library(treemapify)

plotdata <- mpg %>% count(manufacturer)      #1

ggplot(plotdata,      #2
       aes(fill = manufacturer,
           area = n,
           label = manufacturer)) +
  geom_treemap() +
  geom_treemap_text() +
  theme(legend.position = "none")
```



First we calculate the frequency counts

- for each level the manufacturer variable #1.
- This information is passed to ggplot2
 - to create the graph #2.
- In the `aes()` function,
 - fill refers to the categorical variable,
 - area is the count for level,
 - and label is the option variable used to label the cells.
- The `geom_treemap()` function
 - creates the tree map
 - and the `geom_tree_text()` function adds the labels to each cell.
- The `theme()` function is used to suppress the legend,
 - which is redundant here, since each cell is labeled.

In the next example, a second variable is added drivetrain.

- The number of cars by manufacturer is plotted
 - for front-wheel, rear-wheel, and four-wheel drives.

#1 Compute cell counts #2 Provide better labels for drivetrains #3 Create tree map

```

plotdata <- mpg %>%      #1
  count(manufacturer, drv)
plotdata$drv <- factor(
  #2
  plotdata$drv,
  #3
  levels = c("4", "f", "r"),
  labels = c("4-wheel", "front-wheel", "rear")
)

ggplot(plotdata,
  aes(
    fill = manufacturer,
    area = n,
    label = manufacturer,
    subgroup = drv
  )) +
  geom_treemap() +
  geom_treemap_subgroup_border() +
  geom_treemap_subgroup_text(
    place = "middle",
    colour = "black",
    alpha = 0.5,
    grow = FALSE
  ) +
  geom_treemap_text(colour = "white",
    place = "centre",
    grow = FALSE) +
  theme(legend.position = "none")

```



6.1.2.5 Histograms

- Histograms display the distribution of a continuous variable

- by dividing the range of scores
- into a specified number of bins on the x-axis
- and displaying the frequency of scores in each bin on the y-axis.

```
ggplot(data, aes(x = contvar)) + geom_histogram()
```

You can create histograms using `geom_histogram`

- where data is a data frame
- and contvar is a continuous variable.

Using the mpg data set in the ggplot package,

- we'll examine the distribution of city miles per gallon (cty)
 - for 117 automobile configurations in 2008.

#1 Simple histogram #2 Colored histogram with 20 bins #3 Histogram with percentages #4 Histogram with density curve


```
library(ggplot2)
library(scales)

data(mpg)
cars2008 <- mpg[mpg$year == 2008, ]

ggplot(cars2008, aes(x = hwy)) +      #1
  geom_histogram() +
  labs(title = "Default histogram")
```

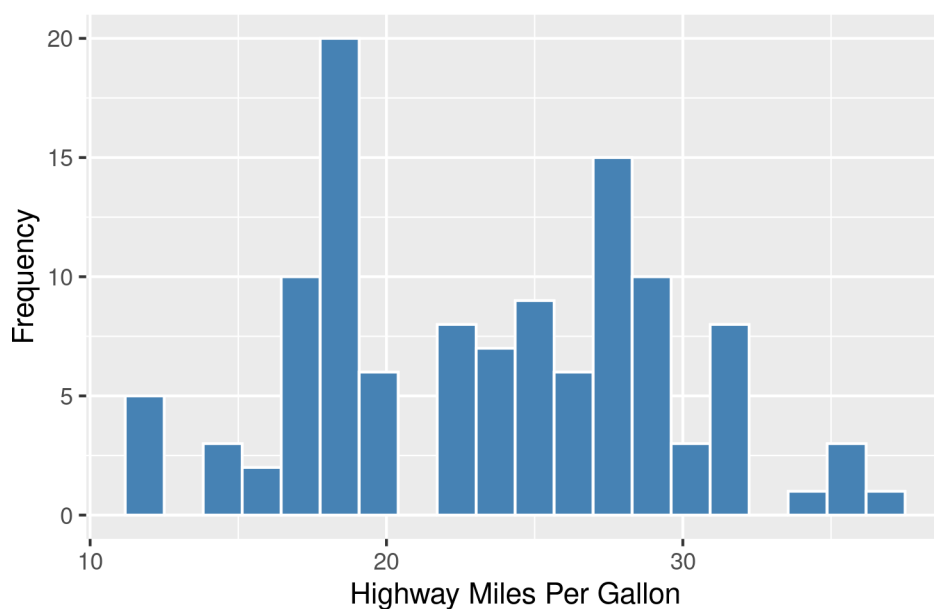
`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

Default histogram



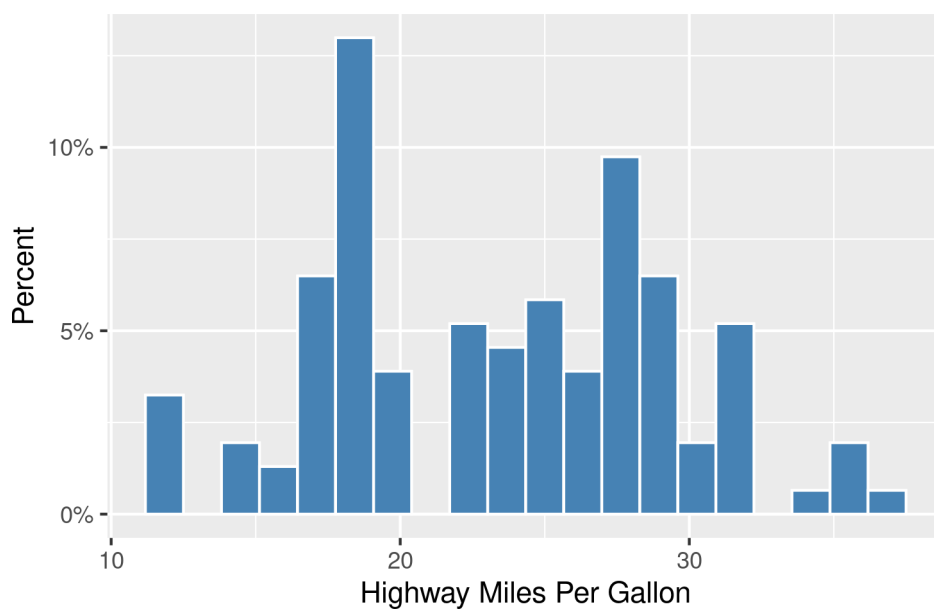
```
ggplot(cars2008, aes(x = hwy)) +
  geom_histogram(bins = 20, color = "white", fill = "steelblue") +      #2
  labs(title = "Colored histogram with 20 bins",
        x = "Highway Miles Per Gallon",
        y = "Frequency")
```

Colored histogram with 20 bins



```
ggplot(cars2008, aes(x = hwy, y = ..density..)) +  
  geom_histogram(bins = 20, color = "white", fill = "steelblue") + #3  
  scale_y_continuous(labels = scales::percent) +  
  labs(title = "Histogram with percentages",  
       y = "Percent",  
       x = "Highway Miles Per Gallon")
```

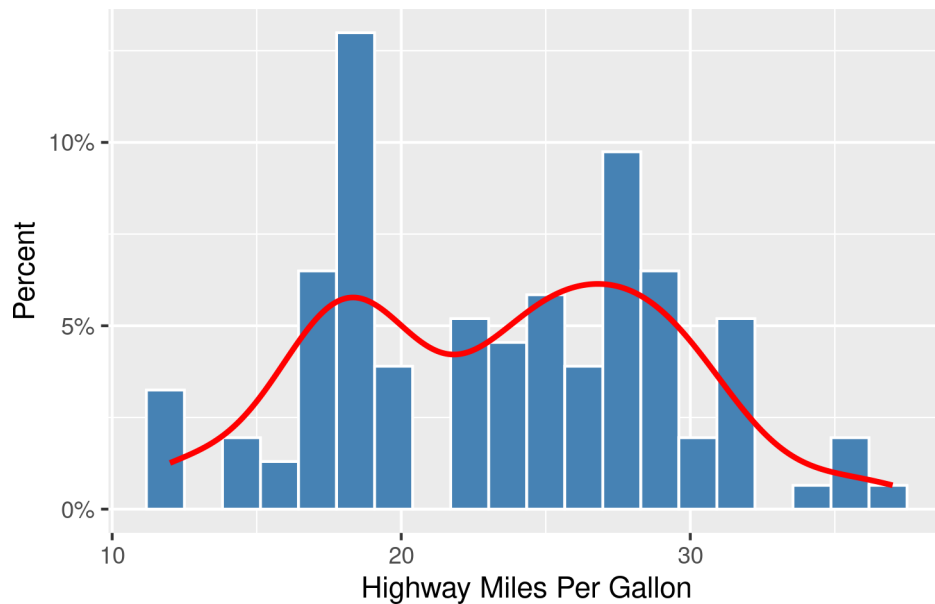
Histogram with percentages



```
ggplot(cars2008, aes(x = hwy, y = ..density..)) +  
  geom_histogram(bins = 20, color = "white", fill = "steelblue") + #4  
  scale_y_continuous(labels = scales::percent) +  
  geom_density(color = "red", size = 1) +  
  labs(title = "Histogram with density curve",
```

```
y = "Percent" ,
x = "Highway Miles Per Gallon")
```

Histogram with density curve



The first histogram #1

- demonstrates the default plot when no options are specified.
 - In this case, 30 bins are created.
- For the second histogram #2,
 - 20 bins, a steel blue fill, and a white border color are specified.
 - In addition, more informative labels have been added.
- The number of bins can strongly influence the appearance of the histogram.
- It is a good idea to experiment with the bins value
 - until you find one that captures the distribution well.
 - With 20 bins, it appears that there are two peaks to the distribution
 - * one around 13 mpg
 - * and one around 20.5 mpg.
- The third histogram #3 plots the data as percents rather than frequencies.
 - This is accomplished by assigning
 - * the built-in variable **density** to the y axis.
 - The scales package is used to format the y-axis as percents.
 - The fourth histogram #4 is similar to the previous plot,
 - * but adds a density curve.
 - * The density curve is a kernel density estimate
 - * and is described in the next section.
 - It provides a smoother description of the distribution of scores.
 - The `geom_density()` function is used
 - * to plot the kernel curve in a red color
 - * and a width that's slightly larger the default thickness for lines.
 - The density curve also suggests a bimodal distribution (two peaks).

6.1.2.6 Kernel density plots


- Technically, kernel density estimation
 - is a nonparametric method for estimating

- * the probability density function of a random variable.
- Basically, we're trying to draw a smoothed histogram,
 - where the area under the curve equals one.
- Although the mathematics are beyond the scope now,
 - density plots can be an effective way
 - * to view the distribution of a continuous variable.
- The format for a density plot is
 - `ggplot(data, aes(x = contvar)) + geom_density()`
 - * where data is a data frame
 - * and contvar is a continuous variable.

#1 Default density plot #2 Filled density plot #3 Print default bandwidth #4 Density plot with smaller bandwidth


```
data(mpg, package = "ggplot2")
cars2008 <- mpg[mpg$year == 2008,]
```

```
ggplot(cars2008, aes(x = cty)) +      #1
  geom_density() +
  labs(title = "Default kernel density plot")
```



```
ggplot(cars2008, aes(x = cty)) +      #2
  geom_density(fill = "red") +
  labs(title = "Filled kernel density plot")
```

Filled kernel density plot

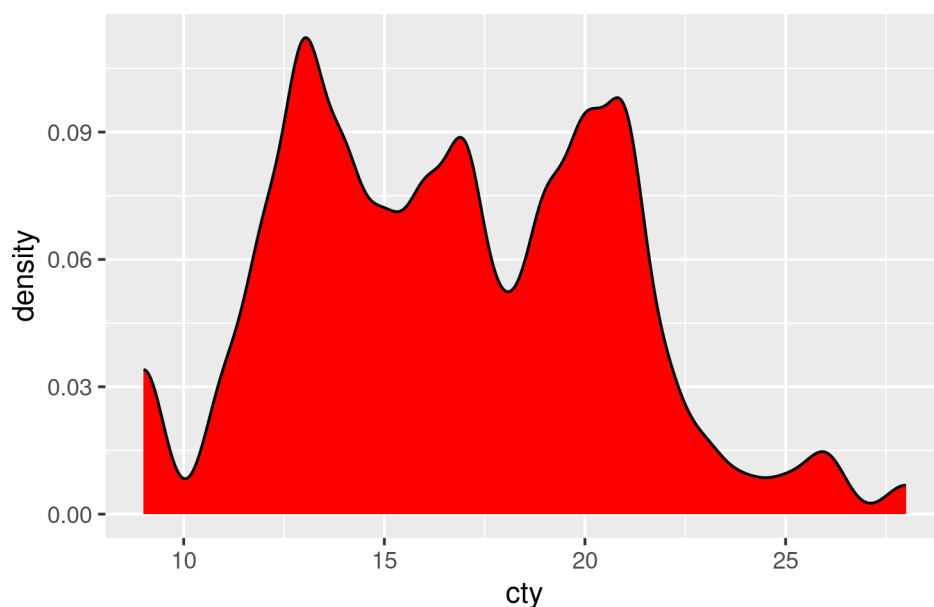


```
bw.nrd0(cars2008$cty)      #3
```

```
## [1] 1.408399
```

```
ggplot(cars2008, aes(x = cty)) +      #4
  geom_density(fill = "red", bw = .5) +
  labs(title = "Kernel density plot with bw=0.5")
```

Kernel density plot with bw=0.5



The default kernel density plot is given first #1.

- In the second example #2, the area under the curve is fill with red.
- The smoothness of the curve is controlled with a bandwidth parameter,
 - which is calculated from the data being plotted.
 - The code `bw.nrd0(cars2008$cty)` displays this value (1.408) #3.

- Using a larger bandwidth will give a smoother curve with less details.
 - A smaller value will give a more squiggly curve.
- The third example uses a smaller bandwidth (`bw=`),
 - allowing us to see more detail #4.
- As with the bins parameter for histograms,
 - it is a good idea to try several bandwidth values
 - to see which value helps you visualize the data most effectively.

Kernel density plots can be used to compare groups.

- This is a highly underutilized approach,
 - probably due to a general lack of easily accessible software.


For this example, we'll compare the 2008 city gas mileage estimates

- for 4-, 6-, and 8- cylinder cars.
- There are only a handful of cars with 5 cylinders
 - so we will drop them from the analyses.


#1 Prepare the data #2 Plots the density curves #3 Plot filled density curves

```
data(mpg, package = "ggplot2")      #1
cars2008 <- mpg[mpg$year == 2008 & mpg$cyl != 5,]
cars2008$Cylinders <- factor(cars2008$cyl)

ggplot(cars2008, aes(
  #2
  x = cty,
  color = Cylinders,
  linetype = Cylinders
)) +
  geom_density(size = 1) +
  labs(title = "Fuel Efficiency by Number of Cylinders",
       x = "City Miles per Gallon")
```



```
ggplot(cars2008, aes(x = cty, fill = Cylinders)) +      #3
  geom_density(alpha = .4) +
  labs(title = "Fuel Efficiency by Number of Cylinders",
       x = "City Miles per Gallon")
```



First, a fresh copy of the data is loaded

- and 2008 data for cars with 4, 6, or 8 cylinders are retained #1.
- The number of cylinders (cyl) is saved as a categorical factor (Cylinders).
- The transformation is required
 - because ggplot2 expects the grouping variable to be categorical
 - (and cyl is stored as a continuous variable).
- A kernel density curve is plotted for each level of the Cylinders variable #2.
 - Both the color (red, green, blue)
 - and line type (solid, dotted, dashed)
 - are mapped to the number of cylinders.
- Finally, the same plot is produced with filled curves #3.
 - Transparency is added ($\alpha = 0.4$),
 - * since the filled curves overlap
 - * and we want to be able to see each one.

Overlapping kernel density plots


- can be a powerful way to compare groups of observations
 - on an outcome variable.
- Here you can see both the shapes of the distributions
 - and the amount of overlap between groups.

6.1.2.7 Box Plots

- A box-and-whiskers plot describes
 - the distribution of a continuous variable
 - * by plotting its five-number summary:
 - the minimum,
 - lower quartile (25th percentile),

- median (50th percentile),
- upper quartile (75th percentile),
- and maximum.
- It can also display observations
 - * that **may be outliers** (values outside the range of $\pm 1.5 \cdot IQR$,
 - where IQR is the interquartile range
 - defined as the upper quartile minus the lower quartile).

```
ggplot(mtcars, aes(x = "", y = mpg)) +
  geom_boxplot() +
  labs(y = "Miles Per Gallon", x = "", title = "Box Plot")
```



```
boxplot.stats(mtcars$mpg)
```

```
## $stats
## [1] 10.40 15.35 19.20 22.80 33.90
##
## $n
## [1] 32
##
## $conf
## [1] 17.11916 21.28084
##
## $out
## numeric(0)
```

By default, each whisker extends to the most extreme data point,

- which is no more than 1.5 times the interquartile range for the box.
- Values outside this range are depicted as dots.

For example, in this sample of cars,

- the median mpg is 17,
 - 50% of the scores fall between 14 and 19,
 - the smallest value is 9,
 - and the largest value is 35.
- How did I read this so precisely from the graph?




Figure 1: And annotated


- Issuing `boxplot.stats(mtcars$mpg)`
 - * prints the statistics used to build the graph.
- There are four outliers
 - (greater than the upper hinge of 26).
 - These values would be expected to occur
 - * less than 1% of the time in a normal distribution.

6.1.2.7.1 Using parallel box plots to compare groups

- Box plots are a useful method of comparing
 - the distribution of a quantitative variable
 - * across the levels of a categorical variable.
 - Once again, let's compare city gas mileage for 3-, 6-, and 8-cylinder cars,
 - * but this time use both 1999 and 2008 data.
 - Since there are only a few 5-cylinder cars,
 - * we will delete them.
 - We'll also convert year and cyl
 - * from continuous numeric variables
 - * into categorical (grouping) factors.

```
library(ggplot2)
cars <- mpg[mpg$cyl != 5,]
cars$Cylinders <- factor(cars$cyl)
cars$Year <- factor(cars$year)

ggplot(cars, aes(x = Cylinders, y = cty)) +
  geom_boxplot() +
  labs(x = "Number of Cylinders",
       y = "Miles Per Gallon",
       title = "Car Mileage Data")
```



You can see that there's a good

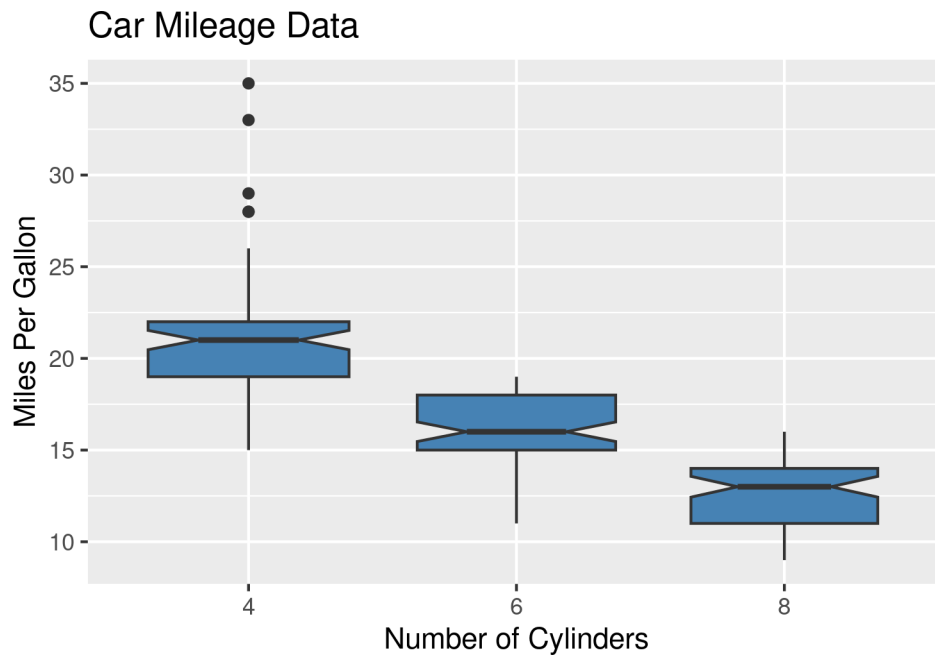
- separation of groups based on gas mileage,

- with fuel efficiency dropping as the number of cylinders increases.
- There are also four outliers
 - (cars with unusually high mileage)
 - in the four-cylinder group.

Box plots are very versatile.

- By adding `notch=TRUE`,
 - you get notched box plots.
- If two boxes' notches don't overlap,
 - there's strong evidence that their medians differ

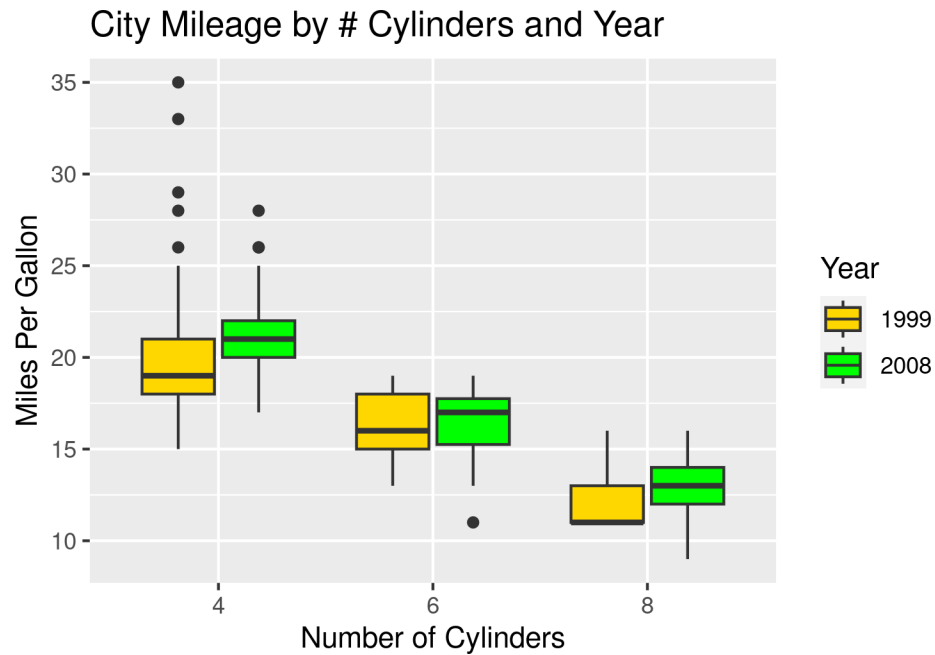
```
ggplot(cars, aes(x = Cylinders, y = cty)) +
  geom_boxplot(notch = TRUE,
               fill = "steelblue",
               varwidth = TRUE) +
  labs(x = "Number of Cylinders",
       y = "Miles Per Gallon",
       title = "Car Mileage Data")
```



The fill option fills the box plots with a red color.

- In a standard box plot,
 - the box width has no meaning.
- Adding `varwidth=TRUE`,
 - draws box widths proportional to the square roots
 - * of the number of observations in each group.
- The `scale_fill_manual()` function has been added
 - in order to customize the fill colors.

```
ggplot(cars, aes(x = Cylinders, y = cty, fill = Year)) +
  geom_boxplot() +
  labs(x = "Number of Cylinders",
       y = "Miles Per Gallon",
       title = "City Mileage by # Cylinders and Year") +
  scale_fill_manual(values = c("gold", "green"))
```





6.1.2.7.2 Violin Plots

- Before we end our discussion of box plots,
 - it's worth examining a variation called a **violin plot**.
- A violin plot is a combination
 - of a box plot and a kernel density plot.
- You can create one using the `geom_violin()` function.

```
cars <- mpg[mpg$cyl != 5, ]
cars$Cylinders <- factor(cars$cyl)

ggplot(cars, aes(x = Cylinders, y = cty)) +
  geom_boxplot(width = 0.2,
               fill = "green") +
  geom_violin(fill = "gold",
              alpha = 0.3) +
  labs(x = "Number of Cylinders",
       y = "City Miles Per Gallon",
       title = "Violin Plots of Miles Per Gallon")
```



The width of the box plots

- are set to 0.2
 - so that they will fit inside the violin plots.
- The violin plots are set with a transparency level of 0.3
 - so that the box plots are still visible.

Violin plots are basically kernel density plots

- superimposed in a mirror-image fashion over box plots.
 - The middle lines are the medians,
 - the black boxes range from the lower to the upper quartile,
 - and the thin black lines represent the whiskers.
 - Dots are outliers.
- The outer shape provides the kernel density plot.
 - Here we can see that the distribution of
 - * gas mileage for 8-cylinder cars
 - may be bimodal
 - * a fact that is obscured by using box plots alone.

6.1.2.8 Dot Plots

- Dot plots provide a method of plotting
 - a large number of labeled values on a simple horizontal scale.
- You create them with the `dotchart()` function,
 - using the format
 - `ggplot(data, aes(x=contvar, y=catvar)) + geom_point()`

Here's an example using the highway gas mileage


- for the 2008 automobiles in the mpg dataset.
- Highway gas mileage is averaged by car model.

```
plotdata <- mpg %>%
  filter(year == "2008") %>%
  group_by(model) %>%
```

```
summarize(meanHwy = mean(hwy))
plotdata
```


```
## # A tibble: 38 x 2
##   model          meanHwy
##   <chr>          <dbl>
## 1 4runner 4wd      18.5
## 2 a4              29.3
## 3 a4 quattro      26.2
## 4 a6 quattro      24
## 5 altima          29
## 6 c1500 suburban 2wd 18
## 7 camry           30
## 8 camry solara    29.7
## 9 caravan 2wd     22.2
## 10 civic          33.8
## # ... with 28 more rows
```

```
ggplot(plotdata, aes(x = meanHwy, y = model)) +
  geom_point() +
  labs(x = "Miles Per Gallon",
       y = "",
       title = "Gas Mileage for Car Models")
```



```
ggplot(plotdata, aes(x = meanHwy, y = reorder(model, meanHwy))) +
  geom_point() +
  labs(
    x = "Miles Per Gallon",
    y = "",
    title = "Gas Mileage for Car Models",
    subtitle = "with standard error bars"
  )
```

Gas Mileage for Car Models with standard error bars

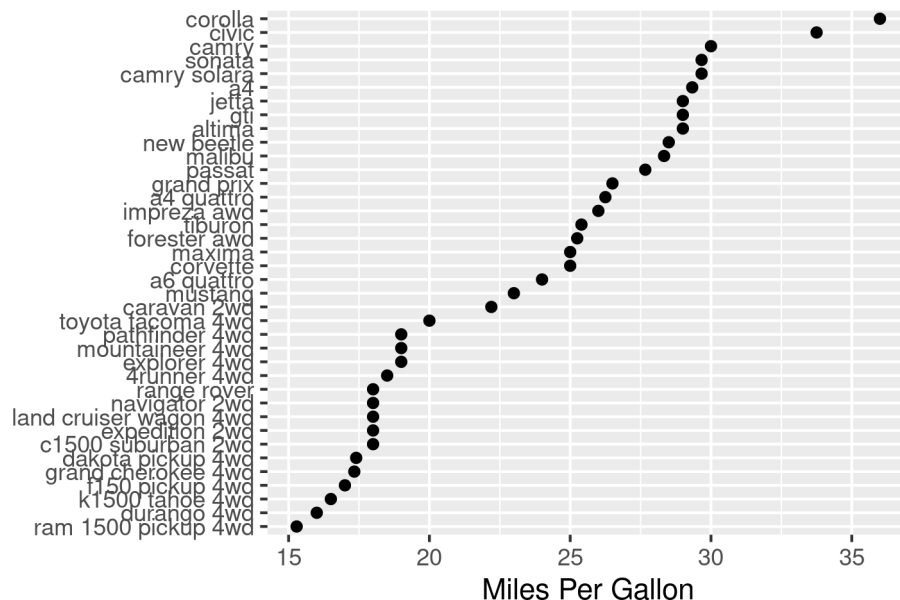


This graph allows you to see

- the mpg for each car model on the same horizontal axis.
- Dot plots typically become most useful
 - when they're sorted.
- The following code sorts the cars
 - from lowest to highest mileage.

```
ggplot(plotdata, aes(x = meanHwy, y = reorder(model, meanHwy))) +
  geom_point() +
  labs(x = "Miles Per Gallon",
       y = "",
       title = "Gas Mileage for Car Models")
```

Gas Mileage for Car Models



You can gain significant insight from the dot plot in this example

- because each point is labeled,
 - the value of each point is inherently meaningful,
 - and the points are arranged in a manner that promotes comparisons.
- But as the number of data points increases,
 - the utility of the dot plot decreases.

6.1.2.9 Summary

- 1. Bar charts (and to a lesser extent pie charts and tree maps)
 - can be used to gain insight into **the distribution of a categorical variable**.
- 2. Stacked, grouped, and filled bar charts
 - can help you understand **how groups differ on a categorical outcome**.
- 3. Histograms, box plots, violin plots, and dot plots
 - can help you **visualize the distribution of continuous variables**.
- 4. Overlapping kernel density plots and parallel box plots
 - can help you **visualize group differences on a continuous outcome variable**.

6.1.2.10 Links

- Robert I. Kabacoff, R in Action, 3rd Edition, Manning Publications 2020