

2201-353-353m-453-PCA-PCR-PLS-FA-TidyModeling Class 07a-p

Profs: R. H. French, L. S. Bruckman, P. Leu, K. Davis, S. Cirlos

TAs: W. Oltjen, K. Hernandez, M. Li, M. Li, D. Colvin

07 March, 2023

Contents

7.1.2.1	Class Readings, Assignments, Textbooks Syllabus Topics	1
7.1.2.1.1	Reading, Lab Exercises, SemProjects	1
7.1.2.1.2	Textbooks	2
7.1.2.1.3	Tidyverse Cheatsheets, Functions and Reading Your Code	2
7.1.2.1.4	Syllabus	3
7.1.2.2	ISLR6 is Model and Variable Selection	3
7.1.2.3	A systematic/programmatic approach to data-driven modeling	3
7.1.2.4	Tidy Modeling	3
7.1.2.4.1	Introduction	3
7.1.2.5	An Example	6
7.1.2.5.1	Load only the <code>tidymodels</code> library	6
7.1.2.5.2	Pre-process	7
7.1.2.5.3	Data Sampling	7
7.1.2.5.4	Pre-process interface	8
7.1.2.5.5	Execute the pre-processing	10
7.1.2.5.6	Model training	11
7.1.2.5.7	Predictions	12
7.1.2.5.8	Model Validation	13
7.1.2.5.9	Per classifier metrics	14
7.1.2.5.10	Summary	18
7.1.2.6	Links	18

7.1.2.1 Class Readings, Assignments, Textbooks Syllabus Topics

7.1.2.1.1 Reading, Lab Exercises, SemProjects

- Readings:
 - For today: ISLR8, DL08,09
 - For next class: ISLR9, (R4DS26-30)
- Laboratory Exercises:
 - LE4 given out
 - LE4 due March 9th
- Office Hours: (Class Canvas Calendar for Zoom Link)
 - Wednesdays @ 4:00 PM to 5:00 PM
 - Saturdays @ 3:00 PM to 4:00 PM
 - **Office Hours are on Zoom, and recorded**
- Semester Projects
 - Office Hours for SemProjs: Mondays at 4pm on Zoom

- DSCI 453 Students Biweekly Updates Due
 - * Update #4 is Due **Friday March 10th**
- DSCI 453 Students
 - * Next Report Out # is Due ** **
- All DSCI 353/353M/453, E1453/2453 Students:
 - * **Peer Grading of Report Out #1 is Due Thursday March 2nd**
- Exams
 - * MidTerm: **Thursday March 9th**, in class or remote, 11:30 - 12:45 PM
 - **CWRU Spring Break is March 13th to March 17, so NO CLASS**
 - * Final: **Thursday May 4th**, 2023, 12:00PM - 3:00PM, Nord 356 or remote

7.1.2.1.2 Textbooks

- Text Books for DSCI353/353M/453
 - [R4DS: Wickham: R for Data Science](#)
 - [ISLR: Intro to Statistical Learning with R, 2nd Ed.](#)
 - DLwR: Deep Learning with R 2nd. Ed., Chollet, Kalinowski, Allaire,
 - [DLGB: Deep Learning, Goodfellow, Bengio, Courville](#)
- Magazine Articles about Deep Learning
 - DL1 to DL12 are “Deep Learning” articles in 3-readings/2-articles/
- Books from DSCI351/351M/451
 - [Peng: R Programming for Data Science](#)
 - [Peng: Exploratory Data Analysis with R](#)
 - [Open Intro Stats, v4](#)
 - [R4DS: Wickham: R for Data Science](#)

7.1.2.1.3 Tidyverse Cheatsheets, Functions and Reading Your Code

- Look at the Tidyverse Cheatsheet
 - **Tidyverse For Beginners Cheatsheet**
 - * In the Git/20s-dsci353-353m-453-prof/3-readings/3-CheatSheets/ folder
 - **Data Wrangling with dplyr and tidyr Cheatsheet**

Tidyverse Functions & Conventions

- The pipe operator `%>%`
- Use `dplyr::filter()` to subset data row-wise.
- Use `dplyr::arrange()` to sort the observations in a data frame
- Use `dplyr::mutate()` to update or create new columns of a data frame
- Use `dplyr::summarize()` to turn many observations into a single data point
- Use `dplyr::arrange()` to change the ordering of the rows of a data frame
- Use `dplyr::select()` to choose variables from a tibble,
 - * keeps only variables you mention
- Use `dplyr::rename()` keeps all the variables and renames variables
 - * `rename(iris, petal_length = Petal.Length)`
- These can be combined using `dplyr::group_by()`
 - * which lets you perform operations “by group”.
- The `%in%` matches conditions provided by a vector using the `c()` function
- The **forcats** package has tidyverse functions
 - * for factors (categorical variables)
- The **readr** package has tidyverse functions
 - * to `read_...`, `melt_...`, `col_...`, `parse_...` data and objects

Reading Your Code: Whenever you see

- The assignment operator `<-`, think “**gets**”
- The pipe operator, `%>%`, think “**then**”

7.1.2.1.4 Syllabus

7.1.2.2 ISLR6 is Model and Variable Selection

- So Regularization is reducing the dimensions / degrees of freedom in a model
 - By choosing fewer variables (e.g. variable selection)
 - By forcing β coefficients towards or to zero (e.g. Lasso)
 - Using constrained functions (e.g. natural splines)
 - By combining variables (e.g. PCA)

So its useful to recognize this need

- And learn all the ways we try to accomplish it

[Principal component analysis](#)

[Principal component regression](#)

[Partial least squares regression](#)

- PCA, PCR, PLS are covered in
 - ISLR 6.3
 - ISLR 10.2

[Factor analysis](#)

- Factor Analysis is covered in ESL Chapter 14.7

The similarities and differences among these approaches

- Can be subtle, or fine
- They are very powerful approaches.

7.1.2.3 A systematic/programmatic approach to data-driven modeling

- Its not just building one model
- Its building a series of models
- Using training and testing splits for Validation

So in Lab Exercises

- It was written so that you would do a large number of models
- And be discussing all the modeling choices
- And be able to demonstrate why you chose each choice
- And how, by quantitative statistical metrics,
 - The one you choose is the best one

7.1.2.4 Tidy Modeling

- Using Tidy approaches to enable efficient meta-modeling and model selection

7.1.2.4.1 Introduction

- Lets look at the `tidymodels` R package

Let frame where `tidymodels` fit

Day:Date	Foundation	Practicum	Readings(optional)	Due(optional)
w01a:Tu:1/17/23 w01b:Th:1/19/23	Markov Cluster Stat. Learning, Approach	R, Rstudio IDE, Git Bash, Git, Class Repo	ISLR1,2 (R4DS-1-3)	(LE0)
w02a:Tu:1/24/23 w02b:Th:1/26/23	Lin. Regr. Bias-Var. Train/Test, Bias vs. Vari.	SemProjs; Regr. Ovrw Tidyverse Review	ISLR3,(R4DS-4-6) DL01 DL02 (R4DS-7,8)	(LE0:Due) LE1
w02Pr:Fr:1/27/23	ADD DROP	DEADLINE		453 Update 1
w03a:Tu:1/31/23	Logistic Regr. Classif	Pred. Analytics, Regr.	DL03,ISLR4	
w03b:Th:2/2/23 w03Sa:2/4/23	LDA/QDA	ggPlot2, Code Expect.	DL04, DL05	LE1:Due, LE2 LE1:Due
w04a:Tu:2/7/23 w04b:Th:2/9/23	Resample Cross-Valid. DL, ML Overview	ggplot Multilevel Mod.	ISLR5 ISLR6 (R4DS9-16)	
w04Pr:Fr:2/10/23				453 Update 2
w05a:Tu:2/14/23	Resampling: Bootstrap	Bootstrap Mixed Effects	DL2R1, DL06,07	LE2:Due, LE3
w05b:Th:2/16/23 w05Pr:Fr:2/17/23	Subset Selec., Shrink.	Dim. Red. PCA	DLwR2	453 Rep. Out 1
w06a:Tu:2/21/23 w06b:Th:2/23/23	ML with NNs Beyond Linear Modls	ggplot, clustering Feature Select., Caret	DLwR3 ISLR7 (R4DS22-25)	LE3:Due, LE4
w06Pr:Fr:2/24/23				453 Update 3
w07a:Tu:2/28/23	Dec. Trees, Rand. Forest	Tidy Modeling	ISLR8, DL08,09	
w07b:Th:3/2/23	MidTerm Review, SVM	SVM, SVR, ROC	ISLR9 (R4DS26-30)	Peer Review 1
w08a:Tu:3/7/23	ML Overview	, Keras/TF2, Torch	ISLR10.1,10.2	
w08b:Th:3/9/23 w08Pr:Fr:3/10/23	MIDTERM EXAM		DL10,11	LE4:Due LE5 453 Update 4
Tu:3/14/23 Th:3/16/23	SPRING SPRING	BREAK BREAK	ISLR10.3,10.4 ISLR10.5,10.6,	
w09a:Tu:3/21/23	Deep Learning	TF2 Keras Intro		ISLR10.7,10.8, DLwR3
w09b:Th:3/23/23 w09Pr:Fr:3/24/23	Computer Vision, CNN	CNN w/TF2, Overfit	DLwR4, DL12,13	453 Rep. Out 2
w10a:Tu:3/28/23	Deep Learn Intro	NN Types	DLwR5 Hinton ImageNet	
w10b:Th:3/30/23 w10Pr:Fr:3/31/23	DL CNN,RNN ImageNet	NN Types, CNN w/TF2		453 Upd.5 & PrRev 2
Sa:4/1/23				LE5:Due LE6
w11a:Tu:4/4/23	Fitting NNs	AUC,Prec,Recall Fruit		
w11b:Th:4/6/23	NLP, Graphs & ML		LeCun DL Rev. 2015	
w12a:Tu:4/11/23	Graphs & ML	NLP with sequences	DLwR6	
w12b:Th:4/13/23	NLP w attention	Graph Repr Proc Wrk-flw		LE6:Due LE7
w13a:Tu:4/18/23	DL Frameworks	Explaining DL w Lime		
w13b:Th:4/20/23 w13Pr:Fr:4/21/23	Linux Distros XGBoost	Explain Preds	Deep Dream	453 Rep. Out 3 Due
w14a:Tu:4/25/23 w14b:Th:4/27/23	Transformers Final Exam Review	Torch NN & DeepLearn		LE7:Due
w14Pr:Fr:4/28/23				Peer Rev 3 Due
	FINAL EXAM	Th. 5/4/23, 12-3pm	Nord 356 & Zoom	
	453 Final PDF Report	Fr. 4/29, 11:59pm		

Table 1: DSCI353-353M-453 Weekly Syllabus. R4DS-x.y, OISx.y, ISLRx.y, DLwRx.y, DLGBx.y refers to chapters and sections assigned as reading in our textbooks. DLx are deep learning articles. March 2, 2023 4

Figure 1: DSCI351-351M-451 Syllabus

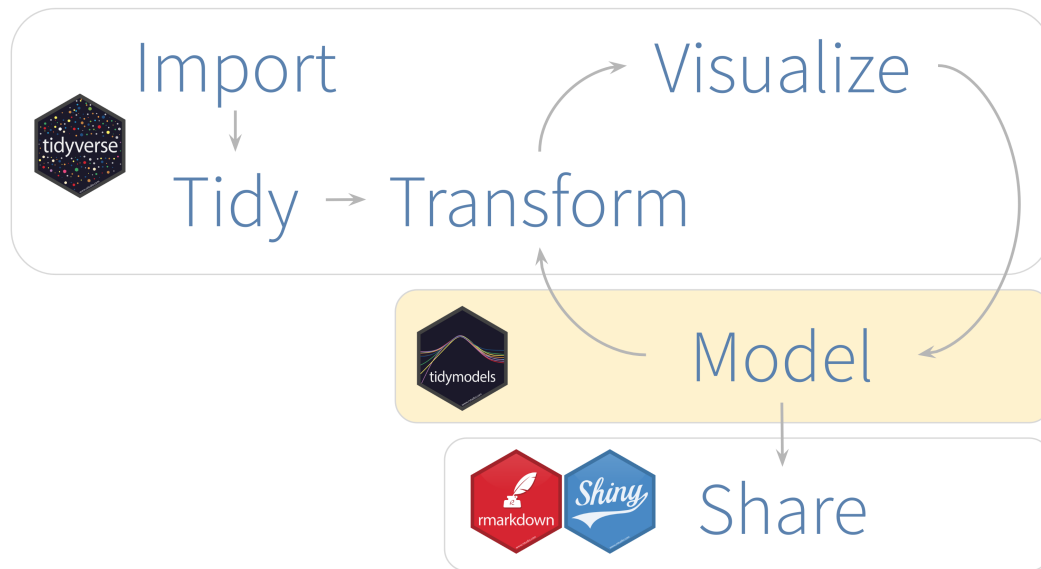


Figure 2: Figure 1. Tidy Model Work Flow, from R4DS

- in a data-driven analysis

Lets illustrates what step each package covers.

Developing models is a bit of a more complex

- developing models can benefit from having a `tidyverse`-friendly interface
- that is where `tidymodels` package can come in

It is important to clarify that

- the group of packages that make up `tidymodels`
 - do not implement statistical models themselves.
- Instead, they focus on making all the tasks
 - around fitting the model much easier.
- Those tasks are
 - data pre-processing and
 - results validation.

So `tidymodels` is an R metapackage

- Similar in this sense to the `caret` package
- [caret Classification and Regression Training](#)

In a way, the Model step itself has sub-steps.

- For these sub-steps, `tidymodels` provides one or several packages.
- Here we'll look at four `tidymodels` packages:

The Tidymodel packages

- `rsample` - Different types of re-samples
- `recipes` - Transformations for model data pre-processing
- `parsnip` - A common interface for model creation
- `yardstick` - Measure model performance

The following diagram

- illustrates each modeling step, and
- lines up the `tidymodels` packages we'll discuss

Pre-Process → Train → Validate



Figure 3: The modeling steps

In a given analysis, a `tidyverse` package may or may not be used.

- For example, not all projects need to work with time variables,
 - so there is no need to use functions from the `hms` package.
- The same idea applies to `tidymodels`.
- Depending on what type of modeling is going to be done,
 - only functions from some its packages will be used.

7.1.2.5 An Example

- Lets use the `iris` data set for an example.
 - Its data is already imported
 - and sufficiently tidy to move directly to modeling.

7.1.2.5.1 Load only the `tidymodels` library

- We can just load the `tidymodel` metapackage
 - Apart from loading its core modeling packages,
 - `tidymodels` also conveniently loads some `tidyverse` packages,
 - * including `dplyr` and `ggplot2`.
 - Throughout this exercise, we will use some functions out of those packages,
 - * but we don't have to explicitly load them into our R session.

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --
## v ggplot2 3.4.1      v purrr   1.0.1
## v tibble  3.1.8      v dplyr  1.1.0
## v tidyr   1.3.0      v stringr 1.5.0
## v readr   2.1.3      v forcats 1.0.0

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()

library(tidymodels)

## -- Attaching packages ----- tidymodels 1.0.0 --
## v broom      1.0.3      v rsample      1.1.1
## v dials      1.1.0      v tune         1.0.1
```

```
## v infer          1.0.4      v workflows      1.1.3
## v modeldata      1.1.0      v workflowsets 1.0.0
## v parsnip        1.0.4      v yardstick   1.1.0
## v recipes        1.0.4

## -- Conflicts ----- tidymodels_conflicts() --
## x scales::discard() masks purrr::discard()
## x dplyr::filter()   masks stats::filter()
## x recipes::fixed() masks stringr::fixed()
## x dplyr::lag()      masks stats::lag()
## x yardstick::spec() masks readr::spec()
## x recipes::step()   masks stats::step()
## * Dig deeper into tidy modeling with R at https://www.tmw.r.org
```

7.1.2.5.2 Pre-process

- This step focuses on making data
 - suitable for modeling by using data transformations.

All transformations can be accomplished

- with `dplyr`, or other `tidyverse` packages

Consider using `tidymodels` packages

- when model development is more heavy and complex.

7.1.2.5.3 Data Sampling

- The `rsample::initial_split()` function is specially built
 - to separate the data set into a training and testing set.
 - By default,
 - * it holds 3/4 of the data for training
 - * and the rest for testing.
 - That can be changed by passing the `prop` argument.
 - This function generates an `rsplit` object,
 - * not a data frame.
 - The printed output shows the row count
 - * for testing, training, and total.

```
iris_split <- initial_split(iris, prop = 0.6)
iris_split
```

```
## <Training/Testing/Total>
## <90/60/150>
```

```
## <90/60/150>
```

To access the observations reserved for training,

- use the `training()` function.

Similarly, use `testing()`

- to access the testing data.

```
iris_split %>%
  training() %>%
  glimpse()
```

```
## Rows: 90
## Columns: 5
## $ Sepal.Length <dbl> 5.4, 6.0, 6.1, 6.0, 6.6, 5.8, 6.3, 5.0, 6.2, 6.7, 5.0, 4.~
## $ Sepal.Width <dbl> 3.9, 2.2, 3.0, 3.0, 3.0, 2.6, 2.5, 3.4, 3.4, 3.1, 3.4, 3.~
## $ Petal.Length <dbl> 1.3, 4.0, 4.9, 4.8, 4.4, 4.0, 5.0, 1.6, 5.4, 5.6, 1.5, 1.~
## $ Petal.Width <dbl> 0.4, 1.0, 1.8, 1.8, 1.4, 1.2, 1.9, 0.4, 2.3, 2.4, 0.2, 0.~
## $ Species <fct> setosa, versicolor, virginica, virginica, versicolor, ver~

## Observations: 90
## Variables: 5
## $ Sepal.Length <dbl> 5.1, 4.9, 4.7, 4.6, 5.0, 5.4, 4.6, 5.0, 4.9, 5.4, 4...
## $ Sepal.Width <dbl> 3.5, 3.0, 3.2, 3.1, 3.6, 3.9, 3.4, 3.4, 3.1, 3.7, 3...
## $ Petal.Length <dbl> 1.4, 1.4, 1.3, 1.5, 1.4, 1.7, 1.4, 1.5, 1.5, 1.5, 1...
## $ Petal.Width <dbl> 0.2, 0.2, 0.2, 0.2, 0.2, 0.4, 0.3, 0.2, 0.1, 0.2, 0...
## $ Species <fct> setosa, setosa, setosa, setosa, setosa, setosa, set...
```

These sampling functions

- are courtesy of the `rsample` package,
- which is part of `tidymodels`.

7.1.2.5.4 Pre-process interface

- In `tidymodels`, the `recipes` package
 - provides an interface that specializes in data pre-processing.

Within the package,

- the functions that start, or execute,
 - the data transformations are named after cooking actions.
- That makes the interface more user-friendly.

For example:

- `recipe()` - Starts a new set of transformations to be applied,
 - similar to the `ggplot()` command.
 - Its main argument is the model's formula.
- `prep()` - Executes the transformations
 - on top of the data that is supplied
 - (typically, the training data).

Each data transformation is a step.

- Functions correspond to specific types of steps,
 - each of which has a prefix of `step_`

There are several `step_` functions;

- in this example, we will use three of them:
- `step_corr()`
 - Removes variables
 - that have large absolute correlations with other variables
- `step_center()`
 - Normalizes numeric data to have a mean of zero
- `step_scale()`
 - Normalizes numeric data to have a standard deviation of one

Another nice feature is that the step can be applied

- to a specific variable,

- groups of variables,
- or all variables.

The `all_outcomes()` and `all_predictors()` functions

- provide a very convenient way to specify groups of variables.

For example, if we want the `step_corr()`

- to only analyze the predictor variables,
- we use `step_corr(all_predictors())`.
 - This capability saves us from having to enumerate each variable.

In the following example,

- we will put together
 - the `recipe()`, `prep()`, and `step` functions
 - to create a recipe object.
- The `training()` function
 - is used to extract that data set
 - from the previously created split sample data set.

```
iris_recipe <- training(iris_split) %>%
  recipe(Species ~.) %>%
  step_corr(all_predictors()) %>%
  step_center(all_predictors(), -all_outcomes()) %>%
  step_scale(all_predictors(), -all_outcomes()) %>%
  prep()
```

If we call the `iris_recipe` object,

- it will print details about the recipe.
- The Operations section describes what was done to the data.
- One of the operations entries in the example
 - explains that the correlation step
 - removed the `Petal.Length` variable.

```
iris_recipe

## Recipe
##
## Inputs:
##
##      role #variables
##  outcome      1
## predictor      4
##
## Training data contained 90 data points and no missing data.
##
## Operations:
##
## Correlation filter on Petal.Length [trained]
## Centering for Sepal.Length, Sepal.Width, Petal.Width [trained]
## Scaling for Sepal.Length, Sepal.Width, Petal.Width [trained]
## Data Recipe
##
## Inputs:
```

```
##      role #variables
## outcome      1
## predictor     4
##
## Training data contained 90 data points and no missing data.
##
## Operations:
##
## Correlation filter removed Petal.Length [trained]
## Centering for Sepal.Length, Sepal.Width, Petal.Width [trained]
## Scaling for Sepal.Length, Sepal.Width, Petal.Width [trained]
```

7.1.2.5.5 Execute the pre-processing

- The testing data can now be transformed
 - using the exact same
 - * steps,
 - * weights,
 - * and categorization
 - used to pre-process the training data.

To do this, another function

- with a cooking term is used: `bake()`.
- Notice that the `testing()` function is used
 - in order to extract the appropriate data set.

```
iris_testing <- iris_recipe %>%
  bake(testing(iris_split))

glimpse(iris_testing)
```

```
## Rows: 60
## Columns: 4
## $ Sepal.Length <dbl> -1.37182722, -1.49717184, -1.74786107, -1.12113798, -1.24~
## $ Sepal.Width <dbl> 0.21695305, 0.66582143, -0.45634952, -0.00748114, -0.2319~
## $ Petal.Width <dbl> -1.2513777, -1.1236861, -1.2513777, -1.3790693, -1.379069~
## $ Species <fct> setosa, setosa, setosa, setosa, setosa, setosa, setosa, s~
## Observations: 60
## Variables: 4
## $ Sepal.Length <dbl> -1.597601746, -1.138960096, 0.007644027, -0.7949788...
## $ Sepal.Width <dbl> -0.41010139, 0.71517681, 2.06551064, 1.61539936, 0.1...
## $ Petal.Width <dbl> -1.2085003, -1.2085003, -1.2085003, -1.0796318, -1.1...
## $ Species <fct> setosa, setosa, setosa, setosa, setosa, setosa, set...
```

Performing the same operation

- over the training data is redundant,
 - because that data has already been prepped.

To load the prepared training data

- into a variable,
 - we use `juice()`.
- It will extract the data
 - from the `iris_recipe` object.

```
iris_training <- juice(iris_recipe)

glimpse(iris_training)

## Rows: 90
## Columns: 4
## $ Sepal.Length <dbl> -0.49441489, 0.25765283, 0.38299745, 0.25765283, 1.009720~
## $ Sepal.Width <dbl> 1.78799238, -2.02738885, -0.23191533, -0.23191533, -0.231~
## $ Petal.Width <dbl> -0.99599447, -0.22984488, 0.79168791, 0.79168791, 0.28092~
## $ Species <fct> setosa, versicolor, virginica, virginica, versicolor, ver~
## Observations: 90
## Variables: 4
## $ Sepal.Length <dbl> -0.7949789, -1.0242997, -1.2536205, -1.3682809, -0.~
## $ Sepal.Width <dbl> 0.94023245, -0.18504575, 0.26506553, 0.04000989, 1.~
## $ Petal.Width <dbl> -1.2085003, -1.2085003, -1.2085003, -1.2085003, -1.~
## $ Species <fct> setosa, setosa, setosa, setosa, setosa, setosa, set...
```

7.1.2.5.6 Model training

- In R, there are multiple packages that fit the same type of model.
 - It is common for each package to provide a unique interface.
 - In other words,
 - * things such as an argument for the same model attribute
 - * is defined differently for each package.

For example, the `ranger` and `randomForest` packages

- fit Random Forest models.
- In the `ranger()` function,
 - to define the number of trees we use `num.trees`.
- In `randomForest` package,
 - that argument is named `ntree`.
- It is not easy to switch between packages
 - to run the same model.

Instead of replacing the modeling package,

- `tidymodels` replaces the interface.

Better said, `tidymodels`

- provides a single set of functions and arguments
 - to define a model.
- It then fits the model
 - against the requested modeling package.

In the example below,

- the `rand_forest()` function is used
 - to initialize a Random Forest model.
 - To define the number of trees,
 - the `trees` argument is used.
- To use the `ranger` version of Random Forest,
 - the `set_engine()` function is used.
- Finally, to execute the model,
 - the `fit()` function is used.
- The expected arguments are the formula and data.

- Notice that the model runs
 - on top of the juiced trained data.

```
iris_ranger <- rand_forest(trees = 100, mode = "classification") %>%
  set_engine("ranger") %>%
  fit(Species ~ ., data = iris_training)
```

The payoff is that

- if we now want to run the same model
 - against randomForest,
- we simply change the value in set_engine()
 - to randomForest.

```
iris_rf <- rand_forest(trees = 100, mode = "classification") %>%
  set_engine("randomForest") %>%
  fit(Species ~ ., data = iris_training)
```

It is also worth mentioning that

- the model is not defined
 - in a single, large function with a lot of arguments.
- The model definition is separated
 - into smaller functions
 - such as fit() and set_engine().
- This allows for a more flexible, and easier to learn, interface.

7.1.2.5.7 Predictions

- Instead of a vector,
 - the predict() function
 - * ran against a parsnip model
 - * returns a tibble.
 - By default, the prediction variable
 - * is called .pred_class.
 - In the example, notice that
 - * the baked testing data is used.

```
predict(iris_ranger, iris_testing)
```

```
## # A tibble: 60 x 1
##   .pred_class
##   <fct>
## 1 setosa
## 2 setosa
## 3 setosa
## 4 setosa
## 5 setosa
## 6 setosa
## 7 setosa
## 8 setosa
## 9 setosa
## 10 setosa
## # ... with 50 more rows
```

```
## # A tibble: 60 x 1
##   .pred_class
```

```
##      <fct>
## 1 setosa
## 2 setosa
## 3 setosa
## 4 setosa
## 5 setosa
## 6 setosa
## 7 setosa
## 8 setosa
## 9 setosa
## 10 setosa
## # ... with 50 more rows
```

It is very easy to add the predictions

- to the baked testing data
- by using `dplyr`'s `bind_cols()` function.

```
iris_ranger %>%
  predict(iris_testing) %>%
  dplyr::bind_cols(iris_testing) %>%
  glimpse()

## Rows: 60
## Columns: 5
## $ .pred_class <fct> setosa, setosa, setosa, setosa, setosa, setosa, setosa, s~
## $ Sepal.Length <dbl> -1.37182722, -1.49717184, -1.74786107, -1.12113798, -1.24~
## $ Sepal.Width <dbl> 0.21695305, 0.66582143, -0.45634952, -0.00748114, -0.2319~
## $ Petal.Width <dbl> -1.2513777, -1.1236861, -1.2513777, -1.3790693, -1.379069~
## $ Species <fct> setosa, setosa, setosa, setosa, setosa, setosa, setosa, s~
## Observations: 60
## Variables: 5
## $ .pred_class <fct> setosa, setosa, setosa, setosa, setosa, setosa, set...
## $ Sepal.Length <dbl> -1.597601746, -1.138960096, 0.007644027, -0.7949788...
## $ Sepal.Width <dbl> -0.41010139, 0.71517681, 2.06551064, 1.61539936, 0...
## $ Petal.Width <dbl> -1.2085003, -1.2085003, -1.2085003, -1.0796318, -1...
## $ Species <fct> setosa, setosa, setosa, setosa, setosa, setosa, set...
```

7.1.2.5.8 Model Validation

- Use the `metrics()` function
 - to measure the performance of the model.

It will automatically choose metrics

- appropriate for a given type of model.

The function expects a tibble

- that contains the actual results (truth)
- and what the model predicted (estimate).

```
iris_ranger %>%
  predict(iris_testing) %>%
  bind_cols(iris_testing) %>%
  metrics(truth = Species, estimate = .pred_class)
```

```
## # A tibble: 2 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 accuracy multiclass 0.917
## 2 kap     multiclass 0.874
```

```
## # A tibble: 2 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 accuracy multiclass 0.917
## 2 kap     multiclass 0.874
```

Because of the consistency of the new interface,

- measuring the same metrics
 - against the `randomForest` model
- is as easy as replacing the model variable
 - at the top of the code.

```
iris_rf %>%
  predict(iris_testing) %>%
  dplyr::bind_cols(iris_testing) %>%
  metrics(truth = Species, estimate = .pred_class)
```

```
## # A tibble: 2 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 accuracy multiclass 0.917
## 2 kap     multiclass 0.874
```

```
## # A tibble: 2 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 accuracy multiclass 0.883
## 2 kap     multiclass 0.824
```

7.1.2.5.9 Per classifier metrics

- It is easy to obtain the probability
 - for each possible predicted value
 - * by setting the `type` argument to `prob`.
 - That will return a tibble
 - * with as many variables
 - * as there are possible predicted values.
 - Their name will default
 - * to the original value name,
 - * prefixed with `.pred_`.

```
iris_ranger %>%
  predict(iris_testing, type = "prob") %>%
  glimpse()
```

```
## Rows: 60
## Columns: 3
## $ .pred_setosa    <dbl> 0.95980556, 0.98407143, 0.93060317, 0.92616270, 0.950~
## $ .pred_versicolor <dbl> 0.02852778, 0.00900000, 0.06939683, 0.06169444, 0.043~
## $ .pred_virginica  <dbl> 0.01166667, 0.006928571, 0.00000000, 0.012142857, 0~
```

```
## Observations: 60
## Variables: 3
## $ .pred_setosa      <dbl> 0.677480159, 0.978293651, 0.783250000, 0.983972...
## $ .pred_versicolor <dbl> 0.295507937, 0.011706349, 0.150833333, 0.001111...
## $ .pred_virginica   <dbl> 0.02701190, 0.01000000, 0.06591667, 0.01491667,...
```

Again, use `dplyr::bind_cols()`

- to append the predictions
- to the baked testing data set.

```
iris_probs <- iris_ranger %>%
  predict(iris_testing, type = "prob") %>%
  dplyr::bind_cols(iris_testing)

glimpse(iris_probs)
```

```
## Rows: 60
## Columns: 7
## $ .pred_setosa      <dbl> 0.95980556, 0.98407143, 0.93060317, 0.92616270, 0.950~
## $ .pred_versicolor <dbl> 0.02852778, 0.00900000, 0.06939683, 0.06169444, 0.043~
## $ .pred_virginica   <dbl> 0.011666667, 0.006928571, 0.000000000, 0.012142857, 0~
## $ Sepal.Length      <dbl> -1.37182722, -1.49717184, -1.74786107, -1.12113798, --
## $ Sepal.Width       <dbl> 0.21695305, 0.66582143, -0.45634952, -0.00748114, -0.~
## $ Petal.Width       <dbl> -1.2513777, -1.1236861, -1.2513777, -1.3790693, -1.37~
## $ Species           <fct> setosa, setosa, setosa, setosa, setosa, setosa, setos~
```

```
## Observations: 60
## Variables: 7
## $ .pred_setosa      <dbl> 0.677480159, 0.978293651, 0.783250000, 0.983972...
## $ .pred_versicolor <dbl> 0.295507937, 0.011706349, 0.150833333, 0.001111...
## $ .pred_virginica   <dbl> 0.02701190, 0.01000000, 0.06591667, 0.01491667,~
## $ Sepal.Length      <dbl> -1.597601746, -1.138960096, 0.007644027, -0.794...
## $ Sepal.Width       <dbl> -0.41010139, 0.71517681, 2.06551064, 1.61539936...
## $ Petal.Width       <dbl> -1.2085003, -1.2085003, -1.2085003, -1.0796318,~
## $ Species           <fct> setosa, setosa, setosa, setosa, setosa, setosa,...
```

Now that everything is in one tibble,

- it is easy to calculate curve methods.
- In this case we are using `gain_curve()`.

```
iris_probs %>%
  gain_curve(Species, .pred_setosa:.pred_virginica) %>%
  glimpse()
```

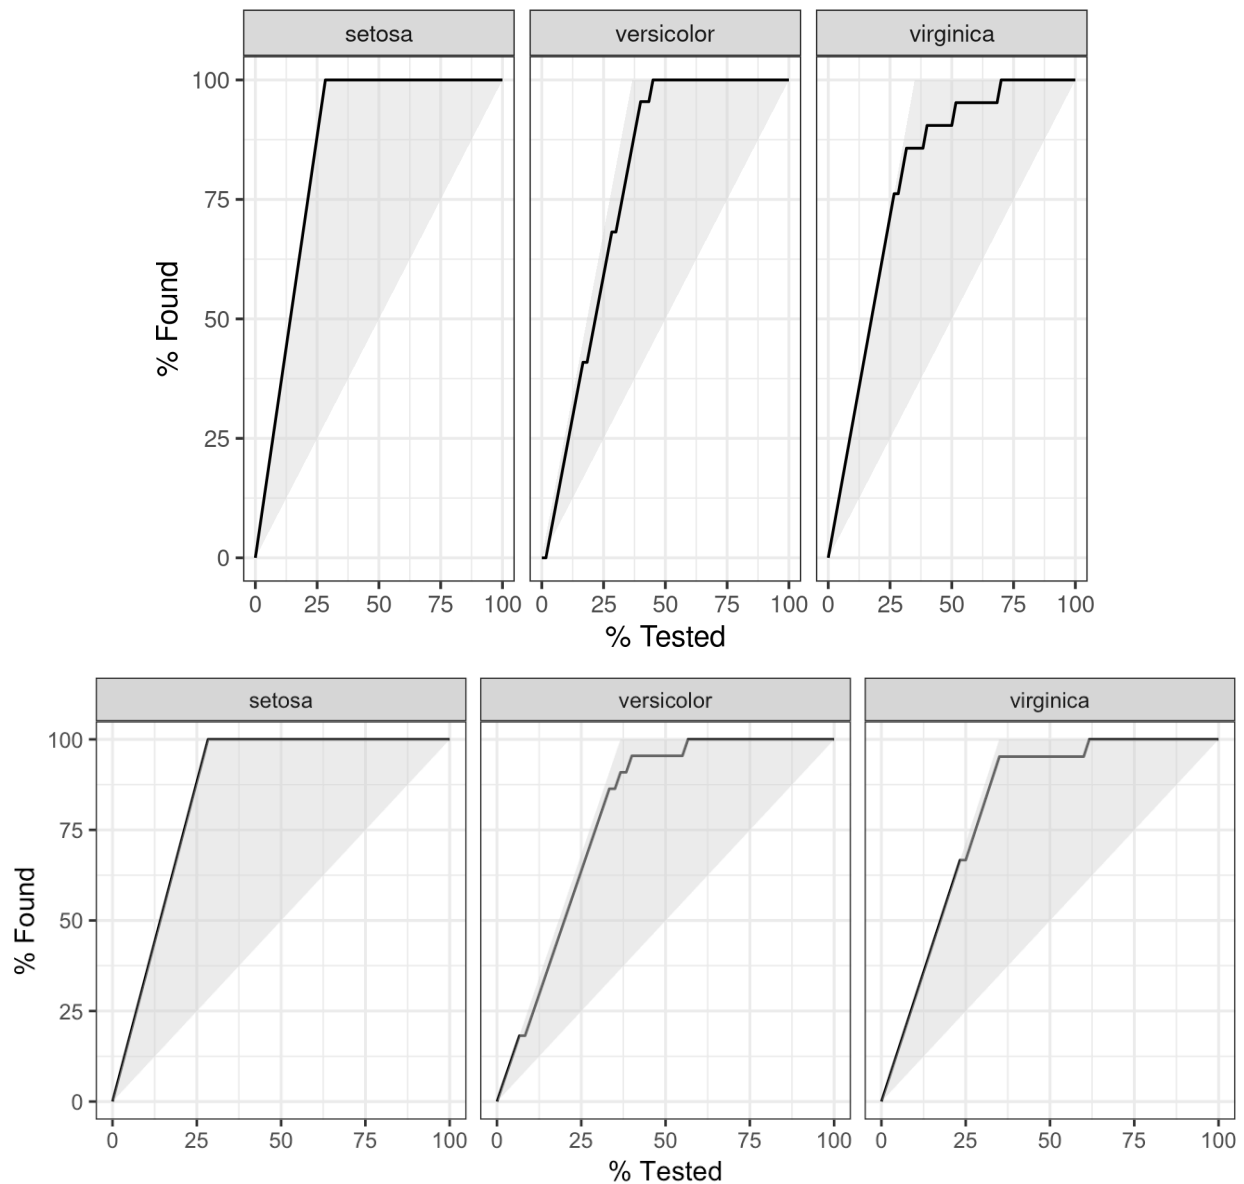
```
## Rows: 137
## Columns: 5
## $ .level           <chr> "setosa", "setosa", "setosa", "setosa", "setosa", "set~
## $ .n               <dbl> 0, 1, 2, 3, 4, 5, 7, 9, 10, 12, 13, 14, 15, 16, 17, 18~
## $ .n_events        <dbl> 0, 1, 2, 3, 4, 5, 7, 9, 10, 12, 13, 14, 15, 16, 17, 17~
## $ .percent_tested  <dbl> 0.000000, 1.666667, 3.333333, 5.000000, 6.666667, 8.33~
## $ .percent_found   <dbl> 0.000000, 5.882353, 11.764706, 17.647059, 23.529412, 2~
## Observations: 141
## Variables: 5
## $ .level           <chr> "setosa", "setosa", "setosa", "setosa", "setosa"...
```

```
## $ .n          <dbl> 0, 1, 3, 4, 5, 7, 8, 9, 10, 12, 13, 14, 15, 16, ...
## $ .n_events   <dbl> 0, 1, 3, 4, 5, 7, 8, 9, 10, 12, 13, 14, 15, 16, ...
## $ .percent_tested <dbl> 0.000000, 1.666667, 5.000000, 6.666667, 8.333333...
## $ .percent_found  <dbl> 0.000000, 5.882353, 17.647059, 23.529412, 29.411...
```

The curve methods

- include an `autoplot()` function
- that easily creates a `ggplot2` visualization.

```
iris_probs %>%
  gain_curve(Species, .pred_setosa:.pred_virginica) %>%
  autoplot()
```

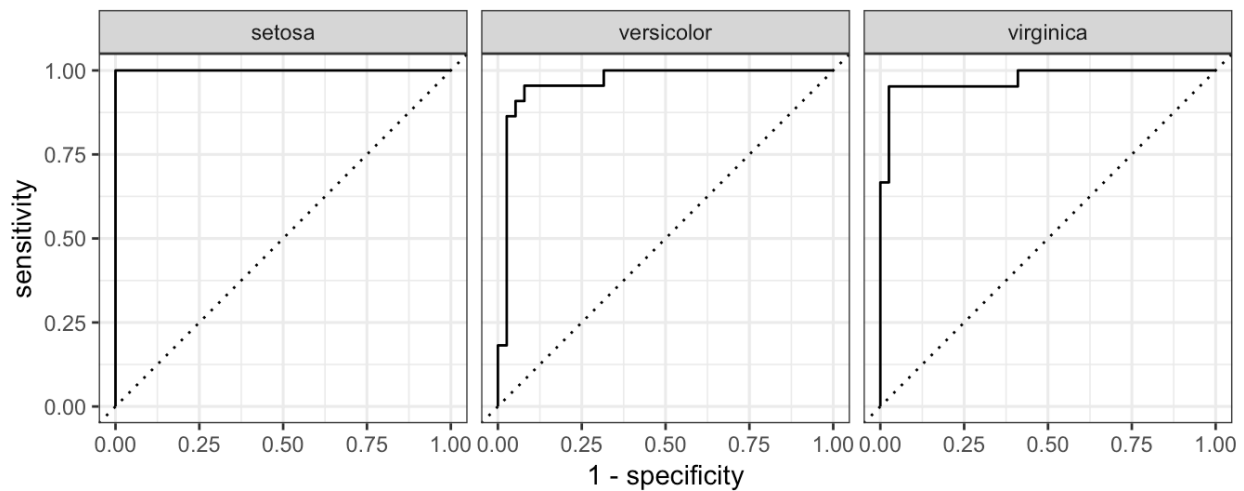



This is an example of a `roc_curve()`.

- roc means “Receiver Operating Characteristic”

- Again, because of the consistency of the interface,
 - only the function name needs to be modified;
 - even the argument values remain the same.

```
iris_probs %>%
  roc_curve(Species, .pred_setosa:.pred_virginica) %>%
  autoplot()
```



To measure the combined

- single predicted value
 - and the probability of each possible value,
- combine the two prediction modes
 - (with and without prob type).

In this example,

- using `dplyr`'s `select()`
- makes the resulting tibble easier to read.

```
predict(iris_ranger, iris_testing, type = "prob") %>%
  bind_cols(predict(iris_ranger, iris_testing)) %>%
  bind_cols(select(iris_testing, Species)) %>%
  glimpse()
```

```
## Rows: 60
## Columns: 5
```

```
## $ .pred_setosa      <dbl> 0.95980556, 0.98407143, 0.93060317, 0.92616270, 0.950~
## $ .pred_versicolor <dbl> 0.02852778, 0.00900000, 0.06939683, 0.06169444, 0.043~
## $ .pred_virginica  <dbl> 0.011666667, 0.006928571, 0.000000000, 0.012142857, 0~
## $ .pred_class      <fct> setosa, setosa, setosa, setosa, setosa, setosa, setos~
## $ Species          <fct> setosa, setosa, setosa, setosa, setosa, setosa, setos~

## Observations: 60
## Variables: 5
## $ .pred_setosa      <dbl> 0.677480159, 0.978293651, 0.783250000, 0.983972...
## $ .pred_versicolor <dbl> 0.295507937, 0.011706349, 0.150833333, 0.001111...
## $ .pred_virginica  <dbl> 0.02701190, 0.01000000, 0.06591667, 0.01491667,...
## $ .pred_class      <fct> setosa, setosa, setosa, setosa, setosa, setosa,...
## $ Species          <fct> setosa, setosa, setosa, setosa, setosa, setosa,...
```

Pipe the resulting table into `metrics()`.

In this case, specify `.pred_class` as the estimate.

```
predict(iris_ranger, iris_testing, type = "prob") %>%
  bind_cols(predict(iris_ranger, iris_testing)) %>%
  bind_cols(select(iris_testing, Species)) %>%
  metrics(Species, .pred_setosa:.pred_virginica, estimate = .pred_class)
```

```
## # A tibble: 4 x 3
##   .metric      .estimator .estimate
##   <chr>        <chr>      <dbl>
## 1 accuracy    multiclass    0.917
## 2 kap         multiclass    0.874
## 3 mn_log_loss multiclass    0.295
## 4 roc_auc     hand_till     0.970
```

```
## # A tibble: 4 x 3
##   .metric      .estimator .estimate
##   <chr>        <chr>      <dbl>
## 1 accuracy    multiclass    0.917
## 2 kap         multiclass    0.874
## 3 mn_log_loss multiclass    0.274
## 4 roc_auc     hand_till     0.980
```

7.1.2.5.10 Summary

- This end-to-end example can serve as a gentle introduction to `tidymodels`.

The number of functions,

- and options of such functions,
- were kept at a minimum for the purposes of this demonstration,
 - but there is much more that can be done with this wonderful metapackages.

Hopefully, this serves as an initial example

- you can build from
- for automating your data-driven modeling

7.1.2.6 Links

- Background on these topics

[Principal Components Analysis](#)

Principal Components Regression

Partial Least Squares Regression

Factor Analysis

Edgar Ruiz, A Gentle Introduction to tidymodels