

DSCI353-353m-453: Class 11b-f-Natural Language Processing

Profs: R. H. French, L. S. Bruckman, P. Leu, K. Davis, S. Cirlos

TAs: W. Oltjen, K. Hernandez, M. Li, M. Li, D. Colvin

06 April, 2023

Contents

0.0.0.0.1	If you are compiling to pdf, for submission,	1
11.2.1.1	Natural Language Processing	1
11.2.1.2	Working with text data	2
11.2.1.2.1	Understanding n-grams and bag-of-words	3
11.2.1.3	One-hot encoding of words and characters	5
11.2.1.3.1	Using word embeddings	6
11.2.1.3.2	Learning word embeddings with an embedding layer	6
11.2.1.3.3	Using Pretrained Word Embeddings, Word2Vec	11
11.2.1.4	Putting it all together: from raw text to word embeddings	12
11.2.1.4.1	Downloading the IMDB data as raw text	12
11.2.1.4.2	Tokenizing the data	12
11.2.1.4.3	Downloading the GLOVE embeddings	13
11.2.1.4.4	Preprocessing the embeddings	13
11.2.1.4.5	Loading the GLOVE embeddings into the model	13
11.2.1.4.6	Training and evaluating the model	14
11.2.1.5	Cites	15

0.0.0.0.1 If you are compiling to pdf, for submission,

- after having run all your code
 - So that the results of each code block is visible
- Then uncomment the FOUR lines above
 - that say `cache = TRUE, eval = TRUE, results = "markup", include = TRUE`



11.2.1.1 Natural Language Processing

- Lets explore deep-learning models that can
 - process text
 - * (understood as sequences of words or sequences of characters),
 - timeseries,
 - and sequence data in general.

The two fundamental deep-learning algorithms for sequence processing are

- Recurrent neural networks (RNNs)
- and 1D convnets, or CNNs
 - the one-dimensional version of the 2D convnets that we've used

Applications of these algorithms include the following:

- Document classification and timeseries classification,
 - such as identifying the topic of an article or the author of a book
- Timeseries comparisons,
 - such as estimating how closely related
 - two documents or two stock tickers are
- Sequence-to-sequence learning,
 - such as decoding an English sentence into French
- Sentiment analysis,
 - such as classifying the sentiment of tweets or movie reviews
 - as positive or negative
- Timeseries forecasting,
 - such as predicting the future weather at a certain location,
 - given recent weather data

These examples will focus on two narrow tasks:

- sentiment analysis on the IMDB dataset,
 - a task we approached earlier,
- and temperature forecasting.

But the techniques demonstrated for these two tasks

- are relevant to all the applications just listed,
- and many more.

11.2.1.2 Working with text data

- Text is one of the most widespread forms of sequence data.

It can be understood as either

- a sequence of characters
 - or a sequence of words,
- but it's most common to work at the level of words.

The deep-learning sequence-processing models introduced here

- can use text to produce a basic form of natural-language understanding,
- sufficient for applications
 - including document classification,
 - sentiment analysis,
 - author identification,
 - and even question answering (QA) in a constrained context.

Keep in mind throughout this chapter that

- none of these deep-learning models
 - truly understand text in a human sense;
- rather, these models can
 - map the statistical structure of written language,
 - which is sufficient to solve many simple textual tasks.

Deep learning for natural-language processing

- is pattern recognition applied to words, sentences, and paragraphs,
- in much the same way that computer vision
 - is pattern recognition applied to pixels.

Like all other neural networks,

- deep-learning models don't take as input raw text:
- they only work with numeric tensors.

Vectorizing text is the process of

- transforming text into numeric tensors.

This can be done in multiple ways:

- Segment text into words,
 - and transform each word into a vector.
- Segment text into characters,
 - and transform each character into a vector.
- Extract n-grams of words or characters,
 - and transform each n-gram into a vector.
 - (N-grams are overlapping groups
 - of multiple consecutive words or characters).

Collectively, the different units into which you can break down text

- (words, characters, or n-grams)
 - are called tokens,
- and breaking text into such tokens
 - is called tokenization.

All text-vectorization processes consist of

- applying some tokenization scheme
 - and then associating numeric vectors
 - with the generated tokens.

These vectors,

- packed into sequence tensors,
- are fed into deep neural networks.

There are multiple ways to associate a vector with a token.

We'll present two major ones:

- one-hot-encoding of tokens, and
- token embedding
 - (typically used exclusively for words, and called word embedding).

Here we'll explain these techniques and

- show how to use them to go
- from raw text to a tensor that
- you can send to a Keras network.

11.2.1.2.1 Understanding n-grams and bag-of-words

- **Word n-grams** are groups of N (or fewer) consecutive words
 - that you can extract from a sentence.
 - The same concept may also be applied to characters
 - * instead of words.

Here's a simple example.

- Consider the sentence "The cat sat on the mat."
- It may be decomposed into the following set of **2-grams**:
 - {"The", "The cat", "cat", "cat sat", "sat",

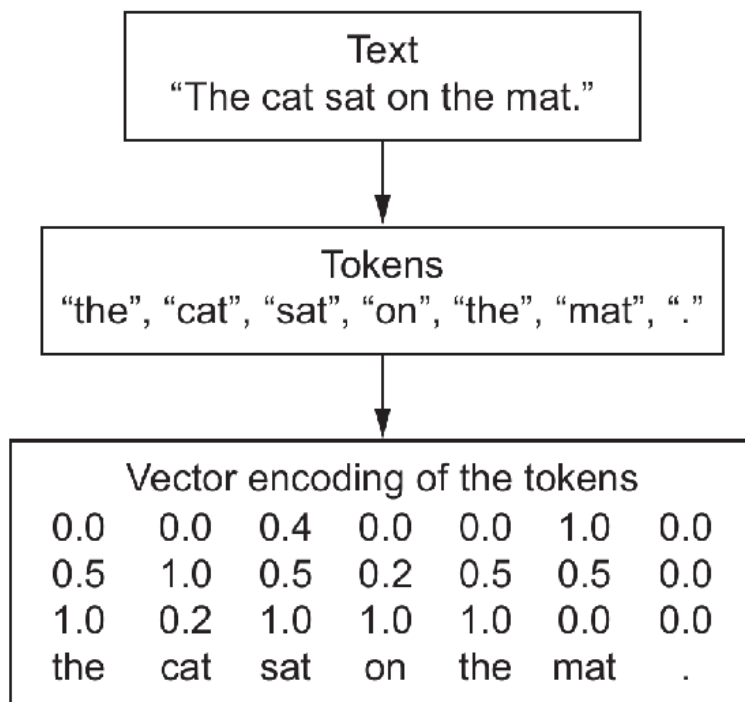


Figure 6.1 From text to tokens to vectors

Figure 1: From text to tokens to vectors

- “sat on”, “on”, “on the”, “the”, “the mat”, “mat”}
- It may also be decomposed into the following set of 3-grams:
 - {“The”, “The cat”, “cat”, “cat sat”, “The cat sat”,
 - “sat”, “sat on”, “on”, “cat sat on”, “on the”, “the”,
 - “sat on the”, “the mat”, “mat”, “on the mat”}
- Such a set is called a
 - bag-of-2-grams or
 - bag-of-3-grams, respectively.

The term bag here refers to the fact that you’re dealing with

- a set of tokens rather than a list or sequence:
 - the tokens have no specific order.
- This family of tokenization methods is called **bag-of-words**.

Because bag-of-words

- isn’t an order-preserving tokenization method
 - (the tokens generated are understood as a set, not a sequence,
 - and the general structure of the sentences is lost),
- it tends to be used in shallow language-processing models
 - rather than in deep-learning models.

Extracting n-grams is a form of feature engineering,

- and deep learning does away with this kind of rigid, brittle approach,
 - replacing it with hierarchical feature learning.
- One-dimensional convnets and recurrent neural networks,
 - introduced later in this chapter,
 - are capable of learning representations for groups of words and characters
 - without being explicitly told about the existence of such groups,
 - by looking at continuous word or character sequences.

For this reason, we won’t cover n-grams any further in this book.

- But do keep in mind that they’re a powerful,
 - unavoidable feature-engineering tool
- when using lightweight, shallow, text-processing models
 - such as logistic regression and random forests.

11.2.1.3 One-hot encoding of words and characters

- One-hot encoding is
 - the most common, most basic way to turn a token into a vector.

You saw it in action in

- the initial IMDB and Reuters examples in chapter 3
- (done with words, in that case).

It consists of associating a unique integer index

- with every word
- and then turning this integer index i
 - into a binary vector of size N
 - (the size of the vocabulary);
- the vector is all zeros
 - except for the i th entry, which is 1.

Of course, one-hot encoding can be done at the character level, as well.

To unambiguously drive home

- what one-hot encoding is and
- how to implement it,
- listings 6.1 and 6.2 in DLwR show two toy examples:
 - one for words, the other for characters.

Note that Keras has built-in utilities for

- doing one-hot encoding of text
 - at the word level or character level,
 - starting from raw text data.
- You should use these utilities, because they take care of
 - a number of important features
 - such as stripping special characters from strings
 - and only taking into account the N most common words in your dataset
 - (a common restriction, to avoid dealing with very large input vector spaces).

11.2.1.3.1 Using word embeddings

- Another popular and powerful way to associate a vector with a word
 - is the use of dense word vectors,
 - * also called word embeddings.
 - Whereas the vectors obtained through one-hot encoding
 - * are binary, sparse (mostly made of zeros),
 - * and very high-dimensional (same dimensionality
 - as the number of words in the vocabulary),
 - word embeddings are low-dimensional floating-point vectors
 - * (that is, dense vectors, as opposed to sparse vectors)
 - * see figure 6.2.

Unlike the word vectors obtained via one-hot encoding,

- **word embeddings** are learned from data. It's common to see word embeddings
 - that are 256-dimensional, 512-dimensional, or 1,024-dimensional,
 - when dealing with very large vocabularies.
- On the other hand, one-hot encoding words
 - generally leads to vectors that are 20,000-dimensional or greater
 - (capturing a vocabulary of 20,000 tokens, in this case).
- So, word embeddings pack more information into far fewer dimensions.

There are two ways to obtain word embeddings:

- Learn word embeddings jointly with the main task you care about
 - (such as document classification or sentiment prediction).
 - In this setup, you start with random word vectors and
 - then learn word vectors in the same way
 - * you learn the weights of a neural network.
- Load into your model word embeddings that were precomputed
 - using a different machine-learning task
 - * than the one you're trying to solve.
 - These are called pretrained word embeddings.


11.2.1.3.2 Learning word embeddings with an embedding layer

- The simplest way to



One-hot word vectors:

- Sparse
- High-dimensional
- Hardcoded



Word embeddings:

- Dense
- Lower-dimensional
- Learned from data

Figure 6.2 Whereas word representations obtained from one-hot encoding or hashing are sparse, high-dimensional, and hardcoded, word embeddings are dense, relatively low-dimensional, and learned from data.

Figure 2: Figure 6.2 Whereas word representations obtained from one-hot encoding or hashing are sparse, high-dimensional, and hardcoded, word embeddings are dense, relatively low-dimensional, and learned from data.

- associate a dense vector with a word
 - * is to choose the vector at random.
- The problem with this approach is that
 - * the resulting embedding space has no structure:
 - * for instance, the words accurate and exact
 - may end up with completely different embeddings,
 - even though they’re interchangeable in most sentences.
- It’s difficult for a deep neural network
 - * to make sense of such a noisy, unstructured embedding space.

To get a bit more abstract,

- the geometric relationships between word vectors
 - should reflect the semantic relationships between these words.
- Word embeddings are meant to map human language
 - into a geometric space.
- For instance, in a reasonable embedding space,
 - you would expect synonyms to be embedded into similar word vectors;
 - and in general, you would expect the geometric distance
 - * (such as L2 distance)
 - between any two word vectors
 - * to relate to the semantic distance
 - * between the associated words
 - (words meaning different things
 - * are embedded at points far away from each other,
 - * whereas related words are closer).
- In addition to distance,
 - you may want specific directions in the embedding space to be meaningful.

To make this clearer, let’s look at a concrete example.

In figure 6.3, four words are embedded on a 2D plane:

- cat, dog, wolf, and tiger.
- With the vector representations we chose here,
 - some semantic relationships between these words
 - can be encoded as geometric transformations.
- For instance, the same vector allows us
 - to go from cat to tiger
 - and from dog to wolf:
- This vector could be interpreted as a
 - “from pet to wild animal” vector.
- Similarly, another vector lets us go
 - from dog to cat and
 - from wolf to tiger,
- which could be interpreted as a
 - “from canine to feline” vector.

In real-world word-embedding spaces,

- common examples of meaningful geometric transformations are
 - “gender” vectors and
 - “plural” vectors.
- For instance, by adding a “female” vector
 - to the vector “king,”
 - we obtain the vector “queen.”
- By adding a “plural” vector,




Figure 6.3 A toy example of a word-embedding space

Figure 3: Fig. 6.3 a toy example of a word-embedding space

- we obtain “kings.”
- Word-embedding spaces typically feature
 - thousands of such interpretable and potentially useful vectors.

Is there some ideal word-embedding space

- that would perfectly map human language and
 - could be used for any natural-language-processing task?
- Possibly, but we have yet to compute anything of the sort.

Also, there is no such a thing as human language

- there are many different languages,
 - and they aren’t isomorphic,
- because a language is the reflection of
 - a specific culture and
 - a specific context.

But more pragmatically,

- what makes a good word-embedding space
 - depends heavily on your task:
- the perfect word-embedding space
 - for an English-language movie-review sentiment-analysis model
- may look different from the perfect embedding space
 - for an English-language legal-document-classification model,
 - because the importance of certain semantic relationships
 - varies from task to task.

It's thus reasonable to learn a new embedding space with every new task.

Fortunately, backpropagation makes this easy,

- and Keras makes it even easier.
- It's about learning the weights of a layer
 - using `layer_embedding`.

The embedding layer takes at least two arguments:

- the number of possible tokens (here, 1,000) and
 - the dimensionality of the embeddings (here, 64).
- `layer_embedding` is best understood as a dictionary
 - that maps integer indices (which stand for specific words)
 - to dense vectors.
- It takes integers as input,
 - it looks up these integers in an internal dictionary, and
 - it returns the associated vectors.
- It's effectively a dictionary lookup (see figure 6.4).

Word index \longrightarrow Embedding layer \longrightarrow Corresponding word vector

Figure 6.4 An embedding layer

Figure 4: Fig. 6.4 an embedding layer

An embedding layer takes as input

- a 2D tensor of integers,
 - of shape (samples, `sequence_length`),
 - where each entry is a sequence of integers.
- It can embed sequences of variable lengths:
 - for instance, you could feed into the embedding layer in listing 6.5
 - batches with shapes (32, 10)
 - * (batch of 32 sequences of length 10)
 - or (64, 15)
 - * (batch of 64 sequences of length 15).
- All sequences in a batch must have the same length,
 - though (because you need to pack them into a single tensor),
 - so sequences that are shorter than others should be padded with zeros,
 - and sequences that are longer should be truncated.

This layer returns a 3D floating-point tensor

- of shape (samples, `sequence_length`, `embedding_dimensionality`).
- Such a 3D tensor can then be processed by
 - an RNN layer or a 1D convolution layer
 - (both will be introduced in the following sections).

When you instantiate an embedding layer,

- its weights (its internal dictionary of token vectors) are initially random,
 - just as with any other layer.
- During training, these word vectors
 - are gradually adjusted via backpropagation,
 - structuring the space into something the downstream model can exploit.

Once fully trained, the embedding space

- will show a lot of structure
 - a kind of structure specialized for the specific problem
 - for which you’re training your model.

11.2.1.3.3 Using Pretrained Word Embeddings, Word2Vec

- Sometimes, you have so little training data available
 - that you can’t use your data alone
 - * to learn an appropriate task-specific embedding of your vocabulary.
 - What do you do then?

Instead of learning word embeddings

- jointly with the problem you want to solve,
- you can load embedding vectors from a precomputed embedding space
 - that you know is highly structured and
 - exhibits useful properties
 - that captures generic aspects of language structure.

<https://en.wikipedia.org/wiki/Word2vec>

And the 2013 Word2Vec paper cited below [1]

The rationale behind

- using pretrained word embeddings in natural-language processing
- is much the same as for using pretrained convnets in image classification:
- you don’t have enough data available
 - to learn truly powerful features on your own,
 - but you expect the features that you need to be fairly generic
 - that is, common visual features or semantic features.
- In this case, it makes sense to reuse features
 - learned on a different problem.

Such word embeddings are generally computed

- using word-occurrence statistics
 - (observations about what words co-occur in sentences or documents),
 - using a variety of techniques,
 - * some involving neural networks, others not.
- The idea of a dense, low-dimensional embedding space for words,
 - computed in an unsupervised way,
 - was initially explored by Bengio et al. in the early 2000s,
- But it only started to take off in research and industry applications
 - after the release of one of the most famous and successful
 - word-embedding schemes: [the word2vec algorithm](#),
 - developed by Tomas Mikolov at Google in 2013.

Word2vec dimensions

- capture specific semantic properties,
- such as gender.

There are various precomputed databases of word embeddings

- that you can download and use in a Keras embedding layer.
- Word2vec is one of them.
- Another popular one is called [Global Vectors for Word Representation \(GloVe\)](#)

- which was developed by Stanford researchers in 2014.
- GloVe is based on factorizing a matrix
 - of word co-occurrence statistics.
 - Its developers have made available precomputed embeddings
 - for millions of English tokens,
 - obtained from Wikipedia data
 - and Common Crawl data.

Let's look at how you can get started

- using GloVe embeddings in a Keras model.
- The same method is valid for Word2vec embeddings
 - or any other word-embedding database.

11.2.1.4 Putting it all together: from raw text to word embeddings

- You'll use a model similar to the one we just went over:
 - embedding sentences in sequences of vectors,
 - flattening them, and
 - training a dense layer on top.

But you'll do so using pretrained word embeddings;

- and instead of using the pretokenized IMDB data packaged in Keras,
- you'll start from scratch by downloading the original text data.

11.2.1.4.1 Downloading the IMDB data as raw text

- Note: Since in DLR First Edition, which we are using
 - Is written for TF version 1.15
 - * Not TF2 version 2.7
 - So the code in our book doesn't work
 - * DLR 2nd Edition is coming out soon
 - * And will have working code
 - So I don't include code in the practicum.
 - * Since it won't run

First, download the raw IMDB dataset

- from <http://mng.bz/0tIo>.
- Uncompress it.

Now, let's collect the individual training reviews

- into a list of strings,
 - one string per review.

You'll also collect the review labels

- (positive/negative)
- into a labels list.

DLwR: Listing 6.8

11.2.1.4.2 Tokenizing the data

- Let's vectorize the text and prepare a training and validation split,
 - using the concepts introduced earlier in this section.
 - Because pretrained word embeddings are meant to be particularly useful

- * on problems where little training data is available
- * (otherwise, task-specific embeddings are likely to outperform them),
- we’ll add the following twist: restricting the training data to the first 200 samples.
 - * So you’ll learn to classify movie reviews
 - after looking at just 200 examples.

DLwR: Listing 6.9

11.2.1.4.3 Downloading the GLOVE embeddings

- Go to <https://nlp.stanford.edu/projects/glove>,
 - and download the precomputed embeddings from 2014 English Wikipedia.
 - It’s an 822 MB zip file called glove.6B.zip,
 - * containing 100-dimensional embedding vectors
 - * for 400,000 words (or nonword tokens).
 - Unzip it.

11.2.1.4.4 Preprocessing the embeddings

- Let’s parse the unzipped file (a .txt file)
 - to build an index that maps words (as strings)
 - to their vector representation (as number vectors).

DLwR: Listing 6.10

Next, you’ll build an embedding matrix

- that you can load into an embedding layer.
- It must be a matrix of shape (max_words, embedding_dim),
 - where each entry i contains the embedding_dim-dimensional vector
 - for the word of index i
 - in the reference word index (built during tokenization).
- Note that index 1 isn’t supposed to stand for any word or token
 - it’s a placeholder.

DLwR: Listing 6.11

Defining a model

You’ll use the same model architecture as before.

DLwR: Listing 6.12

11.2.1.4.5 Loading the GLOVE embeddings into the model

- The embedding layer has a single weight matrix:
 - a 2D float matrix
 - * where each entry i is the word vector
 - * meant to be associated with index i.
 - Simple enough.
 - * Load the GloVe matrix you prepared into the embedding layer,
 - * the first layer in the model.

DLwR: Listing 6.13

Additionally, you’ll freeze the weights of the embedding layer,

- following the same rationale you’re already familiar with
- in the context of pretrained convnet features:

- when parts of a model are pretrained (like your embedding layer)
- and parts are randomly initialized (like your classifier),
- the pretrained parts shouldn't be updated during training,
 - to avoid forgetting what they already know.
- The large gradient updates triggered by the randomly initialized layers
 - would be disruptive to the already-learned features.

11.2.1.4.6 Training and evaluating the model

- Compile and train the model.

DDLwR: Listing 6.14

Now, plot the model's performance over time (see figure 6.5).

DLwR: Listing 6.15




Figure 6.5 Training and validation metrics when using pretrained word embeddings

Figure 5: Training and Validation Metrics when using pretrained word embeddings

The model quickly starts overfitting,

- which is unsurprising given the small number of training samples.
- Validation accuracy has high variance for the same reason,
 - but it seems to reach the high 50s.

Note that your mileage may vary:

- because you have so few training samples,
- performance is heavily dependent on exactly which 200 samples you choose
 - and you’re choosing them at random.
- If this works poorly for you,
 - try choosing a different random set of 200 samples,
 - * for the sake of the exercise
 - (in real life, you don’t get to choose your training data).

You can also train the same model

- without loading the pretrained word embeddings and
 - without freezing the embedding layer.
- In that case, you’ll learn a task-specific embedding of the input tokens,
 - which is generally more powerful than pre-trained word embeddings
 - when lots of data is available.
- But in this case, you have only 200 training samples.

11.2.1.5 Cites

- This notebook is based on Chapter 6, [Deep Learning with R](#).

[1] T. Mikolov, K. Chen, G. S. Corrado, and J. Dean, “Efficient estimation of word representations in vector space.” 2013 [Online]. Available: <http://arxiv.org/abs/1301.3781>